

Contents to be covered

- Need of a database
- Different types of database,
- Relational vs Non-relational database



- In our day to day life as an end user we discover online applications which try to make our lives more convenient. As we know about these just get ourselves register ourselves for that application.
- After the one-time registration, whenever we want to use that app again, we just need to log in with user name and password sometime just a mobile based OTP, on the other end the application system automatically stores all our data that was provided during the registration process and submitted subsequently.



- All this is possible because of the database in which all our information or data gets stored when we register for any application. It may be our Facebook, Whatsapp, Instagram, email services or airline or hospitality application etc.
- when we browse various ecommerce websites and view millions of product items available on various online shopping applications like Amazon, Flipkart etc or post our status, photos videos on Facebook, twitter, youtube etc to let others view them, it is only possible because of the databases.



 At present all the applications weather a desktop or mobile one, databases acts as a key component for collecting, storing and retrieving of these data in a safe and efficient manner with outmost security and privacy. The task of choosing the right option is tedious process.



- In physical world and go back and analyses the past, before the implementation of computers and databases, departments like electricity boards, insurance offices, Banks, Public health departments etc stores a large heaps of files containing the data of all of their users/ customers.
- It was a quite troublesome to get details pertaining to a customer, e.g. consumption
 of electricity, payments made or pending towards bills generated, etc, if the names
 were not listed alphabetically or data in any other order like in ascending order of
 account number.
- Even the management of the physical data was done is just like it's done in computer databases but not in all cases. And if the data is not present or not in an organized form, then the getting any information from this physical data is taking so long time.



- The data stored in a computer database is organized in such a way that the stored data helps in easy retrieval of information along with storage and editing. The data stored in physical files can get lost when the papers of these files get older and, hence, get destroyed.
- But in a database, we can store as long as we want in a very small space with security and this computer based Data will get lost only if the system crashes. To overcome this, we keep a backup that may be in form of backup data on external storage or in an online mode depending upon the application requirement as well as the budget.



- Depending upon the usage requirements, there are following types of databases available in the market –
 - Centralized database.
 - Distributed database.
 - Personal database.
 - NoSQL database.
 - Operational database.
 - Relational database.
 - Cloud database.
 - Object-oriented database.
 - Graph database.



1. Centralized Database

The information(data) is stored at a centralized location and the users from different locations can access this data. This type of database contains application procedures that help the users to access the data even from a remote location.

Various kinds of authentication procedures are applied for the verification and validation of end users, likewise, a registration number is provided by the application procedures which keep a track and record of data usage. The local area office handles this thing.



2. Distributed Database

Just opposite of the centralized database concept, the distributed database has contributions from the common database as well as the information captured by local computers also. The data is not at one place and is distributed at various sites of an organization. These sites are connected to each other with the help of communication links which helps them to access the distributed data easily.



- You can imagine a distributed database as a one in which various portions of a database are stored in multiple different locations (physical) along with the application procedures which are replicated and distributed among various points in a network.
- There are two kinds of distributed database, viz. homogenous and heterogeneous.
 The databases which have same underlying hardware and run over same operating systems and application procedures are known as homogeneous DDB, for eg. All physical locations in a DDB. Whereas, the operating systems, underlying hardware as well as application procedures can be different at various sites of a DDB which is known as heterogeneous DDB.



3. Personal Database

Data is collected and stored on personal computers which is small and easily manageable. The data is generally used by the same department of an organization and is accessed by a small group of people.



4.NoSQL Database

These are used for large sets of distributed data. There are some big data performance issues which are not effectively handled by relational databases, such kind of issues are easily managed by NoSQL databases. There are very efficient in analyzing large size unstructured data that may be stored at multiple virtual servers of the cloud.



5.Operational Database

Information related to operations of an enterprise is stored inside this database. Functional lines like marketing, employee relations, customer service etc. require such kind of databases.



6.Relational Databases

These databases are categorized by a set of tables where data gets fit into a pre-defined category. The table consists of rows and columns where the column has an entry for data for a specific category and rows contains instance for that data defined according to the category.

The Structured Query Language (SQL) is the standard user and application program interface for a relational database.



There are various simple operations that can be applied over the table which makes these databases easier to extend, join two databases with a common relation and modify all existing applications.

7.Cloud Databases

Now a day, data has been specifically getting stored over clouds also known as a virtual environment, either in a hybrid cloud, public or private cloud. A cloud database is a database that has been optimized or built for such a virtualized environment.



- There are various benefits of a cloud database, some of which are the ability to pay
 for storage capacity and bandwidth on a per-user basis, and they provide scalability
 on demand, along with high availability.
- A cloud database also gives enterprises the opportunity to support business applications in a software-as-a-service deployment.

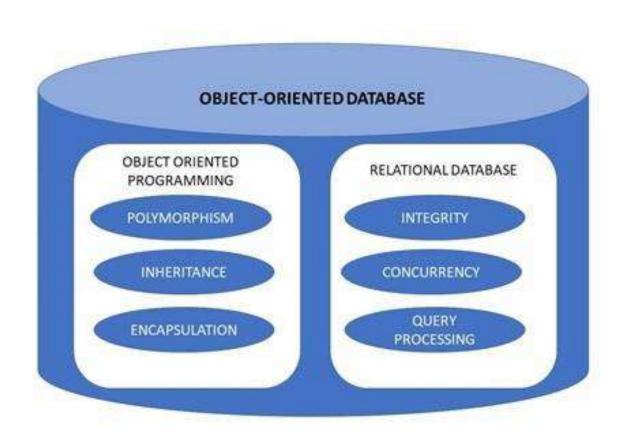


8.Object-Oriented Databases

An object-oriented database is a collection of object-oriented programming and relational database. There are various items which are created using object-oriented programming languages like C++, Java which can be stored in relational databases, but object-oriented databases are well-suited for those items.

An object-oriented database is organized around objects rather than actions, and data rather than logic. For example, a multimedia record in a relational database can be a definable data object, as opposed to an alphanumeric value.



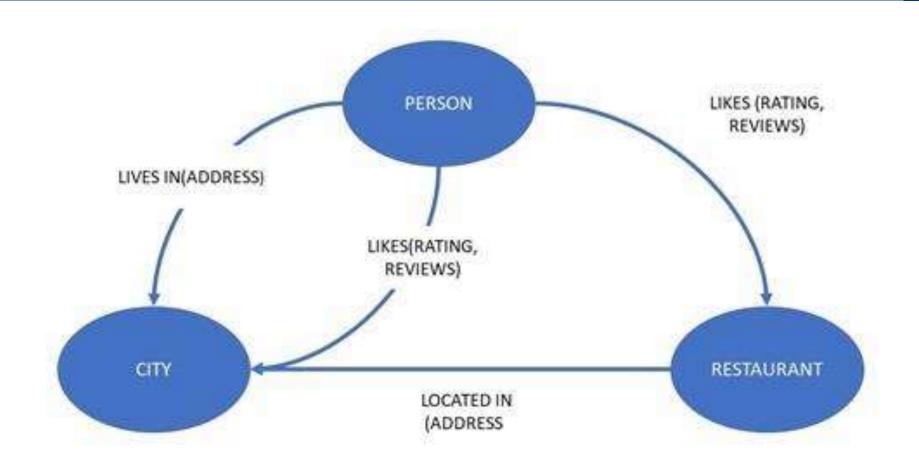




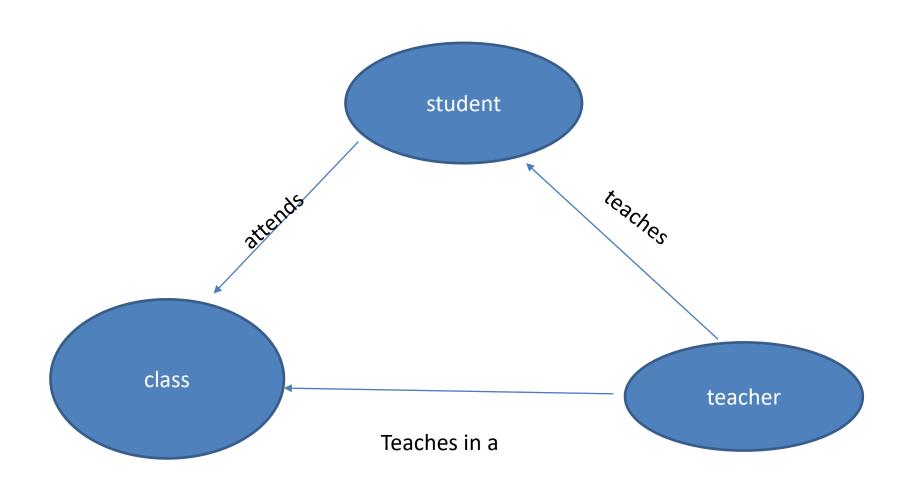
9. Graph Databases

- The graph is a collection of nodes and edges where each node is used to represent an entity and each edge describes the relationship between entities. A graph-oriented database, or graph database, is a type of NoSQL database that uses graph theory to store, map and query relationships.
- Graph databases are basically used for analyzing interconnections.
 For example, companies might use a graph database to mine data about customers from social media.











Relational vs. non-relational database

- Relational databases are the most prominent modern types of databases. Oracle, MySQL, PostgreSQL, or SQL Server are some examples of RDBMS. To access and manipulate the data in RDBMS, SQL (Structured Query Language) is required.
- SQL has well-established standards and allows your data to be easily portable. In relational database data is stored in tables consisting of columns and rows.
- Every row represents an individual record, and a column stands for a field with a data type assigned to it. Primary and foreign keys are used to link Tables having related information.



• Student

Rollno(PK)	Name	Phone number
1	Aaa	56654545656
2	Bbb	546545454

Rollno(FK)	marks	result
1	568	Pass
2	345	pass



Relational vs. non-relational database

- Now-a-days non-relational databases are also gaining popularity and being used so commonly. The main reason for this is the growing need for unstructured data storage like data from Social Media etc. In the era of big data, diversified information is used. diversity.
- Data now also mean images, videos, and even posts on social media networks and to work with non-tabular data, nonrelational database is needed. These are also referred as NoSQL databases as these do not support SQL queries.



RDBMS

Student rollno	Name	Marks
1	Kkk	56
2	LII	85

Non-relational

- {"rollno":1,"name":rrr","marks":67}
- {"rollno":2,"address":"mumbai","phoneno":2343423424}
- {"rollno":3,"hobbies":"cricket","awards":"volleyball"}



 Relational databases built around relational algebra and tuple relational calculus and are optimized for writes, consistency and availability.

Advantages:

- Simplicity, ease of data retrieval, data integrity,
- Single uniform language (DDL) for different roles (developer, user, DBA).
- Single standardized language for different RDBMS.
- sticks to ACID principles (atomicity, consistency, isolation, durability), and thus ensures stability, security, and predictability both for the entire database and every individual transaction



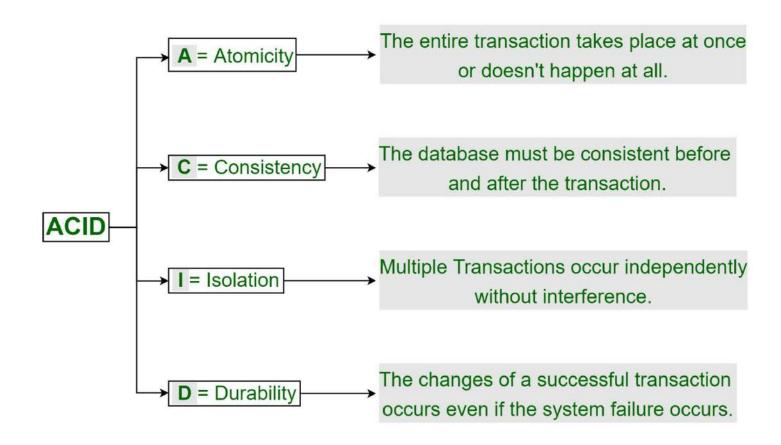
ACID Properties in DBMS

 A <u>transaction</u> is a single logical unit of work which accesses and possibly modifies the contents of a database. Transactions access data using read and write operations.

In order to maintain consistency in a database, before and after the transaction, certain properties are followed. These are called **ACID** properties.



ACID Properties in DBMS





Relational databases

Disadvantages:

- expensive to set up and maintain
- have limits to field lengths. This can be cumbersome for storing a large amount of information in one field.
- very difficult to scale as much as a database grows larger
- Multiple databases can easily become islands of information and difficult to connect the databases where they can talk to each other.



 Non-relational databases are schema-free and built on distributed systems, which makes it easy to scale and shard.
 These are optimized for reads.

Advantages:

- Flexible to store large volumes of structured, semi-structured, and unstructured data.
- Can scale out architecture efficiently without expensive overhead.
- Available as open source and free of cost in many cases



• Disadvantages:

- ACID transactions are not supported by non-relational databases. Instead, they rely on "eventual consistency". The performance benefits of these databases mean there's a cost of consistency.
- There are many types of NoSQL databases but there is almost no uniformity among them.
- No specific programming interface to the different databases. Each one varies in query language with another.
- Not all non-relational databases are good at automating the process of sharding, or spreading the database across multiple nodes.



Contents to be covered

- Introduction to NoSQL Databases
- NoSQL features,
- Different types of NoSQL databases
- Introduction to MongoDB,
- MongoDB architecture,



NoSQL

- NoSQL is a non-relational DBMS, that does not require a fixed schema, avoids joins, and is easy to scale. The purpose of using a NoSQL database is for distributed data stores with large amount of data storage needs.
- NoSQL is used for Big data and real-time web apps. For example, companies like Twitter, Facebook, Google collect terabytes of user data every single day.
- NoSQL database stands for "Not Only SQL" or "Not SQL." Though a better term would be "NoREL", NoSQL caught on. Carl Strozz introduced the NoSQL concept in 1998.



RDBMS

Student name	Rollno	Phone number
Aaa	1	98564758 96
Bbb	2	56465864 65

Rollno	Name	marks
1	Aaa	45
2	Bbb	56

NoSQL

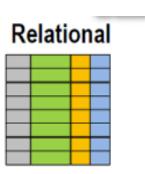
```
name:"aaa",
Rollno:1,
Phoneno.:34534534554,
Academic: {
         marks:45 },
```



NoSQL

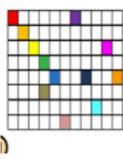
- Traditional RDBMS uses SQL syntax to store and retrieve data for further insights.
- Instead, a NoSQL database system encompasses a wide range of database technologies that can store structured, semi-structured, unstructured and polymorphic data.





NoSQL Databases

Column-Family



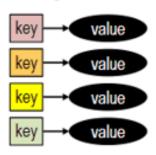
Graph



Document



Key-Value





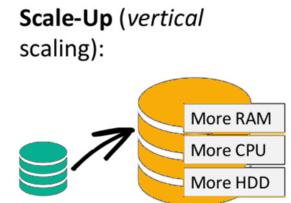
Why NoSQL?

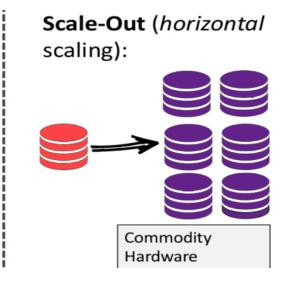
- The concept of NoSQL databases became popular with Internet giants like Google, Facebook, Amazon, etc. who deal with huge volumes of data. The system response time becomes slow when you use RDBMS for massive volumes of data.
- To resolve this problem, we could "scale up" our systems by upgrading our existing hardware. This process is expensive.



Why NoSQL?

 The alternative for this issue is to distribute database load on multiple hosts whenever the load increases. This method is known as "scaling out."







Why NoSQL?

 NoSQL database is non-relational, so it scales out better than relational databases as they are designed with web applications in mind.



History of NoSQL Databases

- 1998- Carlo Strozzi use the term NoSQL for his lightweight, open-source relational database
- 2000- Graph database Neo4j is launched
- 2004- Google BigTable is launched
- 2005- CouchDB is launched
- 2007- The research paper on Amazon Dynamo is released
- 2008- Facebooks open sources the Cassandra project
- 2009- The term NoSQL was reintroduced



Non-relational

- NoSQL databases never follow the relational model
- Never provide tables with flat fixed-column records
- Doesn't require data normalization
- No complex features like query languages, query planners, referential integrity joins, ACID



Rollno	Name	Marks	Phone number
1	Rahul	45	9888254586, 4568954254
2	Smita	56	9856745896, 2568935984

Rollno	Name	Marks	Phone number
1	Rahul	45	9888254586
1	Rahul	45	4568954254
2	Smita	56	9856745896
2	Smita	56	2568935984



Schema-free

- NoSQL databases are either schema-free or have relaxed schemas
- Do not require any sort of definition of the schema of the data
- Offers heterogeneous structures of data in the same domain



Simple API

- Offers easy to use interfaces for storage and querying data provided
- APIs allow low-level data manipulation & selection methods
- Mostly used no standard based query language



Distributed

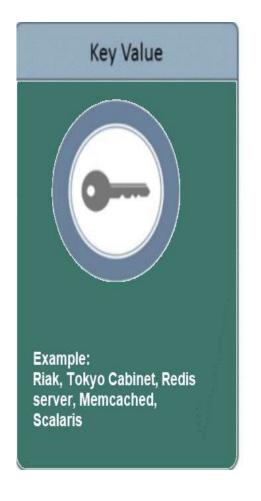
- Multiple NoSQL databases can be executed in a distributed fashion
- Offers auto-scaling
- Often ACID concept can be sacrificed for scalability and throughput



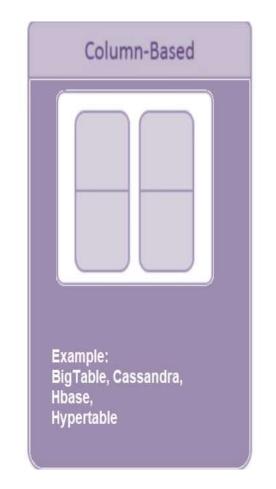
Types of NoSQL databases

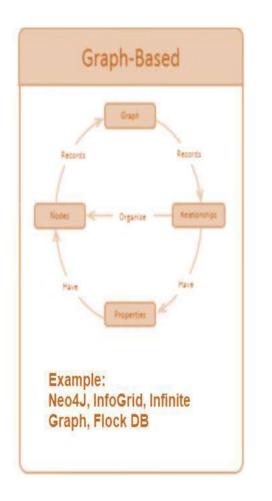
- There are mainly four categories of NoSQL databases. Each of these categories has its unique attributes and limitations. No specific database is better to solve all problems. select a database based on your product needs.
 - Key-value Pair Based
 - Column-oriented Graph
 - Graphs based
 - Document-oriented













Types of NoSQL databases

Key Value Pair Based

- Data is stored in key/value pairs. It is designed in such a way to handle lots of data and heavy load.
- Key-value pair storage databases store data as a hash table where each key is unique, and the value can be a JSON, BLOB(Binary Large Objects), string, etc.
- For example, a key-value pair may contain a key like "student name" associated with a value like "rahul".



Key Value Pair Based

 It is one of the most basic types of NoSQL databases. This kind of NoSQL database is used as a collection, dictionaries etc. Key value stores help the developer to store schema-less data. They work best for shopping cart contents.

```
    Shopping cart
        {
                  product name : soap
                  Product qty : 5
                  Product price : 15 }
```

 Redis, Dynamo, Riak are some examples of key-value store Databases.



Types of NoSQL databases

Column-based

 Column-oriented databases work on columns and are based on BigTable paper by Google. Every column is treated separately. Values of single column databases are stored contiguously.



Column-based

Column family							
R	Column name						
0	Key	Key	Key				
W k	Value Value						
е	Column name						
У	Key	Key	Key				
	Value	Value	value				

student									
ID:1	Personal data								
	Name	Gende	Gender		Ac	Address		Phone no.	
	Rahul	Male		23	dsfsd		895768		
	Academic r	ecord							
	First semester		Second semester		Third semes	st	total		
	75		85			56		Sum(
ID:2	Personal data								
	Name		Gender ag		ge				
	Raj		Male		34				
	Academic r								
	First semes	ter Secon		d semester		Third semester			
	75		45			59			



Column-based

- They deliver high performance on aggregation queries like SUM,
 COUNT, AVG, MIN etc. as the data is readily available in a column.
- Column-based NoSQL databases are widely used to manage data warehouses, business intelligence, Library card catalogs,
- HBase, Cassandra, Hypertable are examples of column based database.



C language						
Name	Author	Price	Qty	Id		
let us c	Yashwant	88	5	A01		
C++						
Name	AUTHOR	QTY	ID			
C++	ROBERT	4	dasfaf			



Types of NoSQL databases

Document-Oriented:

– Document-Oriented NoSQL DB stores and retrieves data as a key value pair but the value part is stored as a document. The document is stored in JSON or XML formats. The value is understood by the DB and can be queried.



Document-Oriented:

RDBMS

Col1	Col2	Col3	Col4
Data	Data	Data	Data
Data	Data	Data	Data
Data	Data	Data	Data

Name	Age	Design ation	salary	incent ive
Rahul	23	Mana ger	2000	
Vijay	25	Officer	3000	
Raj	25	Accou nt officer	4000	4000

Document1

{
Name: "rahul",
Age:23,
Designation:"manager",
Salary:2000
}

Document2

{
Name: "vijay",
Age:25
Designation:
"officer",
Salary:3000
}

Document3

{
Name:"raj"
Age:25
Designation:"accoun
t officer"
Salary:4000
Incentive:4000
}



Relational Vs. Document

- In the left diagram there are rows and columns, and in the right, a document database which has a similar structure to JSON.
- Now for the relational database, there is the need to know what columns are there and so on. However, for a document database, there is data store like JSON object. There is no requirement to define which make it flexible.



Relational Vs. Document

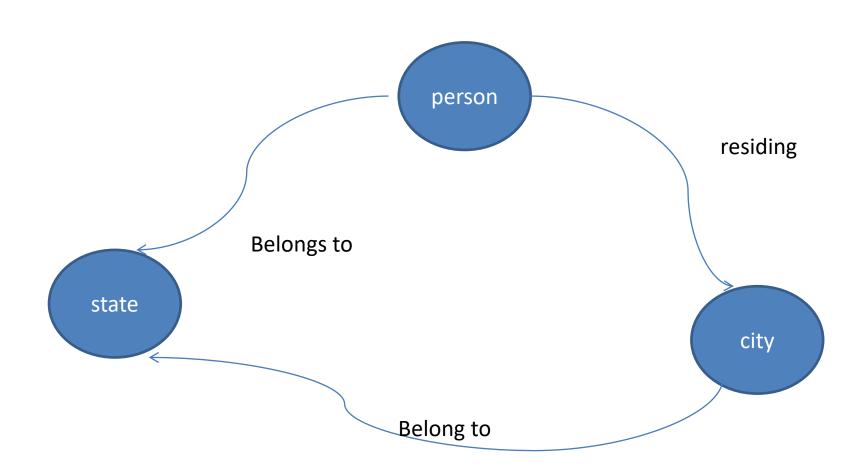
- The document type is mostly used for blogging platforms & ecommerce applications. It should not use for complex transactions which require multiple operations or queries against varying aggregate structures.
- Amazon SimpleDB, CouchDB, MongoDB, Lotus Notes are popular Document originated DBMS systems.



Graph-Based

- A graph type database stores entities as well the relations amongst those entities. The entity is stored as a node with the relationship as edges. An edge gives a relationship between nodes. Every node and edge has a unique identifier.
- Compared to a relational database where tables are loosely connected, a Graph database is a multi-relational in nature.
 Traversing relationship is fast as they are already captured into the DB, and there is no need to calculate them.







Introduction to MongoDB,

- MongoDB is a document-oriented NoSQL database used for high volume data storage. Instead of using tables and rows as in the traditional relational databases, MongoDB makes use of collections and documents.
- Documents consist of key-value pairs which are the basic unit of data in MongoDB.
- Collections contain sets of documents and function which is the equivalent of relational database tables.
- MongoDB is a database which came into light around the mid-2000s



MongoDB Features

 Each database contains collections which in turn contains documents. Each document can be different with a varying number of fields. The size and content of each document can be different from each other.



MongoDB Features

- The rows (or documents as called in MongoDB) doesn't need to have a schema defined beforehand. Instead, the fields can be created as needed
- The data model available within MongoDB allows you to represent hierarchical relationships, to store arrays, and other more complex structures more easily.
- Scalability The MongoDB environments are very scalable. Companies across the world have defined clusters with some of them running 100+ nodes with around millions of documents within the database



MongoDB Example

- The below example shows how a document can be modeled in MongoDB.
 - The _id field is added by MongoDB to uniquely identify the document in the collection.
- the Order Data (OrderID, Product, and Quantity) which in RDBMS will normally be stored in a separate table, while in MongoDB it is actually stored as an embedded document in the collection itself. This is one of the key differences in how data is modeled in MongoDB.



document customer

```
_id:1,
CustomerName:indu
Order:
                                             Embedded document
          OrderID:111
          Product:book
          Quantity:5
```



student

```
Name:rahul
Age:23
Rollno:12
Marks:
   { eng:34
     hindi:45
     math:50
```



Key Components of MongoDB Architecture

- **Collection** This is a grouping of MongoDB documents. A collection is the equivalent of a table which is created in any other RDBMS such as Oracle or MS SQL. A collection exists within a single database. As seen from the introduction collections don't enforce any sort of structure.
- Cursor This is a pointer to the result set of a query. Clients can iterate through a cursor to retrieve results.
- Database This is a container for collections like in RDBMS wherein it is a container for tables. Each database gets its own set of files on the file system. A MongoDB server can store multiple databases.



Key Components of MongoDB Architecture

- Document A record in a MongoDB collection is basically called a document. The document, in turn, will consist of field name and values.
- **Field** A name-value pair in a document. A document has zero or more fields. Fields are similar to columns in relational databases.
- The diagram shows an example of Fields with Key value pairs. So in the example below CustomerID and 11 is one of the key value pair's defined in the document.
- Data stored as key-value pair



Key Components of MongoDB Architecture

```
{
        CustomerID:111,
        CustomerName: "indu"
        OrderID: 111
}
```

- **JSON** This is known as JavaScript Object Notation. This is a human-readable, plain text format for expressing structured data. JSON is currently supported in many programming languages.
- The _id field is used to uniquely identify the documents in a collection and is automatically added by MongoDB when the collection is created.



Contents to be covered

- Data modelling in MongoDB,
- Advantages of MongoDB over RDBMS,
- Mongo Shell,
- Configuration file in MongoDB,



Data Modelling in MongoDB

- As we have seen from the Introduction section, the data in MongoDB has a flexible schema.
- in SQL databases, where you must have a table's schema declared before inserting data, MongoDB's collections do not enforce document structure.
- This sort of flexibility is what makes MongoDB so powerful.



Data Modelling in MongoDB

- When modeling data in Mongo, keep the following things in mind
 - What are the needs of the application Look at the business needs of the application and see what data and the type of data needed for the application. Based on this, ensure that the structure of the document is decided accordingly.
 - What are data retrieval patterns If you foresee a heavy query usage then consider the use of indexes in your data model to improve the efficiency of queries.
 - Are frequent inserts, updates and removals happening in the database?
 Reconsider the use of indexes or incorporate sharding if required in your data modeling design to improve the efficiency of your overall MongoDB environment.



Difference between MongoDB & RDBMS

Below are some of the key term differences between MongoDB and RDBMS

RDBMS	MongoDB	Difference
Table	Collection	In RDBMS, the table contains the columns and rows which are used to store the data whereas, in MongoDB, this same structure is known as a collection. The collection contains documents which in turn contains Fields, which in turn are key-value pairs.
Row	Document	In RDBMS, the row represents a single, implicitly structured data item in a table. In MongoDB, the data is stored in documents.
Column	Field	In RDBMS, the column denotes a set of data values. These in MongoDB are known as Fields.
Joins		In RDBMS, data is sometimes spread across various tables and in order to show a complete view of all data, a join is sometimes formed across tables to get the data. In MongoDB, the data is normally stored in a single collection, but separated by using Embedded documents. So there is no concept of joins in MongoDB.

National Institute of Electronics & Information technology

MMMUT campus Deoria Bypass Road Goralhpur-273010

- Apart from the terms differences, a few other differences are shown below
 - Relational databases are known for enforcing data integrity. This is not an explicit requirement in MongoDB.
 - RDBMS requires that data be normalized first so that it can prevent orphan records and duplicates .Normalizing data then has the requirement of more tables, which will then result in more table joins, thus requiring more keys and indexes.
 - As databases start to grow, performance can start becoming an issue.
 Again this is not an explicit requirement in MongoDB. MongoDB is flexible and does not need the data to be normalized first.



Name	Age	Address	Gender	Mobile
Aaa	45	Hghjghj	F	985632147
Aaa	45	Hghjghj	F	845759635



Advantages of MongoDB over RDBMS

- Ease of scale-out MongoDB is easy to scale.
- Uses internal memory for storing the (windowed) working set, enabling faster access of data.



MongoDB

 MongoDB is a cross-platform, document oriented database that provides, high performance, high availability, and easy scalability. MongoDB works on concept of collection and document.

Database

 Database is a physical container for collections. Each database gets its own set of files on the file system. A single MongoDB server typically has multiple databases.



MongoDB

Collection

 Collection is a group of MongoDB documents. It is the equivalent of an RDBMS table. A collection exists within a single database.
 Collections do not enforce a schema. Documents within a collection can have different fields. Typically, all documents in a collection are of similar or related purpose.



MongoDB

Document

– A document is a set of key-value pairs. Documents have dynamic schema. Dynamic schema means that documents in the same collection do not need to have the same set of fields or structure, and common fields in a collection's documents may hold different types of data.



- Student collections
- {"name":"aaa","rollno":1,"marks":45}
- {"name":"bbb","rollno":2,"hobbies":"playing cricket"}



Following table shows the relationship of RDBMS terminology

with MongoDB.

RDBMS	MongoDB			
Database	Database			
Table	Collection			
Tuple/Row	Document			
column	Field			
Table Join	Embedded Documents			
Primary Key	Primary Key (Default key _id			
	provided by			
	mongodb itself)			
Database Server and Client				
Mysqld/Oracle	mongod			
mysql/sqlplus ctronics & Information	mongo technology			

National Institute of Electronics & Information technology



Sample Document

Following example shows the document structure, which is simply a comma separated key value pair. " id":"34893483478965767ashdhsda", "name":"abc", "rollno":1, "phone number":6535675324, "Marks": { "eng":45,"maths":67,"science":67}

ID field

- **_id** is a 12 bytes hexadecimal number which assures the uniqueness of every document. You can provide _id while inserting the document. If you don't provide then MongoDB provides a unique id for every document.
- These 12 bytes first 4 bytes for the current timestamp, next 3 bytes for machine id, next 2 bytes for process id of MongoDB server and remaining 3 bytes are simple incremental VALUE.
- 12 bytes=8*12 bits=96 bits / 4=24 digit hexadecimal
- When you convert 12 bytes to hexadecimal it becomes 24 digit number
- 48/4=12



Install MongoDB Community Edition

Procedure

 Follow these steps to install MongoDB Community Edition using the MongoDB Installer wizard. The installation process installs both the MongoDB binaries as well as the default configuration file <install directory>\bin\mongod.cfg.

Download the installer.

Download the MongoDB Community .msi installer from the following link:



- > go to website mongodb.com
- In the **Version** dropdown, select the version of MongoDB to download.
- In the Platform dropdown, select Windows.
- In the **Package** dropdown, select **msi**.
- Click Download.



Run the MongoDB installer.

For example, from the Windows Explorer/File Explorer:

Go to the directory where you downloaded the MongoDB installer (.msi file). By default, this is your Downloads directory.

Double-click the .msi file.

Follow the MongoDB Community Edition installation wizard.



 The wizard steps you through the installation of MongoDB and MongoDB Compass.

Choose Setup Type

You can choose either the Complete (recommended for most users)
 or Custom setup type. The Complete setup option installs MongoDB and
 the MongoDB tools to the default location. The Custom setup option
 allows you to specify which executables are installed and where.

Service Configuration

 Starting in MongoDB 4.0, you can set up MongoDB as a Windows service during the install or just install the binaries.



MongoDB Service

 Starting in MongoDB 4.0, you can configure and start MongoDB as a Windows service during the install, and the MongoDB service is started upon successful installation.



- Select Install MongoD as a Service
 MongoDB as a service.
 - Select either:
 - Run the service as Network Service user (Default)
 - This is a Windows user account that is built-in to Windows

or

- Run the service as a local or domain user
 - For an existing local user account, specify a period (i.e. .) for the Account
 Domain and specify the Account Name and the Account Password for the user.



- For an existing domain user, specify the Account Domain, the Account Name and the Account Password for that user.
- **Service Name**. Specify the service name. Default name is MongoDB. If you already have a service with the specified name, you must choose another name.
- **Data Directory**. Specify the data directory, which corresponds to the --dbpath. If the directory does not exist, the installer will create the directory and sets the directory access to the service user.
- **Log Directory**. Specify the Log directory, which corresponds to the --logpath. If the directory does not exist, the installer will create the directory and sets the directory access to the service user.



Install MongoDB Compass

- Optional. To have the wizard install MongoDB Compass, select Install MongoDB Compass (Default).
- When ready, click Install.



Mongo Shell

- Mongo shell or MongoDB Shell is an interactive JavaScript interface to MongoDB. It is used to perform administrative operations and also query and update data.
- The mongo shell is as part of the MongoDB Server installation.
- Mongo Shell is also provided as a standalone package by MongoDB.
- Prior to connecting to MongoDB one has to ensure that MongoDB is running.



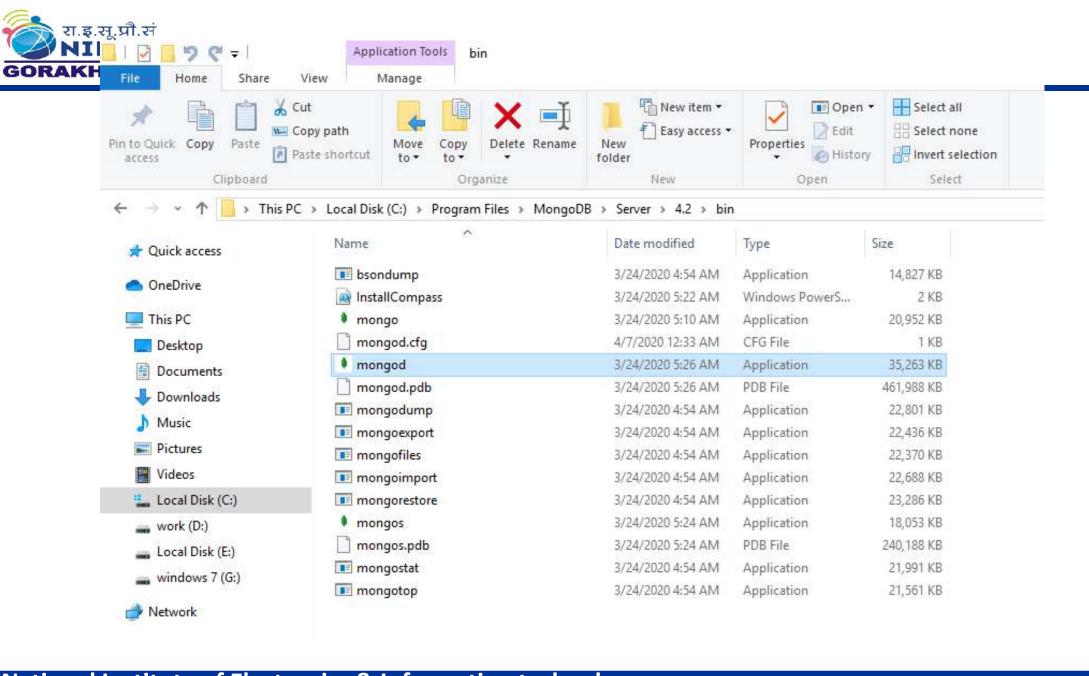
To start MongoDB, Run the following command

```
Command Prompt
 :\Program Files\MongoDB\Server\4.2\bin>mongod.exe
2020-04-28T20:12:28.755+0530 I CONTROL [main] Automatically disabling TLS 1.0, to force-enable TLS 1.0 specify --sslDi
sabledProtocols 'none'
                                         [main] No TransportLayer configured during NetworkInterface startup
2020-04-28T20:12:28.760+0530 W ASIO
                                         [initandlisten] MongoDB starting : pid=9592 port=27017 dbpath=C:\data\db\ 64-bi
2020-04-28T20:12:28.761+0530 I CONTROL
 host=DESKTOP-LK8CJ4M
                                         [initandlisten] targetMinOS: Windows 7/Windows Server 2008 R2
2020-04-28T20:12:28.761+0530 I CONTROL
                                         [initandlisten] db version v4.2.5
2020-04-28T20:12:28.761+0530 I CONTROL
                                         [initandlisten] git version: 2261279b51ea13df08ae708ff278f0679c59dc32
2020-04-28T20:12:28.761+0530 I CONTROL
                                         [initandlisten] allocator: tcmalloc
2020-04-28T20:12:28.762+0530 I CONTROL
                                         [initandlisten] modules: none
2020-04-28T20:12:28.762+0530 I CONTROL
                                         [initandlisten] build environment:
2020-04-28T20:12:28.762+0530 I CONTROL
                                         [initandlisten]
                                                             distmod: 2012plus
2020-04-28T20:12:28.762+0530 I CONTROL
2020-04-28T20:12:28.762+0530 I CONTROL
                                         [initandlisten]
                                                             distarch: x86 64
                                         [initandlisten]
2020-04-28T20:12:28.763+0530 I CONTROL
                                                             target arch: x86 64
                                         [initandlisten] options: {}
2020-04-28T20:12:28.763+0530 I CONTROL
                                         [initandlisten] exception in initAndListen: NonExistentPath: Data directory C:\
2020-04-28T20:12:28.764+0530 I STORAGE
data\db\ not found., terminating
                                         [initandlisten] shutdown: going to close listening sockets...
2020-04-28T20:12:28.765+0530 I NETWORK
                                         [initandlisten] Stopping further Flow Control ticket acquisitions.
2020-04-28T20:12:28.765+0530 I -
                                         [initandlisten] now exiting
2020-04-28T20:12:28.765+0530 I CONTROL
                                         [initandlisten] shutting down with code:100
2020-04-28T20:12:28.765+0530 I CONTROL
C:\Program Files\MongoDB\Server\4.2\bin>
```



OR

• Run the **mongod.exe** file on command prompt (path of the exe file is C:\Program Files\MongoDB\Server\4.2\bin)



National Institute of Electronics & Information technology MMMUT campus Deoria Bypass Road Goralhpur-273010



Running Mongo Shell

- Once you are sure that MongoDB is running,
- Run the **mongo.exe** file from the same path (C:\Program Files\MongoDB\Server\4.2\bin) from command prompt or using Run or using Windows explorer.
- Once executed, the Mongo Shell will open. Now you can execute desired commands.



Example: Running Help Command on the Mongo Shell

Various help options will be displayed as shown in the figure below:

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                 mprovements and to suggest MongoDB products and deployment options to you.
To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
 help
       db.help()
                                    help on db methods
       db.mycoll.help()
                                    help on collection methods
                                    sharding helpers
       sh.help()
       rs.help()
                                    replica set helpers
       help admin
                                    administrative help
       help connect
                                    connecting to a db help
       help keys
                                    key shortcuts
       help misc
                                    misc things to know
       help mr
                                    mapreduce
       show dbs
                                    show database names
       show collections
                                    show collections in current database
       show users
                                    show users in current database
                                    show most recent system.profile entries with time >= 1ms
       show profile
       show logs
                                    show the accessible logger names
                                    prints out the last segment of log in memory, 'global' is default
       show log [name]
                                    set current database
       use <db name>
       db.foo.find()
                                    list objects in collection foo
      db.foo.find( { a : 1 } )
                                    list objects in foo where a == 1
                                    result of the last line evaluated; use to further iterate
       DBQuery.shellBatchSize = x
                                    set default number of items to display on shell
                                    quit the mongo shell
       exit
```



- When we install the full version of MongoDB, configuration file mongod.cfg is automatically created with default setting and is stored in bin folder of the MongoDB server.
- We have installed MongoDb on C:\ drive at the default path and the location of the configuration file is C:\Program Files\MongoDB\Server\4.2\bin.



• If we open this file, the content of the file is as under:

```
# mongod.conf
# for documentation of all options, see:
#http://docs.mongodb.org/manual/reference/configuration-options/
# Where and how to store data.
storage:
dbPath: C:\Program Files\MongoDB\Server\4.2\data
journal:
enabled: true
# engine:
# mmapv1:
# wiredTiger:
```



```
# where to write logging data.
systemLog:
destination: file
logAppend: true
path: C:\Program Files\MongoDB\Server\4.2\log\mongod.log
#processManagement:
#security:
#operationProfiling:
#replication:
#sharding:
```



```
# network interfaces
```

net:

port: 27017

bindlp: 127.0.0.1

Enterprise-Only Options:

#auditLog:

#snmp:

We may edit the path and settings or create our own configuration file with these settings. Storage path and system log path is mandatory.



Contents to be covered

- JSON File format for storing documents
- Introduction to Documents, Collections,
- Database Commands in Mongodb,
- Update and save method
- Quering documnets



JSON File format for storing documents

- JSON stands for JavaScript Object Notation. Json is data-interchange format. It is easy for humans to read and write. It is easy for machines to parse and generate. We call JSON as syntax for storing and exchanging data.
- JSON is a text format which is completely language independent and written with JavaScript object notation. JSON uses conventions that are familiar to programmers of the C- family of languages, including C, C++, C#, Java, JavaScript, Perl, Python, and many others.
- These properties make JSON an ideal data-interchange language.



JSON is built on two structures:

- As a collection of name/value or Attribute/Value pairs. This
 may be an object, record, struct, dictionary, hash table, keyed
 list, or associative array.
- An ordered list of values. This may be an array, vector, list, or sequence.
- JSON is universal data structures which are virtually supported by all modern programming in one form or another. These JSON structure supports the data format which is interchangeable with programming languages.



JSON

- Example 1: An object is an unordered set of name/value pairs. In JSON, an object starts with { left brace and ends with } right brace. Each name is followed by :colon and the name/value pairs are separated by ,comma.
- {"name": "Ajay", "age": 21, "city": "Gorakhpur" }
- { "name": "Mohit", "age": 22, "city": "Deoria" }
- { "name": "Jeetut", "age": 20, "city": "Sonauli" }



JSON

• <u>Example 2</u>: An array is an ordered collection of values. An array begins with *[left bracket* and ends with *] right bracket*. Here, values are separated by *,comma*. Here comes some arrays:

```
["Graduation", "Post Graduation", "NIELIT Course"]
["BSc", "BCom", "BA", "BBA", "BCA"]
["MSc", "MCom", "MA", "MBA", "MCA"]
["CCC", "O Level IT", "A Level IT", "CHM O Level", "MAT O Level"]
```



JSON File format for storing documents

Example 3:

Array of Objects

```
[ {"name": "Ajay", "age": 21, "city": "Gorakhpur" }, 
{ "name": "Mohit", "age": 22, "city": "Deoria" }, 
{ "name": "Jeetut", "age": 20, "city": "Sonauli" } ]
```



Exchanging Data with JSON File

- It is worth to mention here that exchange of data between a browser and a server can only be text.
- JSON is text format and hence we can convert any JavaScript object into JSON, and send JSON to the server.
- Similarly, we may also convert any JSON received from the server into JavaScript objects.
- With this technique, we may work with the data with JSON file as JavaScript objects, with no complicated parsing and translations.



Introduction to Documents, Collections,

MongoDB stores data records as documents (specifically BSON documents) which are gathered together in collections.
 A database stores one or more collections of documents.

Databases

In MongoDB, databases hold one or more collections of documents.

Collections

 MongoDB stores documents in collections. Collections are analogous to tables in relational databases.



Commands in MongoDB

- All the commands will be written in MongoDB shell.
 Alternatively the operations may be executed in Compass GUI.
 We are using commands in shell.
- OPEN mongo shell by using the shell application file mongo.exe from at C:\Program Files\MongoDB\Server\4.2\bin (or the directory you have specified)
- Mongo screen will open



```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                MongoDB shell version v4.2.6
connecting to: mongodb://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("c31fa2e5-6e02-4bbf-9d9b-05d38ab1e81c") }
NongoDB server version: 4.2.6
Server has startup warnings:
2020-04-28T09:41:18.954+0530 I CONTROL [initandlisten]
                                        [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-04-28T09:41:18.954+0530 I CONTROL
2020-04-28T09:41:18.954+0530 I CONTROL [initandlisten] **
                                                                    Read and write access to data and configuration is
unrestricted.
2020-04-28T09:41:18.964+0530 I CONTROL [initandlisten]
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).
The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.
To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
Music
```



Create database

- All commands can be written and executed on this MongoDB Shell prompt.
 - use DATABASE_NAME command is used to create database.
 - This will open/ or switch the specified database if exists otherwise will create a database with the given name and switch to that database.
- Syntax

use DATABASE_NAME

Example:

Use abc



Create database

```
To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()

"> use abc
switched to db abc

"> =

MongoDB-Indexing

>db
model
```



Show database

- Show database (show dbs) command
 - To display the list of databases, show command is used,

Syntax

Show dbs

- This will list all the non-empty databases.
- To display database, you need to insert at least one document into it.
- Just created database (abc) is not present in list shown below as it is not having any document in it.



Show database



The dropDatabase() Method

- MongoDB db.dropDatabase() command is used to drop a existing database.
- Syntax
- Basic syntax of dropDatabase() command is as follows db.dropDatabase()
- This will delete the selected database. If you have not selected any database, then it will delete default 'test' database.



createCollection() method is used to create collection in the selected database.

Syntax

db.createCollection(name, options)

Example

To run this method / command database must be in use. We are using our newly created **abc** database.

```
db.createCollection("college")
this will result status as
{ "ok": 1 }
```



Parameter	Туре	Description
Name	String	Name parameter specifies the name of the collection to be created. It is of String type and should be specified in quotes.
Options (it is optional parameter)	Document	Options parameter is used to specify options about memory size and indexing. This parameter is optional.



```
> db.createCollection(college)
2020-05-04T10:42:41.353+0530 E QUERY [js] uncaught exception: ReferenceError: college is not defined:
@(shell):1:1
> db.createCollection("college")
{ "ok" : 1 }
> Activate \
```



- **Note:** In the above screen, 1st we write collection name without quotes **db.createcollection(college)** which is wrong and in that case some exceptions has been raised.
- In MongoDB, collection creation is not necessary before inserting document. MongoDB creates collection automatically, when you insert some document.



The drop() Method

- MongoDB's db.collection.drop() is used to drop a collection from the database.
- Syntax
- Basic syntax of drop() command is as follows –
- db.COLLECTION_NAME.drop()



Show collections command

- Show collections command
 - Show collections is used to display / list the created collection in the selected database.
- Syntax

Show collections

 If we run this command, it will display our recently created collection "college".



Show collections command

```
> show collections college
```

are -

- **String** This is the most commonly used datatype to store the data. String in MongoDB must be UTF-8 valid.
- Integer This type is used to store a numerical value. Integer can be 32 bit or 64 bit depending upon your server.
- Boolean This type is used to store a boolean (true/ false) value.
- Double This type is used to store floating point values.

are -

- Min/ Max keys This type is used to compare a value against the lowest and highest BSON elements.
- Arrays This type is used to store arrays or list or multiple values into one key.
- Object This datatype is used for embedded documents.



Inserting and Saving Documents

• Insert() command

 Insert command is used to insert document in form of name/value combination to the specified collection. If collection is not already created, it will create the specified collection.

Syntax

db.COLLECTION_NAME.insert (document)



Inserting and Saving Documents

where

- COLLECTION_NAME is the name of the collection in which document has to be inserted.
- document is document(key/value combination) or array of documents to insert into the collection.



Inserting and Saving Documents

_id Field

- We may give an _id filed while inserting the document and this shall be unique within the collection to avoid duplicate key error.
- If no _id field is specified, then MongoDB will add the _id field and assign a unique ObjectId to the document



Example-1:

db.college.insert({"name" : "RCC"})

This will show the following result after successful execution of the command:

WriteResult({ "nInserted" : 1 })



```
> db.college.insert({name:"RCC"})
WriteResult({ "nInserted" : 1 })
> show collections
abc
college
```



Example-2:

db.college.insert({course:"A level",duration:"2 years", fees:"12000"})

This will show the following result after successful execution of

the command:

WriteResult({ "nInserted" : 1 })



Example-2:

```
> db.college.insert({course:"A level",duration:"2 years",fees:"12000"})
WriteResult({ "nInserted" : 1 })
```



Example-3:

• Lets insert a document into a collection which has not been created. For example let collection name as **school**.

```
db.school.insert({"class": "10th", "section": "A section", "student
count": "39"})
```

 Once this command is executed, a new collection named as school will be automatically created.



Example-3:



The insertOne() method

 If you need to insert only one document into a collection you can use this method.

- Syntax
- The basic syntax of insert() command is as follows –

>db.COLLECTION_NAME.insertOne(document)



Batch insert or Inserting Multiple documents into a collection

To Insert multiple documents or Batch insert into a collection insertMany() method is available to do this task.

Syntax:

```
db.collection.insertMany (
  [ <document 1> , <document 2>, ... ]
)
```

Where

Each of these document should be separated by a , comma



Batch insert or Inserting Multiple documents into a collection

- [<document 1>, <document 2>, ...] is an array of the various s documents which needs to be inserted into the collection.
- On Successful execution, it returns a
 - A boolean acknowledged as true if the operation is successful or false if uncessful.
 - An array of _id for each successfully inserted documents



The insertMany() method

- You can insert multiple documents using the insertMany() method. To this method you need to pass an array of documents.
- Example
- Following example inserts three different documents into the empDetails collection using the insertMany() method.



Read Operations

- Read operations retrieve documents from a collection; i.e. query a collection for documents. MongoDB provides the following methods to read documents from a collection:
- db.collection.find()



Find() command

 Find () command is used to see the records/documents in the specified collection.

Syntax

db.COLLECTION_NAME.find()

Example:

db.college.find()



Find example

Lets insert a record with specified _id
 db.school.insert({"_id: "01", "class": "10th", "section": "B section", "student count": "27"})



find example:

 Now a record with _id as 01 (specified in insert command) has been inserted in the collection. Find command used here has been described below.



The pretty() Method

- To display the results in a formatted way, you can use pretty()
 method.
- Syntax
- >db.COLLECTION_NAME.find().pretty()



The findOne() method

• Apart from the find() method, there is findOne() method, that returns only one document.

- Syntax
- >db.COLLECTIONNAME.findOne()



Update and save

- MongoDB's update() and save() methods are used to update document into a collection. The update() method updates the values in the existing document while the save() method replaces the existing document with the document passed in save() method.
- MongoDB Update() Method
- The update() method updates the values in the existing document.
- Syntax
- The basic syntax of update() method is as follows –
- >db.COLLECTION_NAME.update(SELECTION_CRITERIA, UPDATED_DATA)



Save() command

• Save() command may also be used to insert the document. It has the same syntax as of insert() command.

Syntax

db.college.save(document)

- If _id is not specified in the document then save() method will work same as insert() method.
- If _id feild is specified then it will replace whole data of document having the _id as specified in save() method if matched else Insert the document into the collection.



Example-1

Lets save a record in the school collection with _id as 02, and then display the
documents in the collection.

```
db.school.save({"_id":"02","class":"10th","section":"C Section","student count":"37"})
```

this will display

```
WriteResult({ "nMatched" : 0, "nUpserted" : 1, "nModified" : 0, "_id" : "02" })
```

It will display details like "nMatched": 0, i.e. no record is matched with the specified _id, so a new record is "nUpserted": 1 with , "_id": "02".



Example-2

Lets try to save a record in the school collection with _id as 01, and then display the documents in the collection.

```
db.school.save({"_id":"01","class":"10th","section":"D Section","student count":"29"})
```

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified"
: 1 })
```

It will display details like "nMatched": 1 which matched with the specified _id, and the same is modified "nModified": 1.



Quering documents

- Find is used to search for documents
- It can be based on some criteria
- syntax
 db.collection.find(criteria)
 Eg
 db.student.find({name:"tina"})



Select All Documents in a Collection

- To select all documents in the collection, pass an empty document as the query filter parameter to the find method.
 The query filter parameter determines the select criteria:
- db.employee.find({})



Specify Equality Condition

 To specify equality conditions, use <field>:<value> expressions in the query filter document:

db.employee.find({ designation: "manager" })



comparison query operators for find

Name	Description
\$eq	Matches values that are equal to a specified value.
\$gt	Matches values that are greater than a specified value.
\$gte	Matches values that are greater than or equal to a specified value.
\$1t	Matches values that are less than a specified value.
\$1te	Matches values that are less than or equal to a specified value.



Specifies equality condition.

 The \$eq operator matches documents where the value of a field equals the specified value.

```
{ <field>: { $eq: <value> } }
```

- To find all the doucments where marks is equal to 50
- db.student.find({ marks: { \$eq: 50 } })

\$gt

- \$gt selects those documents where the value of the field is greater than (i.e. >) the specified value.
- syntax:

```
{field: {$gt: value} }
```

- Eg:
- db.student.find({ marks: { \$gt: 20 } })



\$gte

• \$gte selects the documents where the value of the field is greater than or equal to (i.e. >=) a specified value (e.g. value.)

• Syntax:

{field: {\$gte: value} }

db.student.find({ marks: { \$gte: 20 } })



\$in

The \$in operator selects the documents where the value of a field equals any value in the specified array. To specify an \$in expression, use the following prototype
{ field: { \$in: [<value1>, <value2>, ... <valueN>] } }

db.student.find({ marks: {\$in: [50, 75]}})

\$It

- Syntax: {field: {\$lt: value} }
- \$It selects the documents where the value of the field is less than (i.e. <) the specified value.
- db.student.find({ marks: { \$lt: 20 } })

\$Ite

 \$Ite selects the documents where the value of the field is less than or equal to (i.e. <=) the specified value

• Syntax:

```
{ field: { $lte: value} }
```

db.student.find({ marks: { \$lte: 20 } })

\$ne

• \$ne selects the documents where the value of the field is not equal to the specified value. This includes documents that do not contain the field.

• Syntax:

{field: {\$ne: value} }

db.student.find({ marks: { \$ne: 20 } })



logical operators with find

Logical

Name	Description
\$and	Joins query clauses with a logical AND returns all documents that match the conditions of both clauses.
\$or	Joins query clauses with a logical OR returns all documents that match the conditions of either clause.



logical operators

• Syntax:

```
{ $and: [ { <expression1> }, { <expression2> } , ... , { <expression
N> } ] }
```

• \$and performs a logical AND operation on an array of *one or more* expressions (e.g. <expression1>, <expression2>, etc.) and selects the documents that satisfy *all* the expressions in the array



db.student.find({ \$and: [{ marks: { \$ne: 50 } }, { rollno: { \$lt:4 } }] })



logical operators

- The \$or operator performs a logical OR operation on an array of *two or more* <expressions> and selects the documents that satisfy *at least* one of the <expressions>.
- The \$or has the following syntax:

```
{ $or: [ { <expression1> }, { <expression2> }, ... , { <expressionN>
} ] }
```



The Limit() Method

To limit the records in MongoDB, you need to use limit() method.
 The method accepts one number type argument, which is the number of documents that you want to be displayed.

- Syntax
- The basic syntax of limit() method is as follows –

>db.COLLECTION_NAME.find().limit(NUMBER)



MongoDB Skip() Method

Apart from limit() method, there is one more method skip()
which also accepts number type argument and is used to skip
the number of documents.

- Syntax
- The basic syntax of skip() method is as follows –

>db.COLLECTION_NAME.find().skip(NUMBER)



The sort() Method

To sort documents in MongoDB, you need to use sort() method.
The method accepts a document containing a list of fields along
with their sorting order. To specify sorting order 1 and -1 are used.
1 is used for ascending order while -1 is used for descending order.

- Syntax
- The basic syntax of sort() method is as follows –
- >db.COLLECTION_NAME.find().sort({KEY:1})



Contents to be covered

- projection
- Updating Documents
- Operator and Modifiers
- Removing Documents
- Document Replacement



projection

 By default, queries in MongoDB return all fields in matching documents. To limit the amount of data that MongoDB sends to applications, you can include a projection document to specify or restrict field



projection

• If you do not specify a projection document, the db.collection.find() method returns all fields in the matching documents.

Syntax:

db.collectionname.find({criteria},{field:1})

- db.marks.find({ rollno:11 }, { name: 1,marks: 1 })



Updating Documents

- update() Method: Updating the values in document
 - The update() method is used to update the value in the existing document of specified collection.



Syntax

db.COLLECTION_NAME.update(SELECTION_CRITERIA, UPDATED_DATA)

where

SELECTION_CRITERIA -- The selection criteria for the updating the document

UPDATED_DATA -- The modifications to apply to the document.



Operator and Modifiers

Update Operators

Update methods supports various field and array operators. It has the following syntax:

```
Syntax
```

```
{
    <operator1>: { <field1>: <value1>, ... },
    <operator2>: { <field2>: <value2>, ... },
    ...
}
```



Operator and Modifiers

Where operator can be field or array operator.



Field operators

Name	Description
\$set	Sets the value of a field in a document.
\$inc	Increments the value of the specified field by the specified amount.
\$min	if the given value is less than the existing field value, then only updates the field
\$max	if the given value is greater than the existing field value, then only updates the field
\$mul	Used to multiply the value of the specified field by the given amount.
\$rename	Renames the specified field.



\$set

It may be simply given with \$set: parameter.

Example 1

db.school.update({"class":"10th"}, {\$set:{class:"11th"}})

after this command, record with class 10th will be updated to class 11th



```
> db.school.find()
 " id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "10th", "section" : "A section", "student count" : "40"
\{ "id": ObjectId("5eb24d89db4e42dfda383020"), "class": "12th", "section": "C section", "student count": "30"
 "id": ObjectId("5eb24d8fdb4e42dfda383021"), "class": "12th", "section": "B section", "student count": "35"
" id" : ObjectId("5eb24ef5db4e42dfda383022"), "class" : "9th", "section" : "A section", "student count" : "40"
\{ " id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "9th", "section" : "B section", "student count" : "39" \}
> db.school.update({"class":"10th"}, {$set:{class:"11th"}})
WriteResult({    "nMatched" : 1,    "nUpserted" : 0,    "nModified" : 1    })
> db.school.find()
{ " id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40"
{ " id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "12th", "section" : "C section", "student count" : "30"
 "id": ObjectId("5eb24d8fdb4e42dfda383021"), "class": "12th", "section": "B section", "student count": "35"
" id" : ObjectId("5eb24ef5db4e42dfda383022"), "class" : "9th", "section" : "A section", "student count" : "40"
\{ " id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "9th", "section" : "B section", "student count" : "39" \}
```



Example 2:

In case of multiple matching criteria, only 1st document will be updated

db.school.update({"class":"12th"}, {\$set:{class:"9th"}})

after this command, 1st record with class 12th will be updated to class 9th



```
db.school.find()
{ "_id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "12th", "section" : "C section", "student count" : "30" }
{ "_id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24def5db4e42dfda383022"), "class" : "10th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "9th", "section" : "B section", "student count" : "39" }
> db.school.update({"class":"12th"}, {$set:{class:"9th"}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.school.find()
{ "_id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "9th", "section" : "C section", "student count" : "30" }
{ "_id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24dfda383022"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383022"), "class" : "10th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "10th", "section" : "B section", "student count" : "39" }
```

WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1
 }) means 1 record matched and updated one record.



updateMany() Method:

The updateMany() method is used to update all the matching records of collection. All other parameter and system is similar to update() method.

Syntax

db.COLLECTION_NAME.updateMany(SELECTION_CRITERIA, UPDATED_DATA)

Example:

Lets update all the records of the matching criteria, db.school.updateMany({"class":"9th"}, {\$set:{class:"12th"}})



updateMany() Method:

```
db.school.find()
{ "_id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "9th", "section" : "C section", "student count" : "30" }
{ "_id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383022"), "class" : "10th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "9th", "section" : "B section", "student count" : "39" }
> db.school.updateMany({"class":"9th"}, {$set:{class:"12th"}})
{ "acknowledged" : true, "matchedCount" : 2, "modifiedCount" : 2 }
> db.school.find()
{ "_id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40" }
{ "_id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "12th", "section" : "C section", "student count" : "30" }
{ "_id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24d8fdb4e42dfda383022"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "12th", "section" : "B section", "student count" : "39" }
}
```



updateMany() Method:

It results as:

```
{ "acknowledged" : true, "matchedCount" : 2, "modifiedCount" : 2 }
```

Its shows that two records are matched and 2 of them has been modified



\$inc

- Increments the value of the specified field by the specified amount.
- syntax

db.collectionname.update({criteria},{\$inc:{field:value}})

Eg:

db.student.update({rollno:4},{\$inc:{marks:10}})



\$min

- if the given value is less than the existing field value, then only updates the field
- Syntax:

db.collectionname.update({criteria},{\$min:{field:value}})

- Eg
- db.student.update({rollno:4},{\$min:{marks:10}})



\$max

 if the given value is greater than the existing field value, then only updates the field

- db.collectionname.update({criteria},{\$max:{field:value}})
- Eg
- db.student.update({rollno:4},{\$max:{marks:70}})



\$mul

 Used to multiply the value of the specified field by the given amount.

- db.collectionname.update({criteria},{\$mul:{field:value}})
- Eg
- db.student.update({rollno:4},{\$mul:{marks:10}})



\$rename

Renames the specified field.

- db.collectionname.update({criteria},{\$rename:{field:value}})
- Eg
- db.student.update({rollno:4},{\$rename:{"marks":"score"}})



Array Operators

Name	Description
\$addToSet	Adds given elements to the array, but the existing elements will not be added.
\$push	Adds an item to the array.



Modifiers with array operator

Update methods supports following 4 modifiers to update documents based on the criteria.

Modifier name	Description
\$each	If we wish to add multiple items to array in update, \$each is used with \$push and \$addToSet operators
\$position	Used to add element at a particular position in array. It is used with \$push operator. Without \$position, element is added at the end of the array.
\$slice	Used to limit the size of updated arrays when used with \$push, i.e. limit the number of elements.
\$sort	Modifies the order of the elements of an array when used with \$push operation.



Insert values into Array

- Lets insert the multiple values into an array. The arrays are represented in []. For example and array of marks can be [45, 44, 43].
- The values in array can be inserted in [] using simple db.collectionname.insert() method.



Example:

C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
> db.result.find()
>
> db.result.insert({"name": "amit", "marks":[45, 48, 41]})
WriteResult({ "nInserted" : 1 })
> 
> db.result.insert({"name": "indu", "marks":[42, 44]})
WriteResult({ "nInserted" : 1 })
> 
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44 ] }
>
```



operators used with array

- { \$addToSet: { <field>: [<value1>, <value2> ...] } } }
 - \$addToSet array operator to add multiple values to an array <field>. If the values do not exist in the <field> of array it adds to it, if exist do nothing.
- { \$push: { <field>: [<value1>, <value2> ...] } } } \$push operator to append multiple values to an array <field>.

If you use these operator without each then they will add the values as sub array



Syntax of \$each modifier

- Each modifier adds the value in the array
- { \$addToSet: { <field>: { \$each: [<value1>, <value2> ...] } } } \$each modifier can be Used alongwith the \$addToSet array operator to add multiple values to an array <field>. If the values do not exist in the <field> of array it adds to it, if exist do nothing.
- { \$push: { <field>: { \$each: [<value1>, <value2> ...] } } } \$each modifier can be Used with the \$push operator to append multiple values to an array <field>.



\$addToSet operator with \$each modifier:

- Let update the Marks array and add elements to it. This is being done using \$each modifier and \$addToSet operator.
- With this only elements specified in the update, will be appended to the array, if they do not exist in the array, if for existing elements, do nothing.
- In the below example, we are adding 3 elements to the array marks only if the specified elements are not existing in the array for amit.

GORAKHPUS add To Set operator with Seach modifier:

C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
db.result.find()
 "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41 ] }
 "id": ObjectId("5eb8e6d0c92747cab059837a"), "name": "indu", "marks": [ 42, 44, 43, 47, 51 ] }
 db.result.update({"name":"amit"}, {$addToSet: {"marks":{$each: [42,41,48]}}})
VriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
 db.result.find()
 "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41, 42 ] }
 " id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 43, 47, 51 ] }
```



Examples of Update Operator & Modifiers

- **\$each modifier:** Let update the Marks array and add elements to it. This is being done using **\$each** modifier and **\$push** operator. With this all the elements specified in the update, will be appended to the array.
- In the following example, 3 elements has been appended to the marks where name is indu.



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44 ] }
> 
> db.result.update({"name":"indu"}, {$push: {"marks":{$each: [43,47,51]}}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> 
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 43, 47, 51 ] }
> __
```



\$slice and \$sort modifier:

- We may use sort and slice modifier with Update also. Sort may sort the specified array in ascending or descending order and slice will limit the given no of elements in the array.
- In the following example we are appending some elements to marks array, then sorting it in descending order and slicing top 5 records in the collection.



Syntax of \$sort modifier

```
$push: {
    <field>: {
        $each: [ <value1>, <value2>, ... ],
        $sort: <sort specification>
     }
}
```



Syntax of \$sort modifier

where <sort specification> is

- 1 for ascending
- -1 for descending.

It may also be used with field i.e. {field: 1 } or { field: -1 }.



db.student.update({rollno:10},{\$push:{"marks":{\$each:[34,56], \$sort:1}}})



\$slice modifier

• It will slice the array to the specified value



Syntax of \$slice modifier

```
{
    $push: {
        <field>: {
            $each: [ <value1>, <value2>, ... ],
            $slice: <num>
        }
    }
}
```



Syntax of \$slice modifier

- The <num> can be:
 - Zero To update the array <field> to an empty array [].
 - Negative To update the array <field> to contain only the last <num> elements.
 - Positive To update the array <field> contain only the first <num> elements



♦ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 45, 48, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 43, 47, 51 ] }
> db.result.update({name:"amit"},{$push:{marks:{$each:[40,78,89],$sort:-1,$slice:5}}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 89, 78, 48, 45, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 43, 47, 51 ] }
>
```



Syntax of \$position modifier

```
{
    $push: {
        <field>: {
            $each: [ <value1>, <value2>, ... ],
            $position: <num>
        }
    }
}
```

 Where <num> is the position for the element in the array, and array index starts with 0.



\$position modifier:

- Using the \$position modifier, we may specify the element location in the array at which the \$push operator inserts elements. \$position modifier is always used with \$each modifier.
- In the following example, we are adding elements at the position 2 in the array marks for indu.



♦ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 89, 78, 48, 45, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 43, 47, 51 ] }
> 
> db.result.update({name:"indu"},{$push:{marks:{$each:[91,79],$position:2}}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.result.find()
{ "_id" : ObjectId("5eb8e6c6c92747cab0598379"), "name" : "amit", "marks" : [ 89, 78, 48, 45, 41 ] }
{ "_id" : ObjectId("5eb8e6d0c92747cab059837a"), "name" : "indu", "marks" : [ 42, 44, 91, 79, 43, 47, 51 ] }
>
```



• If we specify a negative number in \$position, it will add the element in the array counting the **\$position** from the last element of the array. Example:



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe



Removing Documents,

remove() method to remove documents from collection

remove() method is used to remove a document from collection.

Syntax

db.COLLECTION_NAME.remove(DELLETION_CRITTERIA, justOne)

this method has two parameters.

One is deletion criteria and second is justOne flag.



Removing Documents,

- **DELLETION_CRITTERIA** (Optional) deletion criteria of the documents to be removed. To delete all documents in a collection, pass an empty document ({ }).
- **justOne** (Optional) if set to true or 1, then remove only one document i.e. the 1st document. Optional. If not specified i,e, the default value of false and it will delete all documents matching the deletion criteria.
- Example1:Remove only one document db.school.remove({"class":"12th"},1)
- will remove 1st record of class 12th.



Removing Documents,

```
db.school.find()
   id": ObjectId("5eb0f690ffe73fbdd0e2b1e5"), "class": "10th", "section": "A sction", "student count": "39"}
   id": "01", "class": "10th", "section": "D Section", "student count": "29" }
   id": "02", "class": "10th", "section": "C Section", "student count": "37" }
   id" : ObjectId("5eb2446ddb4e42dfda38301b"), "class" : "12th", "section" : "A section", "student count" : "33" }
  'id": ObjectId("5eb24479db4e42dfda38301c"), "class": "12th", "section": "B section", "student count": "35"
  " id" : ObjectId("5eb2448adb4e42dfda38301d"), "class" : "12th", "section" : "C section", "student count" : "30"
  " id" : ObjectId("5eb244a2db4e42dfda38301e"), "class" : "9th", "section" : "A section", "student count" : "40" }
 db.school.remove({"class":"12th"},1)
WriteResult({ "nRemoved" : 1 })
 db.school.find()
  "id": ObjectId("5eb0f690ffe73fbdd0e2b1e5"), "class": "10th", "section": "A sction", "student count": "39"}
  " id" : "01", "class" : "10th", "section" : "D Section", "student count" : "29" }
   id" : "02", "class" : "10th", "section" : "C Section", "student count" : "37" }
  "id": ObjectId("5eb24479db4e42dfda38301c"), "class": "12th", "section": "B section", "student count": "35" }
  " id" : ObjectId("5eb2448adb4e42dfda38301d"), "class" : "12th", "section" : "C section", "student count" : "30" }
  id" : ObjectId("5eb244a2db4e42dfda38301e"), "class" : "9th", "section" : "A section", "student count" : "40" }
```



Example2:Remove all records of the matching criteria

db.school.remove({"class":"12th"})

will remove all the record of class 12th.



Example2:Remove all records of the matching criteria

```
> db.school.find()
{ "_id" : ObjectId("5eb0f690ffe73fbdd0e2b1e5"), "class" : "10th", "section" : "A sction", "student count" : "39" }
{ "_id" : "01", "class" : "10th", "section" : "D Section", "student count" : "29" }
{ "_id" : "02", "class" : "10th", "section" : "C Section", "student count" : "37" }
{ "_id" : ObjectId("5eb24479db4e42dfda38301c"), "class" : "12th", "section" : "B section", "student count" : "35" }
{ "_id" : ObjectId("5eb2448adb4e42dfda38301d"), "class" : "12th", "section" : "C section", "student count" : "30" }
{ "_id" : ObjectId("5eb244a2db4e42dfda38301e"), "class" : "9th", "section" : "A section", "student count" : "40" }
> db.school.remove({"class":"12th"})
WriteResult({ "nRemoved" : 2 })
> db.school.find()
{ "_id" : ObjectId("5eb0f690ffe73fbdd0e2b1e5"), "class" : "10th", "section" : "A sction", "student count" : "39" }
{ "_id" : "01", "class" : "10th", "section" : "D Section", "student count" : "29" }
{ "_id" : "02", "class" : "10th", "section" : "C Section", "student count" : "37" }
{ "_id" : ObjectId("5eb244a2db4e42dfda38301e"), "class" : "9th", "section" : "A section", "student count" : "40" }
}
```

- WriteResult({ "nRemoved" : 2 })
- Here it has removed all the 2 records of matching criteria.



Example 3: Remove all document if no criteria is specified

db.school.remove({})

will remove all the record from the collection.

WriteResult({ "nRemoved": 2 })

Here it has removed all the 2 records of matching criteria



Example 3: Remove all document if no criteria is specified

```
> db.school.find()
> db.school.find()
{ "_id" : ObjectId("5eb0f690ffe73fbdd0e2b1e5"), "class" : "10th", "section" : "A sction", "student count" : "39" }
{ "_id" : "01", "class" : "10th", "section" : "D Section", "student count" : "29" }
{ "_id" : "02", "class" : "10th", "section" : "C Section", "student count" : "37" }
{ "_id" : ObjectId("5eb244a2db4e42dfda38301e"), "class" : "9th", "section" : "A section", "student count" : "40" }
> 
> db.school.remove({ })
WriteResult({ "nRemoved" : 4 })
> db.school.find()
```

WriteResult({ "nRemoved" : 4 })

Here it has removed all the 4 records from collection and now the collection is empty.



Document Replacement

Replacing Document in collection

- db.collection.replaceOne() method is used to replace the entire content of a document except the _id field.
- Based on the condition, whole document is replaced for the 1st matching record if there are multiple records matching the given criteria.
- The _id field remains unchanged after the replace. This method requires only field/value pairs.



Document Replacement

Syntax

db.COLLECTION_NAME.replaceOne(SELECTION_CRITERIA, Replaced_DATA)

where

SELECTION_CRITERIA -- The selection criteria for the replacing the document Replaced_DATA -- The replacement document on the matching criteria.

- The replacement document may also have different fields from the original document. _id filed may be given in new document to be replaced but it has no significance.
- In the following example, we are replacing document matching with criteria "class":"10th" with a new document having only 2 fields class and student count in school collection:



Document Replacement

```
db.school.find();
__id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40"
id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "12th", "section" : "C section", "student count" : "30"
" id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35"
" id" : ObjectId("5eb24ef5db4e42dfda383022"), "class" : "10th", "section" : "A section", "student count" : "40"
" id" : ObjectId("5eb24f1ddb4e42dfda383023"), "class" : "12th", "section" : "B section", "student count" : "39"
db.school.replaceOne({class:"10th"},{"class":"10th","student count":"NIL"})
"acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
db.school.find()
__id" : ObjectId("5eb24d83db4e42dfda38301f"), "class" : "11th", "section" : "A section", "student count" : "40"
_id" : ObjectId("5eb24d89db4e42dfda383020"), "class" : "12th", "section" : "C section", "student count" : "30"
 _id" : ObjectId("5eb24d8fdb4e42dfda383021"), "class" : "12th", "section" : "B section", "student count" : "35" ]
'id": ObjectId("5eb24ef5db4e42dfda383022"), "class": "10th", "student count": "NIL" }
"id": ObjectId("5eb24f1ddb4e42dfda383023"), "class": "12th", "section": "B\,section", "student\,count": "39"}
```

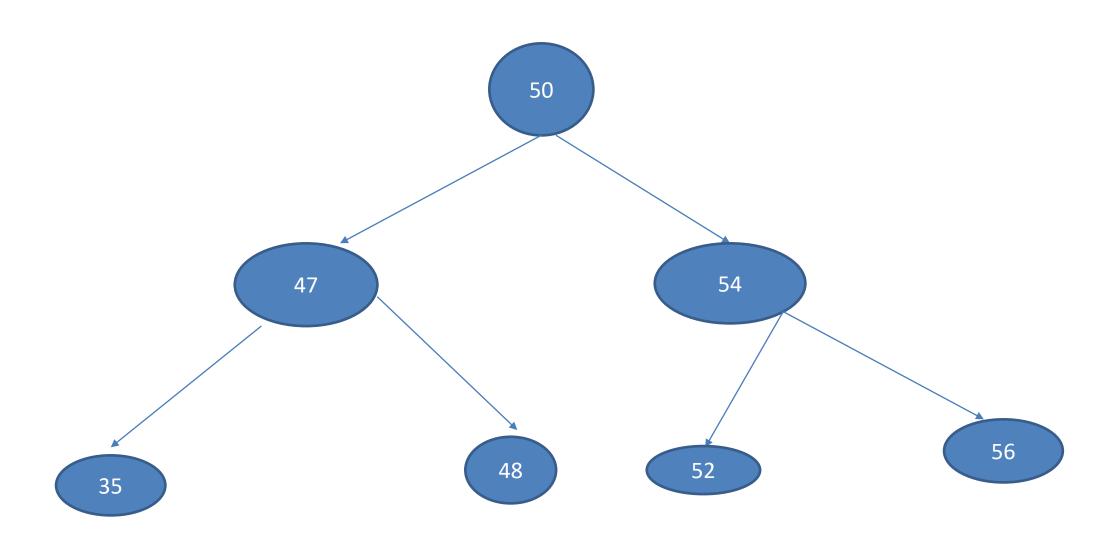


Indexing in MongoDb

 In MongoDB, index is a special type of data structure which is used to easily and quickly locate the record in a given table of the database without traversing every record in the give table. Indexed may be generated for one or multiple fields of the table. Binary Tree data structures are used by indexes.









Indexes

 plays a vital role in execution of queries in MongoDB. If no index is defined, MongoDB has to scan every record/ document of the specified collection and if index is there, it has to scan less number of documents resulting in a faster result to select those documents that match the query criteria. The index in MongoDB is similar to the indexes used in other relational databases.



Indexing in MongoDb

- In MongoDB, indexes are special data structure, which stores the value of a specific field or set of fields of a collection, ordered by the value of the field.
- The ordering of the index, ascending or descending, supports efficient equality matches and range-based query operations apart from returning sorted results by using the ordering in the index.
- The indexes are defined at collection level in MongoDB on any field or sub-field of the documents. MongoDB indexes also uses a B-tree data structure.
- It is very important and suggested to create the index on the field that will be frequently searched in a collection.



Indexing in MongoDb

Default Index in MongoDB i.e. _id

- By default, MongoDB creates a _id field if not specified in each collection and create a unique index on the _id field during the creation of a collection.
- The _id field acts like Primary Key of other RDBMS and prevents clients from inserting two documents with the same value for the _id field. The index on the _id field automatically created by MongoDB cannot be removed or dropped.



 MongoDb provides the facility to create index on any user defied key in the collection apart from the default _id index. This user created index can be ascending/descending indexes on a single field of a document as defined by the user.

Syntax

db.collection.createIndex({ <field>: <Value>})

Where

Field: is the name of the key in the collection

Value: 1 for ascending, -1 for descending



Example

Command to create a descending order index on field "student count"

db.school.createIndex({"student count": -1})



```
Select C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
db.school.find()
"_id" : ObjectId("5eb97ae983a89f15e9b7e13e"), "class" : "10th", "section" : "A Section", "student count" : 55 }
"_id" : ObjectId("5eb97afd83a89f15e9b7e13f"), "class" : "10th", "section" : "B Section", "student count" : 52 }
"_id" : ObjectId("5eb97b0b83a89f15e9b7e140"), "class" : "10th", "section" : "C Section", "student count" : 42 }
"_id" : ObjectId("5eb97b1a83a89f15e9b7e141"), "class" : "9th", "section" : "A Section", "student count" : 42 }
"_id" : ObjectId("5eb97b2883a89f15e9b7e142"), "class" : "9th", "section" : "B Section", "student count" : 51 }
"_id" : ObjectId("5eb97b3383a89f15e9b7e143"), "class" : "9th", "section" : "C Section", "student count" : 41 }
   id" : ObjectId("5eb97b4683a89f15e9b7e144"), "class" : "11th", "section" : "A Section", "student count" : 56
   id" : ObjectId("5eb97b5283a89f15e9b7e145"), "class" : "11th", "section" : "B Section", "student count" : 48
  id" : ObjectId("5eb97b5e83a89f15e9b7e146"), "class" : "11th", "section" : "C Section", "student count" : 44
  __id" : ObjectId("5eb97b6983a89f15e9b7e147"), "class" : "12th", "section" : "C Section", "student count" : 40 j
  __id" : ObjectId("5eb97b7583a89f15e9b7e148"), "class" : "12th", "section" : "B Section", "student count" : 59 j
 id" : ObjectId("5eb97b8283a89f15e9b7e149"), "class" : "12th", "section" : "A Section", "student count" : 54 }
db.school.createIndex({"student count": -1})
         "createdCollectionAutomatically" : false,
         "numIndexesBefore": 1,
         "numIndexesAfter" : 2,
         "ok" : 1
```



- After executing the command, it displays
 - "createdCollectionAutomatically": false -- It is created by user
 - "numIndexesBefore": 1 -- No of indexes before the command (only _id)
 - "numIndexesAfter": 2 -- No of indexes after the command (_id and the other one)
 - "ok": 1 Command executed Successfully
- Here, an index has been created on "student count" which means when someone searches the document based on the "student count", the search will be faster because the created index will be used for this search.



Finding the indexes in a collection

- getIndexes() method in MongoDB is used to find all the indexes created on a collection.
- Syntax

db.collection_name.getIndexes()

- Example:
 - To get all the indexes of **student** collection, the command will be:
 - > db.school.getIndexes()



Finding the indexes in a collection

C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe db.school.getIndexes() "name" : " id ", "ns" : "abc.school" }, { "key" : "student count" : -1 "name" : "student count -1", "ns" : "abc.school" }, { "key" : "class" : -1, "section" : 1 },
"name" : "class_-1_section_1", "ns" : "abc.school"



Finding the indexes in a collection

- The above example shows that three indexes are there in "student" collection. These are
 - the default index created on _id with name "_id_",
 - The single index we have created on "student count" field with name
 - "student count_-1",
 - The compound index we have created on "class" and "student count" with name "class_-1_section_1",



 A compound index is an index on two or more fields of a collection, and it can support queries based on those fields. In compound indexes, Fields can be sorted inside other fields.

Syntax

db.collection.createIndex({ <field1>: <value>, <field2>: <value>, ... })

Where

Field: is the name of the key in the collection

Value: 1 for ascending, -1 for descending



Example

 We have to create a compound index on two fields, i.e. class (descending order) and section(ascending order), to do so, the command will be

db.school.createIndex({"class": -1, "section": 1})



```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                                   db.school.find()
" id" : ObjectId("5eb97ae983a89f15e9b7e13e"), "class" : "10th", "section" : "A Section", "student count" : 55 }
 '_id" : ObjectId("5eb97afd83a89f15e9b7e13f"), "class" : "10th", "section" : "B Section", "student count" : 52 }
'_id" : ObjectId("5eb97b0b83a89f15e9b7e140"), "class" : "10th", "section" : "C Section", "student count" : 42 }
 id" : ObjectId("5eb97b1a83a89f15e9b7e141"), "class" : "9th", "section" : "A Section", "student count" : 42 }
"_id" : ObjectId("5eb97b2883a89f15e9b7e142"), "class" : "9th", "section" : "B Section", "student count" : 51 }
"_id" : ObjectId("5eb97b3383a89f15e9b7e143"), "class" : "9th", "section" : "C Section", "student count" : 41 }
 _id" : ObjectId("5eb97b4683a89f15e9b7e144"), "class" : "11th", "section" : "A Section", "student count" : 56
 _id" : ObjectId("5eb97b5283a89f15e9b7e145"), "class" : "11th", "section" : "B Section", "student count" : 48
"_id" : ObjectId("5eb97b5e83a89f15e9b7e146"), "class" : "11th", "section" : "C Section", "student count" : 44
 '_id" : ObjectId("5eb97b6983a89f15e9b7e147"), "class" : "12th", "section" : "C Section", "student count" : 40 [
 '_id" : ObjectId("5eb97b7583a89f15e9b7e148"), "class" : "12th", "section" : "B Section", "student count" : 59 j
id" : ObjectId("5eb97b8283a89f15e9b7e149"), "class" : "12th", "section" : "A Section", "student count" : 54 }
db.school.createIndex({"class": -1, "section": 1})
       "createdCollectionAutomatically" : false,
       "numIndexesBefore" : 2,
       "numIndexesAfter" : 3,
       "ok" : 1
```



- After executing the command, it displays
 - "createdCollectionAutomatically": false -- It is created by user
 - "numIndexesBefore": 2 -- No of indexes before the command (_id and "student count" created just before)
 - "numIndexesAfter": 3 -- No of indexes after the command (_id, "student count" and compound index of "class and section")
 - "ok": 1 Command executed Successfully
- This compound index of "class + section" sorts the collection first by "class" and then, within each "class" value, sort by "section".



Drop index(es) in a collection

- Dropping a index:
 - dropIndex() method is used to drop i.e. delete or remove a particular index on a collection.
- Syntax

db.collection_name.dropIndex(index)



Drop index(es) in a collection

- where index can be either "index_name"
 for example, "student count_-1" or "class_-1_section_1",
 OR
- "index specification document"
 for example, { "student count" : -1} or { class" : -1, "section" : 1}



Example:

 db.school.dropIndex("student count_-1") will delete the index with index name "student count_-1"

```
db.school.dropIndex("student count_-1")
"nIndexesWas" : 3, "ok" : 1 }
db.school.getIndexes()
                 "ns" : "abc.school"
                 "key"
                            class" : -1.
                           "section" : 1
                 "name" : "class -1 section 1",
                 "ns" : "abc.school"
```



Example:

- nIndexesWas: 3 --shows how many indexes were there before this command got executed i.e. 3
- ok: 1: -- means the command is executed successfully.
- For dropping a compound index:
 - dropIndexes() method is are used to drop i.e. delete or remove, all the indexes on a collection except the index on the _id field.



Syntax

db.collection_name.dropIndexes()

 without any parameter it will delete all indexes except on _id field, if we specify the "index_name" or "index specification document" it will delete that particular index and act like db.collection_name.dropIndex() method



Example

db.student.dropIndexes() will delete all the indexes.

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
db.school.dropIndexes()
      "nIndexesWas" : 3,
      "msg" : "non-_id indexes dropped for collection",
      "ok" : 1
db.school.getIndexes()
                      "abc.school"
```



- Msg: "non-_id indexes dropped for collection" means that the default index _id will still remain and cannot be dropped.
- nIndexesWas: 3 --shows how many indexes were there before this command got executed i.e. 3
- ok: 1: -- means the command is executed successfully.



Multikey Index

- MongoDB uses multikey indexes to index the content stored in arrays. To index a field that holds an array value, MongoDB creates an index key for each element in the array, i.e. creates separate index entries for every element of the array
- These *multikey* indexes support efficient queries against array fields. Multikey indexes can be constructed over arrays that hold both scalar values (e.g. strings, numbers) *and* nested documents.
- The multikey indexes, created by MongoDB, allow queries to select documents that contain arrays by matching on element or elements of the arrays



Syntax

db.collection_name.createIndex({ <field>: value })

- where
 - Field: is the name of the key in the collection
 - Value: 1 for ascending, -1 for descending
- While Creating an Index, it is not required to explicitly specify the multikey type, MongoDB automatically determines whether to create a multikey index, if the specified indexed field contains an array value.



Example:

• Lets consider a collection result with following documents:

```
{ "id": ObjectId("5ebced72c8f756b82a99743e"), "name":
"indu", "class": "10th", "marks": [45, 48, 43] }.
{ "id": ObjectId("5ebced91c8f756b82a99743f"), "name":
"amit", "class": "10th", "marks": [44, 46, 42]},
{ "id": ObjectId("5ebceda6c8f756b82a997440"), "name":
"geetu", "class": "10th", "marks": [43, 46, 50]},
{ " id" : ObjectId("5ebcedc6c8f756b82a997441"), "name" :
"rohit", "class": "11th", "marks": [41, 45, 49]}
```





- Lets create an index on "marks" key, which is an array on descending order
- > db.result.createIndex({"marks":-1})





- And if we run the command **getindexes**(), it will display all the indexes in the collection **result**
- > db.result.getIndexes()



```
db.result.createIndex({"marks" : -1})
      "createdCollectionAutomatically" : false,
      "numIndexesBefore" : 1,
      "numIndexesAfter" : 2,
      "ok" : 1
db.result.getIndexes()
```

 It shows that an index on marks key has been created with name "marks_-1" and stored internally, but it will be a multikey index as it has been created on an array key.



Compound Multikey Index

 Just like, creating a compound index, we may create a compound multikey index by specifying multiple keys in createIndex method

Syntax

```
db.collection.createIndex( { <field1>: <value>, <field2>: <value>, ... } )
```

where

Field: is the name of the key in the collection

Value: 1 for ascending, -1 for descending



Example

- Lets create a compound multikey index on class and marks key, both in ascending order.
- > db.result.createIndex({"class":1, "marks":1})



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

 This will create a compound index combining both the specified key, both in ascending order as we have specified the value as 1 in both the cases.



 And if we display the created indexes on the result collection, with getIndexes() method, it will display both the created indexes.



Multikey Indexes on array fields with Objects

- We may also create multikey indexes on array fields that contain nested objects i.e. creating Index inside an array object in a document.
- Lets consider the collection stock having the documents as under:



```
{ "_id" : ObjectId("5ebd64f24dd36225ea3a4076"), "item" : "tshirt", "stock" : [ { "size" : "M", "color" : "blue", "qty" : 50 }, { "size" : "S", "color" : "red", "qty" : 43 }, { "size" : "M", "color" : "blue", "qty" : 25 } ] },

{ "_id" : ObjectId("5ebd65f64dd36225ea3a4078"), "item" : "polo", "stock" : [ { "size" : "S", "color" : "white", "qty" : 40 }, { "size" : "S", "color" : "green", "qty" : 15 }, { "size" : "L", "color" : "red", "qty" : 50 }, { "size" : "XL", "color" : "blue", "qty" : 68 } ] },

{ "_id" : ObjectId("5ebd66b54dd36225ea3a4079"), "item" : "shirt", "stock" : [ { "size" : "L", "color" : "red", "qty" : 55 }, { "size" : "XL", "color" : "red", "qty" : 95 }, { "size" : "L", "color" : "blue", "qty" : 30 } ] }
```





- This stock collection has an array stock which is an object having 3 fields "size", "color" and "qty".
- Lets display only one document using findOne() from the stock collection



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
db.store.findOne()
      "_id" : ObjectId("5ebd64f24dd36225ea3a4076"),
      "item" : "tshirt",
      "stock" : [
                      "size" : "M",
                      "color" : "blue",
                      "qty" : 50
              },
{
                      "size" : "S",
                      "color" : "red",
                      "qty" : 43
              },
{
                      "size" : "M",
                      "color" : "blue",
                      "qty" : 25
```



- This clearly shows that the stock is an array object which has size, color and qty (quantity) fields. Store owner always want to query quantity available for each color and size. Therefore, they may use query command on finding stock.qty or stock.size data frequently. In order to meet the demand & supply ratio, it is good to create an index on these fields.
- To create a simple multikey index on the array object field stock.qty, command will be

db.store.createIndex({"stock.qty":1})



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

> db.store.createIndex({"stock.qty":1})
{
 "createdCollectionAutomatically" : false,
 "numIndexesBefore" : 1,
 "numIndexesAfter" : 2,
 "ok" : 1
}

 We can also create a compound multikey index on the stock.size and stock.qty fields on store collection using the following command:

db.store.createIndex({ "stock.size": 1, "stock.qty": 1 })



C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

```
> db.store.createIndex({"stock.size":1, "stock.qty":1})
{
        "createdCollectionAutomatically" : false,
        "numIndexesBefore" : 2,
        "numIndexesAfter" : 3,
        "ok" : 1
}
```

• These kind of compound multikey index may support queries with predicates that include both indexed fields as well as predicates that include only the index prefix "stock.size".



Example-1

db.store.find({"stock.size":"XL"})

```
\times
 C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                  db.store.find()
 : 50 }, { "size" : "S", "color" : "red", "qty" : 43 }, { "size" : "M", "color" : "blue", "qty" : 25 } ] }
 "_id" : ObjectId("5ebd65f64dd36225ea3a4078"), "item" : "polo", "stock" : [ { "size" : "S", "color" : "white", "qty'
: 40 }, { "size" : "S", "color" : "green", "qty" : 15 }, { "size" : "L", "color" : "red", "qty" : 50 }, { "size" :
KL", "color" : "blue", "qty" : 68 } ] }
 55 }, { "size" : "XL", "color" : "red", "qty" : 95 }, { "size" : "L", "color" : "blue", "qty" : 30 } ] }
 db.store.find({"stock.size":"XL"})
 : 40 }, { "size" : "S", "color" : "green", "qty" : 15 }, { "size" : "L", "color" : "red", "qty" : 50 }, { "size" :
XL", "color" : "blue", "qty" : 68 } ] }
 " id" : ObjectId("5ebd66b54dd36225ea3a4079"), "item" : "shirt", "stock" : [ { "size" : "L", "color" : "red", "qty"
 55 }, { "size" : "XL", "color" : "red", "qty" : 95 }, { "size" : "L", "color" : "blue", "qty" : 30 } ] }
```

This is to find all stock of size XL in the collection.



Example-2

db.store.find({"stock.size":"S", "stock.qty":{\$eq:15}})



Contents to be covered

- Aggregation Framework
- Pipeline Operations- \$match, \$sort, \$group



Aggregation Framework

- Aggregation is an operation used to process the data that returns the computed results.
- In Simple words, Aggregation groups the data from multiple documents in a collection and operates in several ways on those grouped data in order to return one combined result i.e. total number(sum), average, minimum, maximum etc out of the group selected.
- In SQL count(*) and with "group by" is an equivalent of MongoDB aggregation. In MongoDB, aggregate() method is used for the aggregation



 The aggregation pipeline is a framework for data aggregation modeled on the concept of data processing pipelines. Documents enter a multi-stage pipeline that transforms the documents into aggregated results.

Pipeline

— The MongoDB aggregation pipeline consists of stages (aggregation states). Each stage transforms the documents as they pass through the pipeline. Pipeline stages do not need to produce one output document for every input document; e.g., some stages may generate new documents or filter out documents based on the various operators and functions etc.



- The most basic pipeline stages provide *filters* that operate like queries and *document transformations* that modify the form of the output document.
- Other pipeline operations provide tools for grouping and sorting documents by specific field or fields as well as tools for aggregating the contents of arrays, including arrays of documents. In addition, pipeline stages can use operators for tasks such as calculating the average, Sum, Min, MAX or concatenating a string also.
- The pipeline provides efficient data aggregation using native operations within MongoDB, and is the preferred method for data aggregation in MongoDB.



Syntax

db.COLLECTION_NAME.aggregate(pipeline)



रा.इ.सू.प्री.सं NIELIT Aggregation pipeline operators

Name	Description
\$match	The \$match operator filters the documents stream to pass only those
	documents that match the specified condition(s) to the next pipeline stage.
	\$match uses standard MongoDB queries. For each input document, outputs
	either one document (a match) or zero documents (no match).
\$group	In MongoDB, the \$group operator groups the input documents by the
	specified expression and groups the document for each distinct grouping. An
	identifier (_id) field in the output documents contains the distinct group by
	key. The output documents can also contain computed fields that hold the
	values of some accumulator expression grouped by the \$group's
	_id(identifier) field.



Different expressions used by Aggregate function

Expression	Description
\$sum	Summates the defined values from all the documents in a collection
\$avg	Calculates the average values from all the documents in a collection
\$min	Return the minimum of all values of documents in a collection
\$max	Return the maximum of all values of documents in a collection



• **Example:** lets take a collection marks, having marks of various subjects for each students in various class given as under:

```
{ "_id" : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit",
"class" : "9th", "rollno" : 3, "sub" : "computer", "marks" : 48 },
{ "_id" : ObjectId("5ec103583b6e4f8f5b4f1149"), "name" : "rohit",
"class" : "9th", "rollno" : 3, "sub" : "english", "marks" : 44 },
```



```
{ "_id" : ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name" : "rohit", "class" : "9th",
    "rollno" : 3, "sub" : "hindi", "marks" : 41 },
{ "_id" : ObjectId("5ec103913b6e4f8f5b4f114b"), "name" : "suman", "class" : "9th",
    "rollno" : 2, "sub" : "computer", "marks" : 41 },
{ "_id" : ObjectId("5ec103a33b6e4f8f5b4f114c"), "name" : "suman", "class" : "9th",
    "rollno" : 2, "sub" : "english", "marks" : 43 },
{ "_id" : ObjectId("5ec103b53b6e4f8f5b4f114d"), "name" : "suman", "class" : "9th",
    "rollno" : 2, "sub" : "hindi", "marks" : 43 },
```



```
{ "_id" : ObjectId("5ec103cf3b6e4f8f5b4f114e"), "name" : "ajay", "class" : "10th", "rollno" : 8, "sub" :
"hindi", "marks" : 45 },
{ "_id" : ObjectId("5ec103e73b6e4f8f5b4f114f"), "name" : "ajay", "class" : "10th", "rollno" : 8, "sub" :
"english", "marks" : 39 },
{ "_id" : ObjectId("5ec103f83b6e4f8f5b4f1150"), "name" : "ajay", "class" : "10th", "rollno" : 8, "sub" :
"computer", "marks" : 44 },
{ "_id" : ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name" : "manoj", "class" : "10th", "rollno" : 9,
"sub" : "hindi", "marks" : 44 },
{ " id" : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" : "manoj", "class" : "10th", "rollno" : 9,
"sub" : "computer", "marks" : 49 },
{ "_id" : ObjectId("5ec104373b6e4f8f5b4f1153"), "name" : "manoj", "class" : "10th", "rollno" : 9,
"sub" : "english", "marks" : 40 }
```



```
      C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
      X

      > db.marks.find()
      { ".id": ObjectId("5ec103443b6e4f8f5b4f1148"), "name": "rohit", "class": "9th", "rollno": 3, "sub": "computer", "marks": 48 }

      { ".id": ObjectId("5ec103583b6e4f8f5b4f1149"), "name": "rohit", "class": "9th", "rollno": 3, "sub": "english", "marks": 44 }

      { ".id": ObjectId("5ec1036c3b6e4f8f5b4f1144"), "name": "rohit", "class": "9th", "rollno": 2, "sub": "computer", "marks": 41 }

      { ".id": ObjectId("5ec103913b6e4f8f5b4f114b"), "name": "suman", "class": "9th", "rollno": 2, "sub": "english", "marks": 43 }

      { ".id": ObjectId("5ec103a33b6e4f8f5b4f114d"), "name": "suman", "class": "9th", "rollno": 2, "sub": "english", "marks": 43 }

      { ".id": ObjectId("5ec103cf3b6e4f8f5b4f114d"), "name": "suman", "class": "9th", "rollno": 2, "sub": "hindi", "marks": 43 }

      { ".id": ObjectId("5ec103cf3b6e4f8f5b4f114d"), "name": "ajay", "class": "10th", "rollno": 8, "sub": "hindi", "marks": 45 }

      { ".id": ObjectId("5ec103cf3b6e4f8f5b4f114f"), "name": "ajay", "class": "10th", "rollno": 8, "sub": "english", "marks": 44 }

      { ".id": ObjectId("5ec103cf3b6e4f8f5b4f114f"), "name": "ajay", "class": "10th", "rollno": 8, "sub": "computer", "marks": 44 }

      { ".id": ObjectId("5ec1041b3b6e4f8f5b4f115f"), "name": "manoj", "class": "10th", "rollno": 9, "sub": "hindi", "marks": 49 }

      { ".id": ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name": "manoj", "class": "10th", "rollno": 9, "sub": "computer", "marks": 49 }

      { ".id": ObjectId("5ec104373b6e4f8f5b4f1152"), "name": "manoj", "class": "10th", "rollno": 9, "sub": "english", "marks": 49 }
```



Now, lets execute the following aggregate command:

```
db.marks.aggregate({$match:{"class":"10th"}},{$group:{_id:"$name",
    "Total_Marks":{$sum:"$marks"}}})
```



- This command executed in two states,
 - first Stage: The \$match stage filters the documents by the status field and passes to the next stage those documents that have class equal to "10th".
 - Second Stage: The \$group stage groups the documents by the sub field to calculate the sum of the amount for marks.
- This resulted into:

```
- { "_id" : "manoj", "Total_Marks" : 133 }
- { "_id" : "ajay", "Total_Marks" : 128 }
```



- Now Run another, aggregate command and see the output
- db.marks.aggregate({\$match:{"class":"9th"}},{\$group:{_id:"\$sub",Max_marks:{\$max:"\$marks"}}})



 Here, in state 1, students of class 10th are filtered out and then Maximum marks of each subject has been computed (filtered). The output is:

```
{ "_id" : "english", "Max_marks" : 44 }
{ "_id" : "computer", "Max_marks" : 48 }
{ "_id" : "hindi", "Max_marks" : 43 }
```



Aggregation pipeline operators

Name	Description
\$sort	\$ sort is used to Reorders the document stream by a specified sort key. It
	only changes the order not the documents. For each input document,
	outputs one document is there



\$sort stage Operator

- \$sort stage operator in aggregation Sorts all the input documents and returns them to the pipeline in sorted order.
- Syntax

```
{ $sort: { <field1>: <sort order>, <field2>: <sort order> ... } }
```

\$sort stage operator takes a document that specifies the field(s) to sort by and the respective sort order.

sort order can be:

- 1 Sort ascending.
- -1 Sort descending.



\$sort stage Operator

- We may use multiple fields to sort the documents in different orders as specified. If sorting on multiple fields, sort order is evaluated from left to right.
- For example, if follow the syntax, documents are first sorted by <field1>, then documents with the same <field1> values are further sorted by <field2> and so on.



Example-1

 Lets sort the documents of marks collection, 1st on subject in ascending order then on marks in descending order within the subject. The command will be

db.marks.aggregate({\$sort:{sub:1, marks:-1}})



Example-1

```
X
 C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                           db.marks.aggregate( [ {$sort:{sub:1, marks:-1}}])
  " id" : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" : "manoj", "class" : "10th", "rollno" : 9,
sub" : "computer", "marks" : 49 }
 " id" : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit", "class" : "9th", "rollno" : 3, "s
ub" : "computer", "marks" : 48 }
{ " id" : ObjectId("5ec103f83b6e4f8f5b4f1150"), "name" : "ajay", "class" : "10th", "rollno" : 8, "s
ub" : "computer", "marks" : 44 }
  " id" : ObjectId("5ec103913b6e4f8f5b4f114b"), "name" : "suman", "class" : "9th", "rollno" : 2, "s
ub" : "computer", "marks" : 41 }
 "id": ObjectId("5ec103583b6e4f8f5b4f1149"), "name": "rohit", "class": "9th", "rollno": 3, "s
ub" : "english", "marks" : 44 }
{ " id" : ObjectId("5ec103a33b6e4f8f5b4f114c"), "name" : "suman", "class" : "9th", "rollno" : 2, "s
ub" : "english", "marks" : 43 }
 "id": ObjectId("5ec104373b6e4f8f5b4f1153"), "name": "manoj", "class": "10th", "rollno": 9,
sub" : "english", "marks" : 40 }
 " id" : ObjectId("5ec103e73b6e4f8f5b4f114f"), "name" : "ajay", "class" : "10th", "rollno" : 8, "s
ub" : "english", "marks" : 39 }
" id" : ObjectId("5ec103cf3b6e4f8f5b4f114e"), "name" : "ajay", "class" : "10th", "rollno" : 8, "s
ub" : "hindi", "marks" : 45 }
 __id" : ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name" : "manoj", "class" : "10th", "rollno" : 9,
sub" : "hindi", "marks" : 44 }
 "id": ObjectId("5ec103b53b6e4f8f5b4f114d"), "name": "suman", "class": "9th", "rollno": 2, "s
ub" : "hindi", "marks" : 43 }
 "id": ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name": "rohit", "class": "9th", "rollno": 3, "s
ub" : "hindi", "marks" : 41 }
```



- The aggregation pipeline can use indexes to improve its performance during some of its stages. Some of the pipeline stages which may take advantage of indexes are as under:
 - \$match : The \$match stage can use an index to filter documents if it occurs at the beginning of a pipeline.
 - \$sort : The \$sort stage can use an index as long as it is not preceded by a \$group stage.



EXAMPLE

- If a collection named marks contains an index { sub: 1}, the following pipeline can use that index to find the first document of each group:
 - db.marks.aggregate({ \$sort:{ sub : 1} })
 - db.marks.aggregate({\$match:{"class":"10th"})



Contents to be covered

- Pipeline Operations-
 - \$project
 - \$unwind
 - \$skip
 - \$limit



pipeline operators

Name	Description
\$project	The \$project function in MongoDB passes along the documents with only the
	specified fields to the next stage in the pipeline, i.e. it Reshapes each document in
	the stream, such as by adding new fields or removing existing fields where the field
	may be the existing fields from the input documents or newly computed fields.
\$unwind	The \$unwind operator is used to deconstructing an array field from the input
	documents to output a document for each element. Each output document replaces
	the array with an element value i.e. Every output document is the input document
	with the value of the array field replaced by the element.



In MongoDB Aggregation, \$project function is passed along the documents with only the specified fields to the next stage in the pipeline. This field can be either the existing fields from the input documents or newly computed fields.

Syntax

```
{ $project: { <specifications> } }
```

Where, the specification contain

- •the inclusion of fields,
- the suppression of the _id field,
- •the addition of new fields, and
- •the resetting the values of existing fields.



Parameters

Specification	Description
<field>: <1 or true></field>	Specify the inclusion of a field.
_id: <0 or false>	Specify the suppression of the _id field. By default _id filed is included in the output document.
<field>: <expression></expression></field>	Add a new field or reset the value of an existing field.



Further,

- A field that does not exist in the document are going to include, the \$project ignores that field inclusion;
- A new field can be added and value of an existing field can be reset by specifying the field name and set its value to some expression.



 For Example 1 & 2 below, Lets have a school collection as follows

```
> db.school.find()
> db.school.find()
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb180"), "stu_id" : "A01", "stu_name" : "vinod", "class" : "2nd", "marks" : 60 }
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb181"), "stu_id" : "A02", "stu_name" : "amit", "class" : "2nd", "marks" : 55 }
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb182"), "Stu_name" : "arun" }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb183"), "stu_id" : "A04", "stu_name" : "bipin", "class" : "2nd", "marks" : 60 }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb184"), "stu_id" : "A05", "stu_name" : "anil", "class" : "3rd", "marks" : 55 }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb185"), "stu_id" : "A06", "stu_name" : "ajay", "class" : "3rd", "marks" : 70 }
```



• **Example-1:** To include some Specific Fields in the Output Documents

```
db.school.aggregate( { $project : { Stu_id: 1 , stu_name : 1 } } )
```

```
db.school.aggregate([{$project:{stu_id:1,stu_name:1}}])

{ "_id" : ObjectId("5ec233f75ddaf5609c6cb180"), "stu_id" : "A01", "stu_name" : "vinod" }

{ "_id" : ObjectId("5ec233f75ddaf5609c6cb181"), "stu_id" : "A02", "stu_name" : "amit" }

{ "_id" : ObjectId("5ec233f75ddaf5609c6cb182") }

{ "_id" : ObjectId("5ec236255ddaf5609c6cb183"), "stu_id" : "A04", "stu_name" : "bipin" }

{ "_id" : ObjectId("5ec236255ddaf5609c6cb184"), "stu_id" : "A05", "stu_name" : "anil" }

**Id" : ObjectId("5ec236255ddaf5609c6cb185"), "stu_id" : "A06", "stu_name" : "ajay" }

**Id" : ObjectId("5ec236255ddaf5609c6cb185"), "stu_id" : "A06", "stu_name" : "ajay" }
```

 The command includes only the _id, std_id, and stu_name fields in its output documents as shown above. _id field is included by default.



Example-2: To Suppress _id Field in the Output Documents
 db.school.aggregate({ \$project : { _id: 0, stu_id : 1, stu_name : 1 } })



• The above command suppresses i.e. excludes the **_id** field but includes the std_id, and stu_name fields in its output documents as we have given value as 0 for _id field.



- To Include Specific Fields from Embedded Documents
- Lets 1st consider the School collection, the same has been updated and added an embedded document fees having tution, admission and books keys. Now the collection looks like:



```
db.school.find()
: 6000, "admission" : 600, "books" : 700 } }
" id" : ObjectId("5ec233f75ddaf5609c6cb181"), "stu name" : "amit", "class" : "2nd", "marks" : 55, "fees" : { "tution"
4000, "admission" : 300, "books" : 300 } }
" id" : ObjectId("5ec233f75ddaf5609c6cb182"), "Stu name" : "arun" }
" id" : ObjectId("5ec236255ddaf5609c6cb183"), "stu name" : "bipin", "class" : "2nd", "marks" : 45, "fees" : { "tution"
: 2000, "admission" : 500, "books" : 700 } }
5000, "admission" : 300, "books" : 900 } }
4500, "admission" : 600, "books" : 200 } }
```



 We wish to include only tuition fee component from the embedded document. We may use dot (.) notation for taking only the tution from fee {fee.tution:1} or we may write it like {fee:{tution:1}}. The command for this will be:

```
db.school.aggregate( { $project: { "fee.tution": 1 } } )
or we may also write it as
db.schools.aggregate( { $project: { fees: { tution: 1 } } } );
```



```
> db.school.aggregate([{$project:{"fees.tution":1}}])
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb180"), "fees" : { "tution" : 6000 } }
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb181"), "fees" : { "tution" : 4000 } }
{ "_id" : ObjectId("5ec233f75ddaf5609c6cb182") }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb183"), "fees" : { "tution" : 2000 } }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb184"), "fees" : { "tution" : 5000 } }
{ "_id" : ObjectId("5ec236255ddaf5609c6cb185"), "fees" : { "tution" : 4500 } }
```



For adding new fields

syntax

db.collectionname.aggregate({\$project:{field:1,"newfield":"\$oldfield"}})

```
Command Prompt - mongo
Pa... )
> db.employee.aggregate([{$project: {emp code:1,employeename: "$empname"}}])
> db.employee.aggregate([{$project: {emp_code:1,employeename: "$emp_name"}}])
12 { "_id" : ObjectId("553f35a0eff0e6345e7c95e7"), "emp_code" : "E001", "employeename" : "smita" }
{ " id" : ObjectId("553f35d1eff0e6345e7c95e8"), "emp code" : "E002", "employeename" : "ankit" }
 "_id" : ObjectId("553f39edeff0e6345e7c95e9"), "emp_code" : "E003", "employeename" : "rohit" }
 " id" : ObjectId("553f3c20eff0e6345e7c95ea"), "emp code" : "E004", "employeename" : "sarita" }
> db.employee.find()
{ "_id" : ObjectId("553f35a0eff0e6345e7c95e7"), "emp_code" : "E001", "emp_name" : "smita", "date_of_join" : "16/10/2010"
 , "salary" : 8000 }
 { " id" : ObjectId("553f35d1eff0e6345e7c95e8"), "emp_code" : "E002", "emp_name" : "ankit", "date_of_join" : "16/06/2011"
  "salary" : 6000 }
 { "_id" : ObjectId("553f39edeff0e6345e7c95e9"), "emp_code" : "E003", "emp_name" : "rohit", "date_of_join" : "25/12/2010"
  "salary" : 8000 }
  "id": ObjectId("553f3c20eff0e6345e7c95ea"), "emp code": "E004", "emp name": "sarita", "date of join": "16/10/2010
 ', "salary" : "8000" }
                                    nielit1.webex.com is sharing your screen.
                                                       Stop sharing
```

IATIONALINSTITUTE OF ELECTRONICS AND IT, MMMUT CAMPUS, DEORIA ROAD, GORAKHPUR, 273010

```
Command Prompt - mongo
{ " id" : ObjectId("553f35a0eff0e6345e7c95e7"), "emp code" : "E001", "employeename" : "smita" }
{ " id" : ObjectId("553f35d1eff0e6345e7c95e8"), "emp code" : "E002", "employeename" : "ankit" }
P{ "id": ObjectId("553f39edeff0e6345e7c95e9"), "emp code": "E003", "employeename": "rohit"}
 " id" : ObjectId("553f3c20eff0e6345e7c95ea"), "emp code" : "E004", "employeename" : "sarita" }
> db.employee.find()
 { "_id" : ObjectId("553f35a0eff0e6345e7c95e7"), "emp_code" : "E001", "emp_name" : "smita", "date_of_join" : "16/10/2010"
 , "salary" : 8000 }
 { "_id" : ObjectId("553f35d1eff0e6345e7c95e8"), "emp_code" : "E002", "emp_name" : "ankit", "date_of_join" : "16/06/2011"
 . "salary" : 6000 }
12{ " id" : ObjectId("553f39edeff0e6345e7c95e9"), "emp code" : "E003", "emp name" : "rohit", "date of join" : "25/12/2010"
 , "salary" : 8000 }
  "_id" : ObjectId("553f3c20eff0e6345e7c95ea"),    "emp_code" : "E004",    "emp_name" : "sarita",    "date_of_join" : "16/10/2010
 ', "salary" : "8000" }
 > db.employee.aggregate([{$project: {emp code:1,employeename:"$emp name"}}])
{ "_id" : ObjectId("553f35a0eff0e6345e7c95e7"), "emp_code" : "E001", "employeename" : "smita" }
  { " id" : ObjectId("553f39edeff0e6345e7c95e9"), "emp code" : "E003", "employeename" : "rohit" }
  "_id" : ObjectId("553f3c20eff0e6345e7c95ea"), "emp_code" : "E004", "employeename" : "sarita" }
```



• \$unwind is another very important operator in Aggregation. It is used to deconstruct an array field from the input documents to output a document for *each* element i.e. each output document is the input document with the value of the array field replaced by the element.

Syntax

 It has two different syntax based on what information is passed. We may pass a field path operand or a document operand to deconstruct an array field.



Field Path Operand

```
{ $unwind: <field path> }
```

 Where, field path is an array field prefix with a dollar sign and enclosed in quotes, passed to \$unwind. And If the field value is null, missing, or an empty array, \$unwind does not output any document.



```
Document Operand with Options
{
    $unwind:
    {
      path: <field path>,
         preserveNullAndEmptyArrays: <boolean>
```



```
}
```

- Where,
 - field path: field path is an array field prefix with a dollar sign and enclosed in quotes
 - preserveNullAndEmptyArrays(Optional): True or False (The default value is false).



- True: If the path is null, missing, or an empty array, \$unwind outputs the document.
- False: If path is null, missing, or an empty array, \$unwind does not output a document.



• Lets consider the stock collection having the following documents, where size is an array field.

```
      ↑ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
      — □ X

      > db.stock.find()
      { "_id" : ObjectId("5ec410e112a2fa6533aa44d7"), "item" : "shirt", "color" : "blue", "size" : [ "XS", "S", "M", "L", "XL", "XXL" ] }

      , "XXL" ] }
      { "_id" : ObjectId("5ec410f712a2fa6533aa44d8"), "item" : "shirt", "color" : "white" }

      >
      >
```



- Now, if we deconstruct the array size, with the following command
 - -db.stock.aggregate({\$unwind:"\$size"})

is smiliar to

-db.stock.aggregate({ \$unwind: { path: "\$size" } })



 All the not Null/missing Array are deconstructed and converted to element in output document. In this case the shirt with white color is not in output as it has missing array field.



• Consider the output of the above example, and if we wish to display the documents with missing array field, we may use document operand with options syntax to display this.

```
db.stock.aggregate([{$unwind:{path:"$size",preserveNullAndEmptyArra
ys: true } } ] )
```



```
      ♦ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
      —
      X

      ♦ db.stock.find()
      { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": [ "XS", "S", "M", "L", "XL" ] } { "_id": ObjectId("5ec410e112a2fa6533aa44d8"), "item": "shirt", "color": "white" } }

      > db.stock.aggregate( [ { $unwind: { path: "$size", preserveNullAndEmptyArrays: true } } ] ) { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XS" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "S" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "M" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d7"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d8"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d8"), "item": "shirt", "color": "blue", "size": "XXL" } { "_id": ObjectId("5ec410e112a2fa6533aa44d8"), "item": "shirt", "color": "white" } } } }
```

 Output documents also displays the input documents where the path field is missing



\$sort stage Operator

Example

 We have stock collection, in the following example we first unwind the collection on Array and then sorted the document on the size element of output documents in descending order.

```
db.stock.aggregate( [ { $unwind: { path: "$size", preserveNullAndEmptyArrays: true } }, {$sort:{size:-1}}])
```



\$sort stage Operator



 Lets consider marks collection, having the following documents:

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                X
 db.marks.find()
 " id" : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit", "class" : "9th", "rollno" : 3, "s
ub" : "computer", "marks" : 48 }
 ub" : "english", "marks" : 44 }
 " id" : ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name" : "rohit", "class" : "9th", "rollno" : 3, "s
ub" : "hindi", "marks" : 41 }
__id__: ObjectId("5ec103913b6e4f8f5b4f114b"), "name__: "suman", "class__: "9th_, "rollno__: 2, "s
ub" : "computer", "marks" : 41 }
 "id": ObjectId("5ec103a33b6e4f8f5b4f114c"), "name": "suman", "class": "9th", "rollno": 2, "s
ub" : "english", "marks" : 43 }
 ub" : "hindi", "marks" : 43 }
 "id": ObjectId("5ec103cf3b6e4f8f5b4f114e"), "name": "ajaγ", "class": "10th", "rollno": 8, "s
ub" : "hindi", "marks" : 45 }
" id" : ObjectId("5ec103e73b6e4f8f5b4f114f"), "name" : "ajay", "class" : "10th", "rollno" : 8, "s
ub" : "english", "marks" : 39 }
" id" : ObjectId("5ec103f83b6e4f8f5b4f1150"), "name" : "ajay", "class" : "10th", "rollno" : 8, "s
ub" : "computer", "marks" : 44 }
 " id" : ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name" : "manoj", "class" : "10th", "rollno" : 9,
sub" : "hindi", "marks" : 44 }
 " id" : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" : "manoj", "class" : "10th", "rollno" : 9,
sub" : "computer", "marks" : 49 }
 "id": ObjectId("5ec104373b6e4f8f5b4f1153"), "name": "manoj", "class": "10th", "rollno": 9,
sub" : "english", "marks" : 40 }
```



pipeline operators

Name	Description
\$skip	\$skip operator Skips the first n documents where n is the specified skip
	number and passes the remaining documents unmodified to the pipeline.
	For each input document, outputs either zero documents (for the first n
	documents) or one document (if after the first n documents).



• \$skip stage operator skips the specified number of documents that pass into the stage and passes the remaining documents to the next stage in the pipeline. \$skip has no effect on the content of the documents it passes along the pipeline.

Syntax

```
{ $skip: <positive integer> }
Where,
```

 Positive integer specifies the maximum number of documents to skip to be passed with \$skip stage operator.



- Example
- Consider the result collection having 4 documents:



- If we want to skip first 3 records of result collection, then the command will be
- db.result.aggregate({ \$skip : 3 })
 - This operation skips the first 3 documents passed to it by the pipeline

```
      ◆ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
      —
      —
      ×

      > db.result.find()
      { "_id" : ObjectId("5ebced72c8f756b82a99743e"), "name" : "indu", "class" : "10th", "marks" : [ 45, 48, 43 ] }
      { "_id" : ObjectId("5ebced91c8f756b82a99743f"), "name" : "amit", "class" : "10th", "marks" : [ 44, 46, 42 ] }
      { "_id" : ObjectId("5ebceda6c8f756b82a997440"), "name" : "geetu", "class" : "10th", "marks" : [ 43, 46, 50 ] }
      { "_id" : ObjectId("5ebcf231c8f756b82a99744a"), "name" : "rohit", "class" : "11th", "marks" : [ 41, 45, 49 ] }

      > db.result.aggregate( { $skip : 3 })
      { "_id" : ObjectId("5ebcf231c8f756b82a99744a"), "name" : "rohit", "class" : "11th", "marks" : [ 41, 45, 49 ] }
```



- If we want to skip only the first records of result collection,
 then the command will be
- db.result.aggregate({ \$skip : 1 })
 - This operation skips the first documents passed to it by the pipeline :



```
      ♦ C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
      —
      X

      ♦ db.result.find()
      ("_id": ObjectId("5ebced72c8f756b82a99743e"), "name": "indu", "class": "10th", "marks": [ 45, 48, 43 ] }
      ("_id": ObjectId("5ebced91c8f756b82a99743f"), "name": "amit", "class": "10th", "marks": [ 44, 46, 42 ] }
      ("_id": ObjectId("5ebceda6c8f756b82a997440"), "name": "geetu", "class": "10th", "marks": [ 43, 46, 50 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 44, 46, 42 ] }
      ("_id": ObjectId("5ebced91c8f756b82a99743f"), "name": "amit", "class": "10th", "marks": [ 44, 46, 42 ] }
      ("_id": ObjectId("5ebceda6c8f756b82a99744a"), "name": "geetu", "class": "10th", "marks": [ 43, 46, 50 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "11th", "marks": [ 41, 45, 49 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "name": "rohit", "class": "10th", "marks": [ 44, 46, 42 ] }
      ("_id": ObjectId("5ebcf231c8f756b82a99744a"), "n
```



pipeline operators

Name	Description
\$limit	\$limit operator Passes the first n documents unmodified to the pipeline
	where n is the specified limit. For each input document, outputs either
	one document (for the first n documents) or zero documents (after the
	first n documents).



\$limit stage operator

• \$limit stage operator returns the first N documents, where N is the specified limit and these output documents may be passed to the next stage in the pipeline.

Syntax

```
{ $limit: <positive integer> }
```

Where,

 positive integer specifies the maximum number of documents to be passed, i.e. no of documents for next stage in the pipeline.

Example

— Consider the result collection having 4 documents:



- If we want to pass only first 3 records of the result collection, then the command will be
- db.result.aggregate({ \$limit : 3 })





- The above command will pass only top 3 documents.
- If we want to pass only first 2 records of the result collection, then the command will be

```
db.result.aggregate( { $limit : 2 })
```





- If we want to pass only first 6 documents of the result collection, then the command will be
- db.result.aggregate({ \$limit : 6 })
 - but if the collection is not having the specified documents, then all the documents available in the collection will be passed, no Error will be displayed. Here there are 4 documents which are passed:





Contents to be covered

- MapReduce,
- Aggregation commands



- **MapReduce** is the data processing mechanism for condensing large amount of data into useful aggregated results. This task is executed by **MapReduce** command which subsequently perform **map** and **reduce** operations.
- MapReduce is used for processing large data sets. In straightforward terms, the MapReduce command takes two primary inputs, the mapper function, and the reducer function.



- The Mapper function reads the collection of data and build the Map with the required fields that which are required to process and group them into one array. Subsequently, this keyvalue pair is fed into the Reducer, which further transform the values. The two phases of Map/Reduce is:
 - map phase: filter / transform / convert data
 - reduce phase: perform aggregations over the data



Syntax



<u>Syntax</u>

- Where,
 - map function maps a value with a key and emits a key-value pair. It's a JavaScript function
 - reduce function reduces or groups all the documents having the same key. It's a JavaScript function
 - out specifies the location of the map-reduce query result



Syntax

- query specifies the optional selection criteria for selecting documents
- The map-reduce function first queries the collection, then maps the result documents to emit key-value pairs, which is then reduced based on the keys that have multiple values.



- Example:
- Lets have the marks collection having the student name, roll, no, class and subject marks in each document.

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
db.marks.find()
     : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit", "class" : "9th", "rollno" : 3, "sub" : "computer", "marks" : 48 }
       ObjectId("5ec103583b6e4f8f5b4f1149"), "name": "rohit", "class": "9th", "rollno": 3, "sub"
       ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name" :
                                                       "rohit", "class" : "9th",
                                                                                 "rollno" : 3, "sub"
                                                       "suman", "class" : "9th",
                                                                                 "rollno" : 2, "sub"
                                                                                                        "computer", "marks"
       ObjectId("5ec103913b6e4f8f5b4f114b"), "name" :
       ObjectId("5ec103a33b6e4f8f5b4f114c"), "name" :
                                                       "suman", "class" : "9th", "rollno" : 2, "sub"
                                                       "suman", "class" : "9th", "rollno" : 2, "sub" :
       ObjectId("5ec103b53b6e4f8f5b4f114d"), "name":
                                                       "ajay", "class" : "10th", "rollno" : 8, "sub" :
       ObjectId("5ec103cf3b6e4f8f5b4f114e"), "name" :
       ObjectId("5ec103e73b6e4f8f5b4f114f"), "name"
                                                       "ajay", "class" : "10th", "rollno" : 8, "sub"
       ObjectId("5ec103f83b6e4f8f5b4f1150"), "name" :
                                                       "ajay",
                                                               "class" : "10th", "rollno" : 8, "sub" :
                                                                "class" : "10th", "rollno" : 9, "sub"
     : ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name" :
                                                       "manoj",
     : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" :
                                                       "manoj", "class" : "10th", "rollno" : 9, "sub" : "computer", "marks" : 49
     : ObjectId("5ec104373b6e4f8f5b4f1153"), "name" : "manoj", "class" : "10th", "rollno" : 9, "sub" : "english", "marks" : 40 }
```



- Now, we will use a mapReduce function on marks collection to select all sum of marks of all the students in class 9th. The candidates of the class will be grouped based on the "name" and then the sum the marks of each student. The output of the operation will be saved as collection "result_9th".
- The code will be



```
>db.marks.mapReduce(
 function() { emit(this.name,this.marks); },
function(key, values) {return Array.sum(values)},
  query:{class:"9th"},
  out:"result 9th"
```



• As a result a new collection "result_9th" will be created and the same can be seen using show collections command. The documents in the newly created collection may be seen using find() command.



- The result of the operation displays that
 - Total 6 documents matched the query (class: "9th"),
 - The map function emitted 6 documents with key-value pairs and
 - Finally the reduce function grouped mapped documents having the same keys into 2 documents.



 We may list the collections, and also see the documents in the collection:

```
> show collections
abc
college
marks
old
result
result_9th
result_totals
school
stock
store
> db.result_9th.find()
{ "_id" : "rohit", "value" : 133 }
{ "_id" : "suman", "value" : 127 }
```



 Lets use mapReduce function on marks collection again to find sum of marks of all the students in class 10th. The candidates of the class will be grouped based on the rollno and then the count their marks of each student. The output of the operation will be saved as collection "result_10th".

>



```
db.marks.mapReduce(
 function() { emit(this.rollno, this.marks); },
function(key, values) {return Array.sum(values)}, {
   query:{class:"10th"},
   out:"result_10th"
```



```
    C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

                                                               \times
db.marks.mapReduce(
     function() { emit(this.rollno, this.marks); },
     function(key, values) {return Array.sum(values)}, {
        query:{class:"10th"},
        out: "result 10th"
      "result" : "result 10th",
      "timeMillis" : 551,
      "counts" : {
               "input" : 6,
               "emit" : 6,
               "reduce" : 2,
               "output" : 2
db.result_10th.find()
" id" : 8, "value" : 128 }
" id" : 9, "value" : 133 }
```



Note:

- In above examples, the resulted documents are saved in the collection and we have to run find() command to show them.
- We can use find () command with the mapReducer command to display the resulted documents after the operation. The command for the this to find the result of 10th class is



```
>db.marks.mapReduce(
  function() { emit(this.rollno, this.marks); },
  function(key, values) {return Array.sum(values)}, {
    query:{class:"10th"},
    out:"result_10th"
    })
```



Aggregation Commands

 Aggregation functions perform operations on groups of documents and return the computed result. Aggregation operation mainly groups the data/ values from the multiple documents and perform various operations on the grouped data and returns a single or multiple results.



Aggregation Commands

- MongoDB performs aggregate operations in one of the following three ways:
 - Pipeline.
 - Map Reduce
 - Single-purpose aggregate methods and commands.
- We have already studied about these. We will emphasis once again on the aggregation commands and methods.



- These are used for specific aggregation operations on sets of data. Single-purpose aggregate commands and methods are less complex but with limited scope compared to pipeline and map reduce operations. Single-purpose aggregate commands provide straightforward semantics for common data processing options. These are
 - db.collection.count()
 - db.collection.distinct()
 - db.collection.estimatedDocumentCount()



• All of these operations aggregate documents from a single collection. These operations provide simple access to common aggregation processes, but not as flexible and capable like aggregation pipeline and map-reduce.



• db.collection.count() i.e. count

- The Count operation returns the count of documents from the collection. It takes a number of documents and depending on the match query returns the count of the documents. It is different from find() as find returns the documents while it returns the number of documents matching the query.

Syntax

– db.collectionName.count(query)

where, **Query** – The selection Criteria and it is of document type



Example:

– Lets Consider the marks collection having following documents:

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                                db.marks.find()
                                                                 "class" : "9th", "rollno" : 3, "sub" : "computer", "marks" : 48 }
  id" : ObjectId("5ec103443b6e4f8f5b4f1148"),
                                               "name" : "rohit",
  id" : ObjectId("5ec103583b6e4f8f5b4f1149"),
                                                                 "class" : "9th", "rollno" : 3, "sub" : "english", "marks" : 44 }
  id" : ObjectId("5ec1036c3b6e4f8f5b4f114a"),
                                                                           "9th",
                                                                                   "rollno" : 3.
        ObjectId("5ec103913b6e4f8f5b4f114b").
                                                                 "class" : "9th", "rollno" : 2, "sub"
      : ObjectId("5ec103a33b6e4f8f5b4f114c"),
                                                                           "9th",
                                                                                   "rollno" : 2,
                                                                 "class" : "9th", "rollno" : 2,
        ObjectId("5ec103b53b6e4f8f5b4f114d"),
  id" : ObjectId("5ec103cf3b6e4f8f5b4f114e"),
                                                                "class" : "10th", "rollno" : 8,
        ObjectId("5ec103e73b6e4f8f5b4f114f"),
                                                                "class" : "10th", "rollno" : 8, "sub"
 id" : ObjectId("5ec103f83b6e4f8f5b4f1150")
                                                                "class" : "10th", "rollno" : 8, "sub" :
                                                        "manoj", "class" : "10th", "rollno" : 9, "sub" : "hindi", "marks" : 44 }
      : ObjectId("5ec1041b3b6e4f8f5b4f1151"),
 id" : ObjectId("5ec104283b6e4f8f5b4f1152"),
                                                                 "class" : "10th", "rollno" : 9, "sub" : "computer", "marks" : 49 }
                                                        "manoj", "class" : "10th", "rollno" : 9, "sub" : "english", "marks" : 40 }
 'id": ObjectId("5ec104373b6e4f8f5b4f1153"),
```



 To find no of documents where the marks is greater than 45. The command will be

```
db.marks.count({"marks":{$gt:45}})
```

Output

2

To find no of documents where the marks is greater than 40 and class is 10th. The command will be

```
db.marks.count({$and:[{"marks":{"$gt":40}},{"class":"10th"}]})
Output
```

4



- db.collection.estimatedDocumentCount()
 - db.collection.estimatedDocumentCount() method returns the count of all documents in a collection or view.

Syntax

db.collectionName.estimatedDocumentCount()

• It does not take a query filter and instead uses metadata to return the count for a collection.



Example:

 Lets Consider the marks collection having following documents:



```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                               db.marks.find()
    " : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit", "class" : "9th", "rollno" : 3, "sub" : "computer", "marks" : 48 }
                                                                                                "sub" : "english", "marks" : 44 }
                                                        "rohit".
       ObjectId("5ec103583b6e4f8f5b4f1149"),
                                              "name"
                                                                "class" : "9th",
                                                                                  "rollno" : 3,
       ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name" :
                                                                "class" :
                                                                           "9th",
                                                                                  "rollno" : 3,
                                                                                                        "hindi", "marks" : 41 }
                                                       "suman", "class" : "9th", "rollno" : 2, "sub" :
       ObjectId("5ec103913b6e4f8f5b4f114b"), "name" :
                                                                                                        "computer", "marks" : 41 }
     : ObjectId("5ec103a33b6e4f8f5b4f114c"),
                                                       "suman", "class" : "9th", "rollno" : 2, "sub" :
                                                                                                        "english", "marks" : 43 }
                                              "name" :
       ObjectId("5ec103b53b6e4f8f5b4f114d"),
                                                        "suman", "class" : "9th", "rollno" : 2,
                                                                                                "sub" :
       ObjectId("5ec103cf3b6e4f8f5b4f114e")
                                               "name"
                                                                "class" : "10th", "rollno" : 8,
                                                        "ajay",
                                                                                  "rollno" : 8,
                                              "name" :
                                                        "ajay",
                                                                                                        "english", "marks" : 39 }
       ObjectId("5ec103e73b6e4f8f5b4f114f"),
                                                                "class" :
                                                                         "10th",
                                                                                                "sub" :
     : ObjectId("5ec103f83b6e4f8f5b4f1150"),
                                              "name" :
                                                        "ajay", "class": "10th", "rollno": 8, "sub": "computer", "marks": 44 }
                                                       "manoj", "class" : "10th", "rollno" : 9, "sub" : "hindi", "marks" : 44 }
     : ObjectId("5ec1041b3b6e4f8f5b4f1151"),
 id" : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" :
                                                        "manoj", "class" : "10th", "rollno" : 9, "sub" : "computer", "marks" : 49 }
 id" : ObjectId("5ec104373b6e4f8f5b4f1153"), "name" : "manoj", "class" : "10th", "rollno" : 9, "sub" : "english", "marks" : 40 }
```



 To find no of documents in the marks collection, the command will be

db.marks.estimatedDocumentCount()

Output

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

> db.marks.estimatedDocumentCount()

12
>
```



db.collection.distinct()

The **db.collection.distinct()** command finds the distinct values for a specified field across a single collection or view and returns the results in an array. It takes a document and depending on the match query returns the unique values for a field.

Syntax

db.collectionName.distinct(field, query)

where,

Field - The field for which to return distinct values.

Query – specifies the documents from which to retrieve the distinct values. It is optional.



- For Array Fields: If the value of the specified field is an array then the db.collectionName.distinct() considers each element of the array as a separate value.
- For Embedded Documents: We may find distinct values from embedded documents also by specifying the specific field e.g. stock.color

Example:

Lets consider the marks collection, having following documents



```
    C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

db.marks.find()
' id" : ObjectId("5ec103443b6e4f8f5b4f1148"), "name" : "rohit", "class" : "9th", "rollno" : 3, "sub" : "computer", "marks" : 48 }
     : ObjectId("5ec103583b6e4f8f5b4f1149"), "name" : "rohit", "class" : "9th", "rollno" : 3, "sub" : "english", "marks" : 44 }
      : ObjectId("5ec1036c3b6e4f8f5b4f114a"), "name" : "rohit", "class" : "9th", "rollno" : 3, "sub" : "hindi", "marks" : 41 }
     : ObjectId("5ec103913b6e4f8f5b4f114b"), "name" : "suman", "class" : "9th", "rollno" : 2, "sub" : "computer", "marks" : 41 }
 id" : ObjectId("5ec103a33b6e4f8f5b4f114c"), "name" :
                                                       "suman", "class" : "9th", "rollno" : 2, "sub" : "english", "marks" : 43 }
       ObjectId("5ec103b53b6e4f8f5b4f114d"),
                                              "name" :
                                                       "suman", "class" : "9th", "rollno" : 2,
     : ObjectId("5ec103cf3b6e4f8f5b4f114e"),
                                              "name"
                                                                "class" : "10th", "rollno" : 8,
                                                       "ajay",
                                                               "class" : "10th", "rollno" : 8,
                                                                                                        "english", "marks" : 39 }
       ObjectId("5ec103e73b6e4f8f5b4f114f"),
                                              "name"
                                                       "ajay",
                                                                                                "sub"
                                                               "class": "10th", "rollno": 8, "sub": "computer", "marks": 44 }
       ObjectId("5ec103f83b6e4f8f5b4f1150"), "name"
                                                       "ajay",
       ObjectId("5ec1041b3b6e4f8f5b4f1151"), "name":
                                                       "manoj", "class": "10th", "rollno": 9, "sub": "hindi", "marks": 44 }
      : ObjectId("5ec104283b6e4f8f5b4f1152"), "name" : "manoj", "class" : "10th", "rollno" : 9, "sub" : "computer", "marks" : 49 }
"id": ObjectId("5ec104373b6e4f8f5b4f1153"), "name": "manoj", "class": "10th", "rollno": 9, "sub": "english", "marks": 40 }
```



 To find the distinct marks among the various documents from the collection, the command is db.marks.distinct("marks")



Output

```
Select C:\Program Files\MongoDB\Server\4.2\bin\mor

built db.marks.distinct("marks")

39, 40, 41, 43, 44, 45, 48, 49]

>
```



 To find the distinct marks greater than 41 among the various documents from the collection, the command is

db.marks.distinct("marks", {"marks":{\$gt:41}})



Output

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

> db.marks.distinct("marks", {"marks":{$gt:41}})

[ 43, 44, 45, 48, 49 ]

>
```



- To find the distinct marks for the students of the class 9th among the various documents matching the query from the collection, the command is
 - db.marks.distinct("marks", {"class":"9th"})
- and for class 10th, it will be db.marks.distinct("marks", {"class":"9th"})



Output

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

> db.marks.distinct("marks", {"class":"9th"})
[ 41, 43, 44, 48 ]
> db.marks.distinct("marks", {"class":"10th"})
[ 39, 40, 44, 45, 49 ]
> ___
```



- Example: To find distinct values in case of Embedded documents
 - Lets consider the store collection having stock as embedded documents.



 To find the distinct values of sizes from the stock, the command will be

db.store.distinct("stock.size")

or

 to find the distinct color from the stock, the command will be db.store.distinct("stock.color")



Output

```
>
> db.store.distinct("stock.size")
[ "L", "M", "S", "XL" ]
>
> db.store.distinct("stock.color")
[ "blue", "green", "red", "white" ]
> __
```



- Example: To find values in array fields
 - Lets consider, the result collection having marks as an array field.

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

db.result.find()
{ "_id" : ObjectId("5ebced72c8f756b82a99743e"), "name" : "indu", "class" : "10th", "marks" : [ 45, 48, 43 ] }
{ "_id" : ObjectId("5ebced91c8f756b82a99743f"), "name" : "amit", "class" : "10th", "marks" : [ 44, 46, 42 ] }
{ "_id" : ObjectId("5ebceda6c8f756b82a997440"), "name" : "geetu", "class" : "10th", "marks" : [ 43, 46, 50 ] }
{ "_id" : ObjectId("5ebcf231c8f756b82a99744a"), "name" : "rohit", "class" : "11th", "marks" : [ 41, 45, 49 ] }
}
```



- To find all distinct marks, the command will be
 - db.result.distinct("marks")

```
> db.result.distinct("marks")
[ 41, 42, 43, 44, 45, 46, 48, 49, 50 ]
> _
```



 To find all distinct marks of Class 11th, the command will be db.result.distinct("marks", {class:"11th"})

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe

db.result.distinct("marks", {class:"11th"})
[ 41, 45, 49 ]
```



Contents to be covered

- Backup and restore
- Export and import of data



Backup and Restore

- MongoDB is one of the most popular NoSQL database engines. It is famous for being scalable, powerful, reliable and easy to use.
- The backup and restore operations create or use MongoDB-specific binary data, which preserves not only the consistency and integrity of your data but also its specific MongoDB attributes.
- Backup is also required to meet any unforeseen circumstances apart from migration of database. For these tasks, backup and restore methods are useful as long as the source and target systems are compatible.



 mongodump command can be used to create database in MongoDB. This command will dump the entire data of your server into the dump directory. The dump directory is automatically created under the current directory i.e. bin. You may carry the .bson files to new system, copy the same and use them further.

Syntax

>mongodump



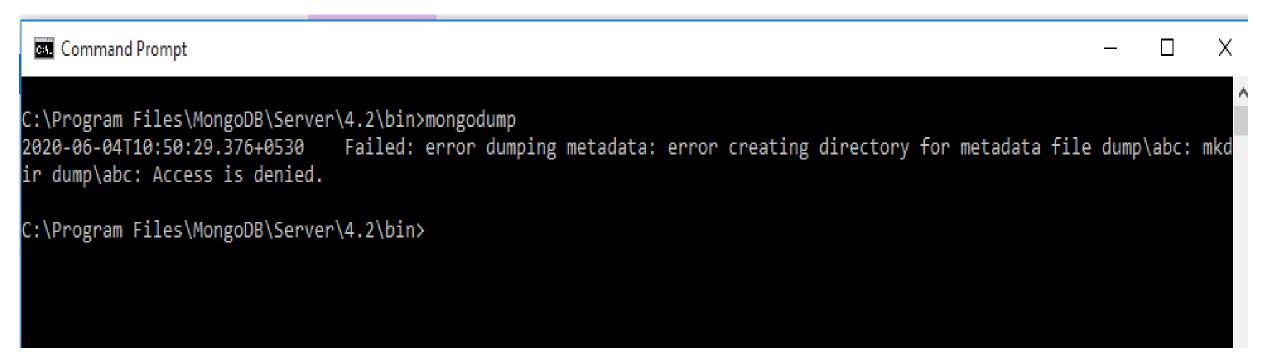
Steps to run the command

- Start mongod server.
- Open Windows command Prompt
- Go the bin directory of MongoDB Server. Where it has been installed.(for example c:\program files\mongodb\server\4.2\bin)



- Now, type the command mongodump
 >mongodump
- By default, this command will connect to the server running at port 27017 (the default port) on local host and back all data of the server to directory /bin/dump/. This command will work only if the directory has full access permission otherwise it will give you an error. If successful, it will create separate folders for each database and these folders will further have dump files for each collections.







On my system, it is giving error due to lack of permissions.
 Solution to this is to grant full permission to the files and folders of your system

```
Microsoft Windows [Version 10.0.10240]
                     (c) 2015 Microsoft Corporation. All rights reserved.
                     C:\Users\ADMIN>D:
                     D:\>CD BIN
                      ):\bin>MONGODUMP
                     2020-06-04T13:13:37.719+0530
                                                   writing admin.system.version to
                                                   done dumping admin.system.version (1 document)
                     2020-06-04T13:13:37.724+0530
                     2020-06-04T13:13:37.724+0530
                                                   writing abc.school to
                     2020-06-04T13:13:37.725+0530
                                                   done dumping abc.school (1 document)
                     D:\bin>
NATIONALINSTITUTE OF ELECTRONICS AND IT
                                                         C:\Program Files\MongoDB\Server\4.2\bin>mongodump --out=d:\dump
MMMUT CAMPUS, DEOKIA KOAD, GORAKHPUK, 27 SUIU
```



 Alternatively just give an output path with mongodump command:

Syntax

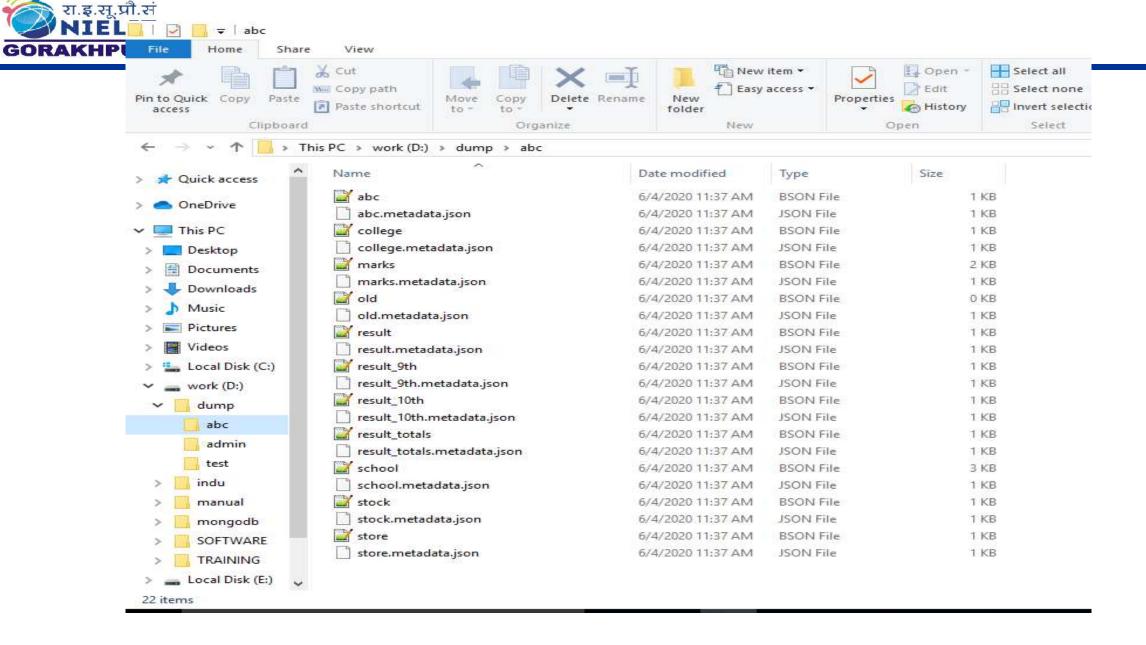
- Mongodump --out=BACKUP_DIRECTORY
- Now run the command with -out parameter and give output path (dump directory as c:\dump). This will take full backup of all the databases. All the command to be run from bin folder only.
- C:\Program Files\MongoDB\Server\4.2\bin>mongodump --out=d:\dump



```
Command Prompt
:\Program Files\MongoDB\Server\4.2\bin>mongodump --out=d:\dump
                                writing admin.system.version to
2020-06-04T11:37:33.431+0530
                                done dumping admin.system.version (1 document)
2020-06-04T11:37:33.436+0530
                                writing abc.marks to
2020-06-04T11:37:33.437+0530
                                done dumping abc.marks (12 documents)
2020-06-04T11:37:33.442+0530
                                writing abc.store to
2020-06-04T11:37:33.443+0530
                                done dumping abc.store (3 documents)
2020-06-04T11:37:33.446+0530
                                writing abc.college to
2020-06-04T11:37:33.447+0530
                                done dumping abc.college (2 documents)
2020-06-04T11:37:33.450+0530
                                writing abc.result 10th to
2020-06-04T11:37:33.451+0530
                                done dumping abc.result 10th (2 documents)
2020-06-04T11:37:33.454+0530
2020-06-04T11:37:33.455+0530
                                writing abc.result 9th to
                                done dumping abc.result 9th (2 documents)
2020-06-04T11:37:33.459+0530
2020-06-04T11:37:33.460+0530
                                writing abc.result totals to
                                done dumping abc.result totals (2 documents)
2020-06-04T11:37:33.463+0530
                                writing abc.stock to
2020-06-04T11:37:33.463+0530
                                done dumping abc.stock (2 documents)
2020-06-04T11:37:33.467+0530
                                writing abc.abc to
2020-06-04T11:37:33.468+0530
                                done dumping abc.abc (2 documents)
2020-06-04T11:37:33.471+0530
                                writing abc.old to
2020-06-04T11:37:33.472+0530
                                done dumping abc.old (0 documents)
2020-06-04T11:37:33.474+0530
                                writing abc.result to
2020-06-04T11:37:33.746+0530
                                writing test.inventory to
2020-06-04T11:37:33.752+0530
                                done dumping abc.result (4 documents)
2020-06-04T11:37:33.754+0530
                                writing abc.school to
2020-06-04T11:37:33.760+0530
                                done dumping test.inventory (10 documents)
2020-06-04T11:37:33.761+0530
                                done dumping abc.school (36 documents)
2020-06-04T11:37:33.766+0530
:\Program Files\MongoDB\Server\4.2\bin>
```



 The command is successful, it will create separate create folders for each database and these folders will further have dump files for each collections.





 Here, we can see that database folders abc, admin and test has been created in dump directory and abc(database folder) has the collections created under abc database.

•

mongoDB provide various options to take customized backup.



- mongodump –host="HOST_NAME:PORT_NUMBER" -out=BACKUP_DIRECTORY
 - This command is used to take backup of any particular Mongod Instance all databases of specified mongod instance. In case of replicaset, it shall be replica set name followed by the port number. -out is optional.
- Example
- mongodump --host="127.0.0.1:27017" --out=d:\dump1



```
\times
 Command Prompt
                                                                                                                 C:\Program Files\MongoDB\Server\4.2\bin>mongodump --host="127.0.0.1:27017" --out=d:\dump1
                                writing admin.system.version to
2020-06-04T11:46:36.172+0530
                                done dumping admin.system.version (1 document)
2020-06-04T11:46:36.175+0530
                                writing abc.school to
2020-06-04T11:46:36.177+0530
                                done dumping abc.school (36 documents)
2020-06-04T11:46:36.182+0530
2020-06-04T11:46:36.184+0530
                                writing abc.store to
                                writing abc.marks to
2020-06-04T11:46:36.189+0530
                                done dumping abc.store (3 documents)
2020-06-04T11:46:36.199+0530
                                done dumping abc.marks (12 documents)
2020-06-04T11:46:36.203+0530
2020-06-04T11:46:36.204+0530
                                writing abc.college to
                                writing abc.result 10th to
2020-06-04T11:46:36.209+0530
2020-06-04T11:46:36.211+0530
                                writing test.inventory to
                                done dumping abc.college (2 documents)
2020-06-04T11:46:36.219+0530
2020-06-04T11:46:36.227+0530
                                done dumping abc.result 10th (2 documents)
                                done dumping test.inventory (10 documents)
2020-06-04T11:46:36.231+0530
2020-06-04T11:46:36.232+0530
                                writing abc.result 9th to
                                writing abc.result to
2020-06-04T11:46:36.235+0530
2020-06-04T11:46:36.238+0530
                                writing abc.result totals to
                                writing abc.stock to
2020-06-04T11:46:36.242+0530
                                done dumping abc.result 9th (2 documents)
2020-06-04T11:46:36.243+0530
                                writing abc.abc to
2020-06-04T11:46:36.247+0530
2020-06-04T11:46:36.248+0530
                                done dumping abc.result (4 documents)
                                done dumping abc.stock (2 documents)
2020-06-04T11:46:36.251+0530
                                done dumping abc.result totals (2 documents)
2020-06-04T11:46:36.253+0530
2020-06-04T11:46:36.255+0530
                                writing abc.old to
2020-06-04T11:46:36.257+0530
                                done dumping abc.abc (2 documents)
                                done dumping abc.old (0 documents)
2020-06-04T11:46:36.260+0530
C:\Program Files\MongoDB\Server\4.2\bin>
```



Creating Backup of particular database

- mongodump --db=DB_name --out=BACKUP_DIRECTORY
 - We may take backup of specific database instead of all the databases,
 by specifying the database name with --db option. For example only abc database.
- Example
- mongodump --db="test" --out=d:\dump2



Creating Backup of particular database

```
C:\Program Files\MongoDB\Server\4.2\bin>mongodump --db="test" --out=d:\dump2
2020-06-04T11:56:10.858+0530 writing test.inventory to
2020-06-04T11:56:10.898+0530 done dumping test.inventory (10 documents)

C:\Program Files\MongoDB\Server\4.2\bin>
```

Creating Backup of particular collection

- mongodump --collection= "COLLECTION_name" --db =DB_NAME --out=BACKUP_DIRECTORY
 - This command can be used if we want to take the backup of specific collection of a particular database.

Example

--mongodump --collection="marks" --db="abc" --out=d:\dump2



```
C:\Program Files\MongoDB\Server\4.2\bin>mongodump --collection="marks" --db="abc" --out=d:\dump2
2020-06-04T11:59:19.736+0530 writing abc.marks to
2020-06-04T11:59:19.742+0530 done dumping abc.marks (12 documents)
```

 In all the cases, if out path is not specified, it will take backup to dump folder if permitted. It is advisable to take backup on a separate drive / folder



Restore data

- MongoDB's mongorestore command is used to restore binary backup created by mongodump. The mongorestore can restore either an entire database backup or a subset of the backup.
- Restore can be performed on the same system or on another system at same location or at different location. It will restore from all the .bson files from dump directory by default. We may copy the data backup to the dump directory and system will restore from that.



Restore data

Syntax

mongorestore

Will restore all databases from the dump folder by default

mongorestore <path-of-dumped-database>

command is used to store from specific path

 Example : Restoring from d:\dump folder mongorestore d:\dump



```
Command Prompt
                                                                                                        :\Program Files\MongoDB\Server\4.2\bin>mongorestore
 020-06-04T12:07:53.014+0530
                             using default 'dump' directory
2020-06-04T12:07:53.016+0530
                             try 'mongorestore --help' for more information
2020-06-04T12:07:53.017+0530
                             Failed: mongorestore target 'dump' invalid: CreateFile dump: The system cannot find the
file specified.
2020-06-04T12:07:53.017+0530
                             0 document(s) restored successfully. 0 document(s) failed to restore.
 :\Program Files\MongoDB\Server\4.2\bin>mongorestore d:\dump
2020-06-04T12:08:14.762+0530
                             preparing collections to restore from
                             restoring to existing collection abc.school without dropping
2020-06-04T12:08:14.768+0530
2020-06-04T12:08:14.769+0530
                             reading metadata for abc.school from d:\dump\abc\school.metadata.json
                             restoring to existing collection abc.marks without dropping
2020-06-04T12:08:14.770+0530
                             restoring to existing collection test.inventory without dropping
2020-06-04T12:08:14.771+0530
2020-06-04T12:08:14.771+0530
                             restoring to existing collection abc.store without dropping
2020-06-04T12:08:14.772+0530
                             reading metadata for abc.marks from d:\dump\abc\marks.metadata.json
2020-06-04T12:08:14.772+0530
                             restoring abc.school from d:\dump\abc\school.bson
                             reading metadata for test.inventory from d:\dump\test\inventory.metadata.json
2020-06-04T12:08:14.772+0530
                             reading metadata for abc.store from d:\dump\abc\store.metadata.json
2020-06-04T12:08:14.775+0530
                             restoring test.inventory from d:\dump\test\inventory.bson
2020-06-04T12:08:14.775+0530
2020-06-04T12:08:14.777+0530
                             restoring abc.marks from d:\dump\abc\marks.bson
                             restoring abc.store from d:\dump\abc\store.bson
2020-06-04T12:08:14.777+0530
                             continuing through error: E11000 duplicate key error collection: abc.store index: _id_ d
2020-06-04T12:08:14.830+0530
2020-06-04T12:08:14.831+0530
                             continuing through error: E11000 duplicate key error collection: abc.store index: id d
2020-06-04T12:08:14.831+0530
                             continuing through error: E11000 duplicate key error collection: abc.store index: id d
up key: { _id: ObjectId('5ebd66b54dd36225ea3a4079') }
2020-06-04T12:08:14.834+0530
                             restoring indexes for collection abc.store from metadata
2020-06-04T12:08:14.842+0530
                             finished restoring abc.store (0 documents, 3 failures)
2020-06-04T12:08:14.842+0530
                             restoring to existing collection abc.result without dropping
                             reading metadata for abc.result from d:\dump\abc\result.metadata.json
2020-06-04T12:08:14.843+0530
                             restoring abc.result from d:\dump\abc\result.bson
2020-06-04T12:08:14.846+0530
2020-06-04T12:08:14.879+0530
                             continuing through error: E11000 duplicate key error collection: abc.result index: id
continuing through error: E11000 duplicate key error collection: abc.result index: _id_
2020-06-04T12:08:14.880+0530
dup key: { _id: ObjectId('5ebced91c8f756b82a99743f') }
2020-06-04T12:08:14.881+0530
                             continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-04T12:08:14.883+0530
                             continuing through error: E11000 duplicate key error collection: abc.result index: _id_
2020-06-04T12:08:14.886+0530
                             restoring indexes for collection abc.result from metadata
3 items 1 item selected
```



Restore data

- Drop option
 - If we have already restored or database is existing, then duplicate key error will occur as shown below.



```
Command Prompt
 :\Program Files\MongoDB\Server\4.2\bin>mongorestore d:\dump
2020-06-04T12:23:26.782+0530
                           preparing collections to restore from
2020-06-04T12:23:26.789+0530
                           restoring to existing collection abc.school without dropping
2020-06-04T12:23:26.790+0530
                           restoring to existing collection test.inventory without dropping
                           reading metadata for abc.school from d:\dump\abc\school.metadata.json
2020-06-04T12:23:26.791+0530
2020-06-04T12:23:26.794+0530
                           restoring to existing collection abc.marks without dropping
2020-06-04T12:23:26.797+0530
                           restoring to existing collection abc.store without dropping
2020-06-04T12:23:26.800+0530
                           reading metadata for test.inventory from d:\dump\test\inventory.metadata.json
                           restoring abc.school from d:\dump\abc\school.bson
2020-06-04T12:23:26.804+0530
2020-06-04T12:23:26.809+0530
                           reading metadata for abc.marks from d:\dump\abc\marks.metadata.json
                           reading metadata for abc.store from d:\dump\abc\store.metadata.json
2020-06-04T12:23:26.833+0530
2020-06-04T12:23:26.834+0530
                           restoring abc.store from d:\dump\abc\store.bson
2020-06-04T12:23:26.834+0530
                           restoring test.inventory from d:\dump\test\inventory.bson
2020-06-04T12:23:26.838+0530
                           restoring abc.marks from d:\dump\abc\marks.bson
2020-06-04T12:23:26.854+0530
                           continuing through error: E11000 duplicate key error collection: abc.marks index: id d
up key: {    _id: ObjectId('5ec103443b6e4f8f5b4f1148')    }
2020-06-04T12:23:26.855+0530
                           continuing through error: E11000 duplicate key error collection: abc.marks index: id d
2020-06-04T12:23:26.856+0530
                           continuing through error: E11000 duplicate key error collection: abc.marks index: id d
up key: {        id: ObjectId('5ec1036c3b6e4f8f5b4f114a')        }
2020-06-04T12:23:26.859+0530
                           continuing through error: E11000 duplicate key error collection: abc.store index: _id_ d
continuing through error: E11000 duplicate key error collection: abc.school index: id
2020-06-04T12:23:26.861+0530
2020-06-04T12:23:26.864+0530
                           continuing through error: E11000 duplicate key error collection: test.inventory index:
2020-06-04T12:23:26.868+0530
                           continuing through error: E11000 duplicate key error collection: abc.marks index: id d
2020-06-04T12:23:26.871+0530
                           continuing through error: E11000 duplicate key error collection: abc.store index: id d
up key: {        id: ObjectId('5ebd65f64dd36225ea3a4078')        }
                           continuing through error: E11000 duplicate key error collection: abc.school index: _id_
2020-06-04T12:23:26.874+0530
continuing through error: E11000 duplicate key error collection: test.inventory index:
2020-06-04T12:23:26.877+0530
id dup key: {    id: ObjectId('5eb8e3340132bb4da6110627')    }
                           continuing through error: E11000 duplicate key error collection: abc.marks index: id d
2020-06-04T12:23:26.880+0530
2020-06-04T12:23:26.883+0530
                           continuing through error: E11000 duplicate key error collection: abc.store index: id d
2020-06-04T12:23:26.886+0530
                           continuing through error: E11000 duplicate key error collection: abc.school index: id
2020-06-04T12:23:26.889+0530
                           continuing through error: E11000 duplicate key error collection: test.inventory index:
 d dup key: { id: ObjectId('5eb8e3340132bb4da6110628') }
```



Restore data

To overcome this issue, drop option is used. Drop is necessary
if we are replacing an existing database.

Syntax

– mongorestore --drop <path-of-dumped-database>

Example:

--drop d:\dump

```
C:\Program Files\MongoDB\Server\4.2\bin>mongorestore --drop d:\dump
2020-06-04T12:26:38.227+0530
                               preparing collections to restore from
                               reading metadata for abc.school from d:\dump\abc\school.metadata.json
2020-06-04T12:26:38.365+0530
                               restoring abc.school from d:\dump\abc\school.bson
2020-06-04T12:26:38.575+0530
                               reading metadata for test.inventory from d:\dump\test\inventory.metadata.json
2020-06-04T12:26:38.765+0530
2020-06-04T12:26:38.809+0530
                               reading metadata for abc.marks from d:\dump\abc\marks.metadata.json
                                no indexes to restore
2020-06-04T12:26:38.894+0530
                               finished restoring abc.school (36 documents, 0 failures)
2020-06-04T12:26:38.895+0530
                               restoring test.inventory from d:\dump\test\inventory.bson
2020-06-04T12:26:39.044+0530
                                no indexes to restore
2020-06-04T12:26:39.049+0530
2020-06-04T12:26:39.049+0530
                               finished restoring test.inventory (10 documents, 0 failures)
                               restoring abc.marks from d:\dump\abc\marks.bson
2020-06-04T12:26:39.203+0530
2020-06-04T12:26:39.221+0530
                               no indexes to restore
                               finished restoring abc.marks (12 documents, 0 failures)
2020-06-04T12:26:39.221+0530
                               reading metadata for abc.stock from d:\dump\abc\stock.metadata.json
2020-06-04T12:26:39.353+0530
                               reading metadata for abc.college from d:\dump\abc\college.metadata.json
2020-06-04T12:26:39.587+0530
                               reading metadata for abc.store from d:\dump\abc\store.metadata.json
2020-06-04T12:26:39.623+0530
                               reading metadata for abc.result from d:\dump\abc\result.metadata.json
2020-06-04T12:26:39.672+0530
                               restoring abc.stock from d:\dump\abc\stock.bson
2020-06-04T12:26:39.747+0530
2020-06-04T12:26:39.766+0530
                                no indexes to restore
                                finished restoring abc.stock (2 documents, 0 failures)
2020-06-04T12:26:39.766+0530
                               restoring abc.college from d:\dump\abc\college.bson
2020-06-04T12:26:40.069+0530
2020-06-04T12:26:40.077+0530
                                no indexes to restore
                                finished restoring abc.college (2 documents, 0 failures)
2020-06-04T12:26:40.078+0530
                               restoring abc.store from d:\dump\abc\store.bson
2020-06-04T12:26:40.147+0530
                               restoring indexes for collection abc.store from metadata
2020-06-04T12:26:40.157+0530
2020-06-04T12:26:40.227+0530
                               restoring abc.result from d:\dump\abc\result.bson
2020-06-04T12:26:40.236+0530
                               restoring indexes for collection abc.result from metadata
                               reading metadata for abc.abc from d:\dump\abc\abc.metadata.json
2020-06-04T12:26:40.252+0530
                               reading metadata for abc.result 9th from d:\dump\abc\result 9th.metadata.json
2020-06-04T12:26:40.568+0530
                               restoring abc.abc from d:\dump\abc\abc.bson
2020-06-04T12:26:40.791+0530
                                no indexes to restore
2020-06-04T12:26:40.795+0530
                               finished restoring abc.abc (2 documents, 0 failures)
2020-06-04T12:26:40.795+0530
                               restoring abc.result_9th from d:\dump\abc\result_9th.bson
2020-06-04T12:26:41.112+0530
                                no indexes to restore
2020-06-04T12:26:41.117+0530
                               finished restoring abc.result 9th (2 documents, 0 failures)
2020-06-04T12:26:41.117+0530
                               reading metadata for abc.result totals from d:\dump\abc\result totals.metadata.json
2020-06-04T12:26:41.302+0530
                               reading metadata for abc.result 10th from d:\dump\abc\result 10th.metadata.json
2020-06-04T12:26:41.489+0530
```



Restoring Specific Database from a .bson file

- Mongorestore –db=<name-your-database-want-to-restoreas> <Name-and-path-of-dumped-database-bson-file/>
 - The command is used to store any specific database from the dump bson file. Drop is optional.

Example

```
mongorestore --drop --db=abc d:\dump\abc\marks.bson In new version,
```

mongorestore -drop -nsInclude=abc.student



```
Command Prompt
C:\Program Files\MongoDB\Server\4.2\bin>mongorestore --drop --db=abc d:\dump\abc\marks.bson
                                checking for collection data in d:\dump\abc\marks.bson
2020-06-04T12:39:31.143+0530
                                reading metadata for abc.marks from d:\dump\abc\marks.metadata.json
2020-06-04T12:39:31.204+0530
                                restoring abc.marks from d:\dump\abc\marks.bson
2020-06-04T12:39:31.341+0530
                                no indexes to restore
2020-06-04T12:39:31.655+0530
                                finished restoring abc.marks (12 documents, 0 failures)
2020-06-04T12:39:31.655+0530
                                12 document(s) restored successfully. 0 document(s) failed to restore.
2020-06-04T12:39:31.656+0530
C:\Program Files\MongoDB\Server\4.2\bin>
```

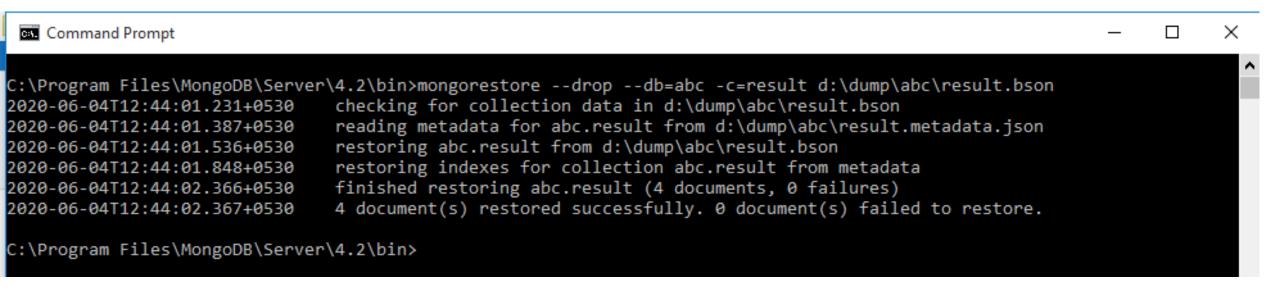


- Mongorestore --db=<name-your-database-want-to-restoreas> --c collection_name <Name-and-path-of-dumpeddatabase-bson-file/>
 - The command is used to store any specific collection from database from the dump bson file. Drop is optional.

Example

 mongorestore --drop --db=abc -c=result d:\dump\abc\result.bson







 And if – drop option is not specified, it will show duplicate error.

```
Select Command Prompt
C:\Program Files\MongoDB\Server\4.2\bin>
C:\Program Files\MongoDB\Server\4.2\bin>mongorestore --db=abc -c=result <u>d:\dump\abc\result.bson</u>
                               checking for collection data in d:\dump\abc\result.bson
2020-06-04T12:45:07.380+0530
                              restoring to existing collection abc.result without dropping
2020-06-04T12:45:07.384+0530
                               reading metadata for abc.result from d:\dump\abc\result.metadata.json
2020-06-04T12:45:07.384+0530
                               restoring abc.result from d:\dump\abc\result.bson
2020-06-04T12:45:07.386+0530
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-04T12:45:07.390+0530
continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-04T12:45:07.391+0530
dup key: {    _id: ObjectId('5ebced91c8f756b82a99743f')    }
                              continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-04T12:45:07.391+0530
dup key: { _id: ObjectId('5ebceda6c8f756b82a997440') }
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-04T12:45:07.392+0530
dup key: { id: ObjectId('5ebcf231c8f756b82a99744a') }
                               restoring indexes for collection abc.result from metadata
2020-06-04T12:45:07.448+0530
                               finished restoring abc.result (0 documents, 4 failures)
2020-06-04T12:45:07.449+0530
                               0 document(s) restored successfully. 4 document(s) failed to restore.
2020-06-04T12:45:07.450+0530
C:\Program Files\MongoDB\Server\4.2\bin>
```



 And specifying the collection, we may restore data from any other collection to the specified collection

```
Command Prompt
C:\Program Files\MongoDB\Server\4.2\bin>mongorestore --drop --db=abc -c=marks d:\dump\abc\marks.bson
2020-06-04T12:55:05.418+0530
                               checking for collection data in d:\dump\abc\marks.bson
2020-06-04T12:55:05.502+0530
                               reading metadata for abc.marks from d:\dump\abc\marks.metadata.json
2020-06-04T12:55:05.651+0530
                               restoring abc.marks from d:\dump\abc\marks.bson
2020-06-04T12:55:05.964+0530
                               no indexes to restore
                               finished restoring abc.marks (12 documents, 0 failures)
2020-06-04T12:55:05.964+0530
2020-06-04T12:55:05.965+0530
                               12 document(s) restored successfully. 0 document(s) failed to restore.
C:\Program Files\MongoDB\Server\4.2\bin>
C:\Program Files\MongoDB\Server\4.2\bin>mongorestore --drop --db=abc -c=marks d:\dump\abc\result.bson
                                checking for collection data in d:\dump\abc\result.bson
2020-06-04T12:55:15.426+0530
                               reading metadata for abc.marks from d:\dump\abc\result.metadata.json
2020-06-04T12:55:15.503+0530
2020-06-04T12:55:15.640+0530
                               restoring abc.marks from d:\dump\abc\result.bson
                               restoring indexes for collection abc.marks from metadata
2020-06-04T12:55:15.951+0530
                               finished restoring abc.marks (4 documents, 0 failures)
2020-06-04T12:55:16.451+0530
2020-06-04T12:55:16.451+0530
                               4 document(s) restored successfully. 0 document(s) failed to restore.
C:\Program Files\MongoDB\Server\4.2\bin>
```



Export and Import in MongoDB

 We may export the full collection or specific field of collection as a CSV or JSON file. And vice versa may import a CSV or JSON file into a collection.

mongoexport command to Export Data from a Collection

 mongoexport command is used to export MongoDB Collection data to a CSV or JSON file. By default, the mongoexport command connects to mongod instance running on the localhost port number 27017.

Syntax

— mongoexport --db DB_NAME --collection COLLECTION_name --fields Field_name(s) --type=[csv or JSON] --out=Name-Path-Output-File



Export and Import in MongoDB

where,

- DB_NAME Name of the Database of the Collection to be exported
- COLLECTION_name Name of Collection of DB_NAME to be exported
- Field_name(s) Name of the Field (or multiple fields separated by comma (,)) to be exported. If not specified, all the fields of the collection will be exported to JSON file.



Type – CSV or JSON format

- Name-Path-Output-File Name and path of the output (exported) file. It is also optional. But it is suggested to properly specify the name and path of the exported file.
- When exporting to CSV format, must specify the fields in the documents to be exported otherwise it will display error.
- When exporting to JSON format, _id field will also be exported by default. While the same will not be exported if exporting as CSV file until not specified in field list



 We have a result collection in abc database having multiple fields as displayed here:



- Lets export the result collection as JSON file result-all.json, means all the fields
 - mongoexport --db abc --collection result --type=json --out d:\result-all.json

```
C:\Program Files\MongoDB\Server\4.2\bin>mongoexport --db abc --collection result --type=json --out d:\result-all.json 2020-06-07T18:43:03.534+0530 connected to: mongodb://localhost/2020-06-07T18:43:03.701+0530 exported 4 records

C:\Program Files\MongoDB\Server\4.2\bin>
```



 Shows that 4 records have been exported. And if we open, the created json file result-all.json, it will show like



- Lets export the only two fields, name and marks of result collection as JSON file result-2.json:
 - mongoexport --db abc --collection result --fields name, marks -type=json --out d:\result-2.json

```
C:\Program Files\MongoDB\Server\4.2\bin>mongoexport --db abc --collection result --fields name,marks --type=json --out d:\result-2.json 2020-06-07T18:48:16.424+0530 connected to: mongodb://localhost/2020-06-07T18:48:16.478+0530 exported 4 records

C:\Program Files\MongoDB\Server\4.2\bin>
```



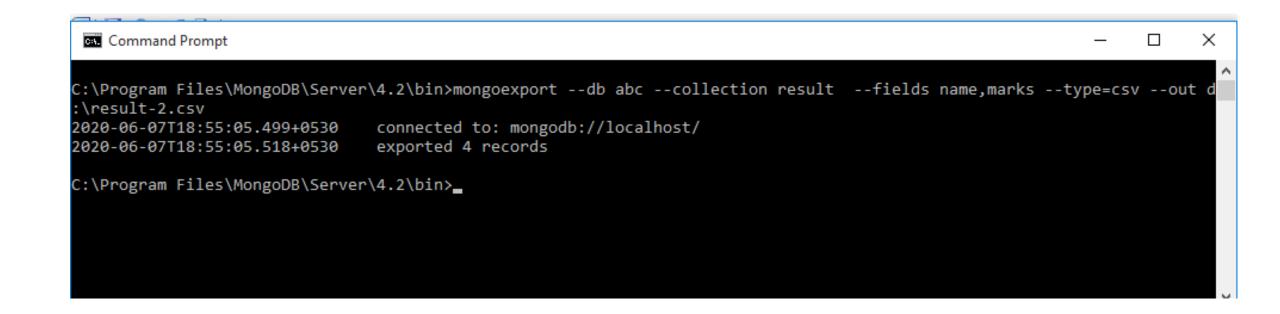
 Shows that 4 records have been exported but here only 2 fields name and marks, will be exported which may be seen in result-

2.json



- _id field will be exported by default irrespective of the number of fields of collection specified for export.
- Lets export the two fields, name and marks of result collection as CSV file result-2.csv:
 - mongoexport --db abc --collection result --fields name,marks -type=csv --out d:\result-2.csv



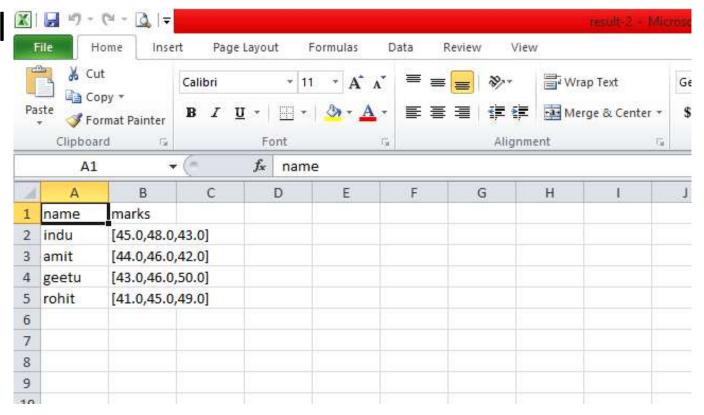




- Shows that 4 records have been exported which may be seen in result-2.csv:
- If opened in notepad



If opened in excel



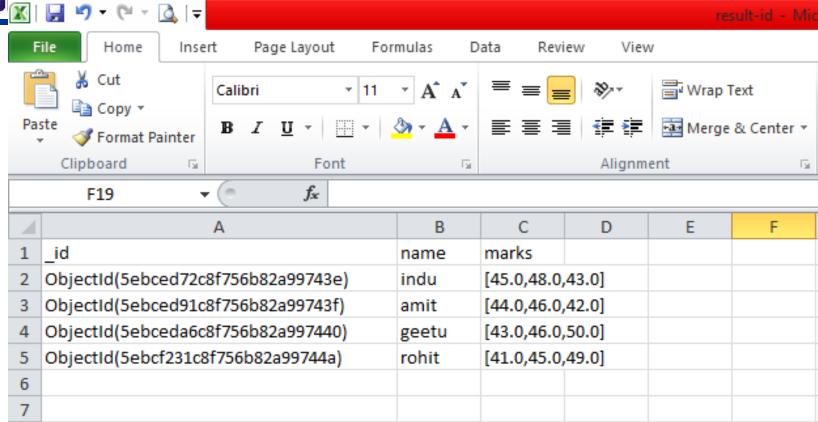


 Note: _id field will not be exported by default until not specified.

```
C:\Program Files\MongoDB\Server\4.2\bin>mongoexport --db abc --collection result --fields _id,name,marks --type=csv --out d:\result-id.csv 2020-06-07T19:00:05.607+0530 connected to: mongodb://localhost/2020-06-07T19:00:05.626+0530 exported 4 records

C:\Program Files\MongoDB\Server\4.2\bin>
```







Contents to be covered

- Importing from JSON file
- Replication,
- Advantages of replication

- mongoimport command is used to restore (import) a database from a backup(export) taken with mongoexport command.
- Importing from JSON File
 - To Import a collection from a JSON file, the syntax is as under
- Syntax

```
mongoimport --db DB_NAME --collection COLLECTION_name --
type=json --file=Name-of-file-to-import
```

where,

DB_NAME – Name of the Database of the Collection to be imported



mongoimport command to Import Collection

- COLLECTION_name Name of Collection of DB_NAME to be imported
- Type –JSON, it is optional. System by default take it as a JSON
- Name-of-file-to-import Name and path of the JSON file to be imported (restore).
- Note: We may specify a new collection name as well as a new database name while importing CSV or JSON file to the collection file.



mongoimport command to Import Collection

Example:

Let's assume that our backup file name is result-all.JSON which is in
 D:\ and the db name is abc and collection name is result.

mongoimport --db abc --collection result --type=json --file d:\result-all.json



```
C:\Program Files\MongoDB\Server\4.2\bin>mongoimport --db abc --collection result --type=json --file d:\result-all.json
                               connected to: mongodb://localhost/
2020-06-08T16:52:22.297+0530
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-08T16:52:22.572+0530
dup key: { _id: ObjectId('5ebced72c8f756b82a99743e') }
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-08T16:52:22.573+0530
dup key: { id: ObjectId('5ebceda6c8f756b82a997440') }
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
2020-06-08T16:52:22.573+0530
dup key: { _id: ObjectId('5ebced91c8f756b82a99743f') }
2020-06-08T16:52:22.576+0530
                               continuing through error: E11000 duplicate key error collection: abc.result index: id
dup key: { id: ObjectId('5ebcf231c8f756b82a99744a') }
                               0 document(s) imported successfully. 4 document(s) failed to import.
2020-06-08T16:52:22.578+0530
```



- In the above example, mongoimport imports the data in the JSON data from the result-all.json file into the collection result in the abc database. The result is showing duplicate key error because the same records are existing in the collection specified.
- We can give a new name to the collection as well as to database to which a JSON data file is imported. In the code below, the same JSON file of above example is used but imported to abcnew database as resultnew collection.
- mongoimport --db abcnew --collection resultnew --type=json --file d:\result-all.json



```
C:\Program Files\MongoDB\Server\4.2\bin>mongoimport --db abcnew --collection resultnew --type=json --file d:\result-all.json 2020-06-08T17:00:43.873+0530 connected to: mongodb://localhost/2020-06-08T17:00:43.879+0530 4 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>_
```



Now list the databases in MongoDB and collection in the abonew database



 mongoimport may also be used to import the csv formatted data into the collection. —headerline option is used to instructs mongoimport to determine the name of the fields using the first line in the CSV file. If the CSV file don't have _id field, then while importing MongoDB automatically insert _id field to all the documents. To Import a collection from a CSV file, the syntax is as under

• Syntax

- mongoimport --db DB_NAME --collection COLLECTION_name -type=csv --headerline --file=Name-of-file-to-import
- where,



- DB_NAME Name of the Database of the Collection to be exported
- COLLECTION_name Name of Collection of DB_NAME to be exported
- Type –csv
- headerline to take 1st record of CSV file as field names.
- Name-of-file-to-import Name and path of the JSON file to be imported (restore).



- If our backup file name is result-2.csv which is in D:\ and the db name is abc and collection name is result.
- mongoimport --db abc --collection result --type=csv -headerline --file d:\result-2.csv



 The result shows the records imported to the collection. As the out CSV file don't have any _id field but while importing MongoDB automatically insert _id field to all the documents.
 We may see that



- Now we are again importing a backup file result-id.csv having
 _id filed from D:\ and the db name is abc and collection name
 is result.
- mongoimport --db abc --collection result --type=csv -headerline --file d:\result-id.csv

```
C:\Program Files\MongoDB\Server\4.2\bin>mongoimport --db abc --collection result --type=csv --headerline --file d:\resultienticsv
2020-06-08T16:56:14.743+0530 connected to: mongodb://localhost/
2020-06-08T16:56:14.781+0530 4 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>
```



 Now if we display the collection result, it will have the records as

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
db.result.find()
db.result.find()
id" : ObjectId("5ede5247591230e9450515b4"), "name" : "amit", "marks" : "[44.0,46.0,42.0]"
"id": ObjectId("5ede5247591230e9450515b5"), "name": "geetu", "marks": "[43.0,46.0,50.0]"
id" : ObjectId("5ede5247591230e9450515b6"), "name" : "rohit", "marks" : "[41.0,45.0,49.0]" "
db.result.find()
" id" : ObjectId("5ede5247591230e9450515b3"), "name" : "indu", "marks" : "[45.0,48.0,43.0]"
 id" : ObjectId("5ede5247591230e9450515b4"), "name" : "amit", "marks" : "[44.0,46.0,42.0]"
     : ObjectId("5ede5247591230e9450515b5"), "name" :
                                                "geetu", "marks" : "[43.0,46.0,50.0]
     : ObjectId("5ede5247591230e9450515b6"), "name" : "rohit", "marks" : "[41.0,45.0,49.0]
 "marks": "[45.0,48.0,43.0]"
    : "ObjectId(5ebced91c8f756b82a99743f)", "name" : "amit", "marks" : "[44.0,46.0,42.0]"
" id" : "ObjectId(5ebceda6c8f756b82a997440)", "name" : "geetu", "marks" : "[43.0,46.0,50.0]"
' id" : "ObjectId(5ebcf231c8f756b82a99744a)", "name" : "rohit", "marks" : "[41.0,45.0,49.0]"
```



Importing from CSV File

- Shows total 8 records, 4 imported in example-3 (system assigned _id) and 4 imported in example-4.
- Now, we have taken database name as abcnew1 and collection name as resultnew1. In this case the collection will be imported as a new collection name and in a different database.
- mongoimport --db abcnew1 --collection resultnew1 -type=csv --headerline --file d:\result-id.csv

```
C:\Program Files\MongoDB\Server\4.2\bin>mongoimport --db abcnew1 --collection resultnew1 --type=csv --headerline --file d:\result-id.csv 2020-06-08T20:38:39.726+0530 connected to: mongodb://localhost/2020-06-08T20:38:39.929+0530 4 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>
```



Importing from CSV File

And if we see the new database and collection, it will be like

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
 show databases
         0.001GB
        0.000GB
abcnew1 0.000GB
        0.000GB
        0.000GB
local
        0.000GB
test
        0.000GB
 use abcnew1
switched to db abcnew1
 show collections
resultnew1
 db.resultnew1.find()
  '_id" : "ObjectId(5ebced72c8f756b82a99743e)", "name" : "indu", "marks" : "[45.0,48.0,43.0]" }
   _id" : "ObjectId(5ebced91c8f756b82a99743f)", "name" : "amit", "marks" : "[44.0,46.0,42.0]" }
  "id" : "ObjectId(5ebceda6c8f756b82a997440)", "name" : "geetu", "marks" : "[43.0,46.0,50.0]" }
   _id" : "ObjectId(5ebcf231c8f756b82a99744a)", "name" : "rohit", "marks" : "[41.0,45.0,49.0]" }
```



Replication

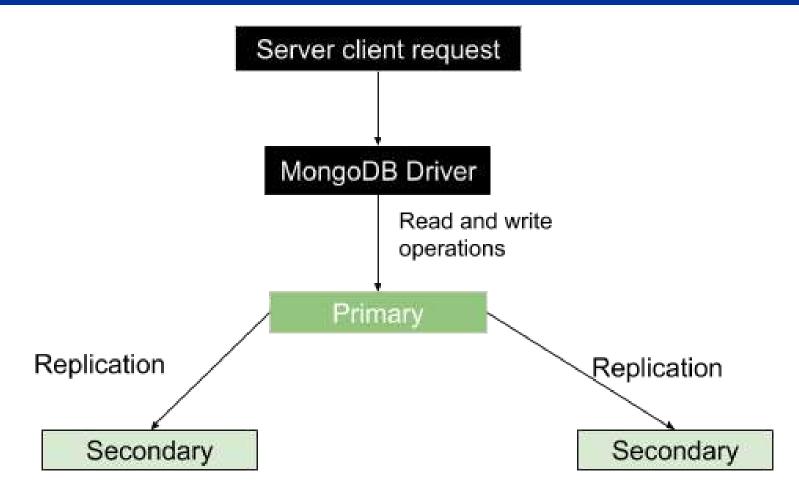
- Replication is the process of synchronizing data or simply duplication (replication) of data across multiple servers. Replication increases data availability with multiple copies of data on different database servers to meet any unforeseen circumstances.
- Replication protects a database from the loss of a single server.
 Replication also allows user to recover from hardware failure and service interruptions.
- With additional copies of the data, we may dedicate servers to disaster recovery, reporting, or backup.



Replication

- Replication is simply the process of ensuring that the same data is available on more than one Mongo DB Server.
- This is sometimes required for the purpose of increasing data availability and to meet any disaster also.
- If there is only one MongoDB Server, and it goes down for any reason, there will be no access to the data. But if data has been replicated to another server or to multiple servers, the data from another server(s) will be accessible even if the primary server fails







Replication

- A very important aspect of replication is the possibility of load balancing. If there are many users connecting to the system, instead of having everyone connect to one system, users can be connected to multiple servers so that there is an equal distribution of the load.
- In MongoDB, multiple MongDB Servers are grouped in sets called Replica sets. The Replica set will have a primary server which will accept all the write operation from clients. All other instances added to the set after this will be called the secondary instances which can be used primarily for all read operations



Advantages of Replication

- The data is safe
- There is availability of data for 24 by 7 (24 x 7)
- There is help in disaster recovery and backup of data.
- There is load balancing
- There is no downtime for maintenance (like backups, index rebuilds, compaction)
- There is read scaling (extra copies to read from)
- Replica set is transparent to the application
- Minimizes downtime for maintenance



Disadvantages of Data Replication

- There is high Cost for Hardware and other infrastructure
- There is redundant data is stored, so more space and server processing required
- It needs connectivity with all the replicated servers



Contents to be covered

- Implementation of replication
- Managing Configuration File in MongoDB
- Setting up replica set in MongoDB,
- Replication configuration ,
- Checking the Configuration
- Adding members

Implementation of Replication in MongoDB

- MongoDB implement replication by the use of replica set. A replica set is a group of Mongod instances that host the same data set.
- In a replica, one node is primary node that receives all write operations and all other instances called secondary instances, apply operations from the primary so that they have the same data set.
- Replica set can have only one primary node and rest all are the secondary nodes. In short,
 - Replica set is a group of two or more nodes. A minimum of 3 nodes are recommended.
 - In a replica set, one node is primary node and remaining nodes are secondary.



Implementation of Replication in MongoDB

- All the data replicates from primary to secondary nodes.
- There are new primary nodes which get elected in case there is automatic maintenance or failover
- Once the failed node is recovered, it again joins the replica set and works as a secondary node.



Features of Replica Set

- There is a cluster of N nodes
- Any one node can be primary
- All the write operations go to primary node
- There is automatic recovery
- There is consensus election of primary



- Before we start creating the Replica Set in MongoDB, it is desired to create separate configuration file for each replica set i.e. for each new port. These replicas will be treated as separate instance of MongoDB Server.
- These server instances can be on the different machines for regular use and it is suggested that the replication servers must be kept apart also for safety and to meet disaster recovery. However, for demo purposes we may create multiple instances of the MongoDB server on the same system also.
- For each case, we have to at least specify the ReplicaSet Name, the server Port number and Bind IP (as applicable), database and logfile in each of the configuration file.



- The configuration file of MongoDB resides in the Bin folder of MongoDB server. If we have installed the MongoDB on a Windows Machine, the configuration file "mongod.cfg" will be available in c:\program files\mongodb\server\4.2\bin directory. This file can be edited in Notepad++ easily.
- Suppose, we want to create Three (3) Replica's, then we have to first create
 directories for database and log files for each of the Secondary(replicated) servers,
 in this case three.
- Create directories
 - c:\data1, c:\data2, c:\data3 for database and
 - c:\data1\log, c:\data2\log,c:\data3\log for log files has been created.
 - C:\data1\config , c:\data2\config c:\data3\config



- Create mongo.cfg files in folders data1,data2,data3 config folder
- Edit Configuration files and specify dbpath logpath and port number
- Lets edit the configuration files: (here mongo.cfg)
 dbPath C:\data1

logpath C:\data1\log\mongod.log

port 27020



- We have to edit all the configuration files in data1, data2 and data3
- For data2
 dbPath C:\data2
 logpath C:\data2\log\mongod.log
 port 27021

Foe data3

dbPath C:\data3 log\mongod.log port 27022



Setting up a Replica Set in MongoDB

- We may convert standalone MongoDB instance to a replica set.
 Following are the steps to convert to replica set,
 - Start the MongoDB server by specifying -- replSet option.
 - Syntax
 - mongod --port "PORT" --dbpath "YOUR_DB_DATA_PATH"-- logpath "logpath" --replSet REPLICA_SET

Where,

- Port is the Port Number
- YOUR_DB_DATA_PATH Is the complete path of the DatabaseName
- REPLICA_SET is the Name of the Replica Set instance
- YOUR_log_File_PATH is the file name and path of Log file.



Setting up a Replica Set in MongoDB

• Example

- mongod --port 27020 --dbpath "c:\data1" --logpath
 "c:\data1\log\mongod.log" --replSet rset
- mongod --port 27021 --dbpath "c:\data2" --logpath
 "c:\data2\log\mongod.log" --replSet rset
- mongod --port 27022 --dbpath "c:\data3" --logpath
 "c:\data3\log\mongod.log" --replSet rset

```
Administrator: Command Prompt - mongod --port 27020 --dbpath c:\data1\db --logpath c:\data1\log\mongod.log -replSet rset
 Microsoft Windows [Version 10.0.19042.1052]
  (c) Microsoft Corporation. All rights reserved.
 C:\Users\admin>mongod --port 27020 --dbpath c:\data1\db --logpath c:\data1\log\mongod.log -replSet rset
       Administrator: Command Prompt - mongod --port 27021 --dbpath c:\data2\db --logpath c:\data2\log\mongod.log -replSet rset
                                                                                                                                      Microsoft Windows [Version 10.0.19042.1052]
      (c) Microsoft Corporation. All rights reserved.
      C:\Users\admin>mongod --port 27021 --dbpath c:\data2\db --logpath c:\data2\log\mongod.log -replSet rset
Administrator: Command Prompt - mongod --port 27022 --dbpath c:\data3\db --logpath c:\data3\log\mongod.log -replSet rset
Microsoft Windows [Version 10.0.19042.1052]
(c) Microsoft Corporation. All rights reserved.
C:\Users\admin>mongod --port 27022 --dbpath c:\data3\db --logpath c:\data3\log\mongod.log -replSet rset
```



Setting up a Replica Set in MongoDB

- This will convert MongoDB instance to replica set with name "rset" and on port 27020, 27021, 27022.
- For the replication all these three instances should be running
- Now start the command prompt and connect to this mongod instance, from the windows command prompt.

mongo --port 27020

it opens a mongo shell

Initiate a replication

• Setting replication: Initiate a replication by running below configuration

rconf={"_id":"rset",members:[{_id:0,host:"127.0.0.1:27020"}]}



- Now, run the command, rs.initiate(rconf)
 - This will initiate a new replica set.

```
14
> rs.initate(rconf)
uncaught exception: TypeError: rs.initate is not a function :
@(shell):1:1
> rs.initiate(rconf)
15{ "ok" : 1 }
rset:SECONDARY>
rset:PRIMARY>
```

 As we run this command, we can see that Prompt shows as rset.SECONDARY>



- You can see secondary connection
- Now enter ,it will turn to primary



- rs.add() commands adds members to the replica set. You can add mongod instance to replica set only if are connected to primary node.
- **db.isMaster()** command is to check whether you are connected to primary or not.
- Syntax

>rs.add(HOST_NAME:PORT)



Example

 Suppose your mongod instance name (HOST Name) is localhost (or 127.0.0.1) and it is running on port 27021. To add this instance to replica set, command will be

>rs.add("localhost:27021")

Or

>rs.add("127.0.0.1:27021")





- Lets add another member,
- >rs.add("127.0.0.1:27022")





Checking the Configuration

- To check the replica set configuration, the command is rs.conf()
 - rs.conf() method (rs.config() may also be used) without any parameter is used to check the present configuration of the application

```
t:SECONDARY>
t:PRIMARY>
t:PRIMARY> rs.config()
     "_id" : "rset",
     "version": 7,
     "term" : 4,
     "members" : [
                     " id" : 0,
                     "host" : "127.0.0.1:27020",
                     "arbiterOnly" : false,
                     "buildIndexes" : true,
                     "hidden" : false,
                     "priority" : 1,
                     "tags" : {
                     },
                     "secondaryDelaySecs" : NumberLong(0),
                     "votes" : 1
             },
                     "_id" : 1,
```

2:14 PM

```
" id" : 1,
        "host": "127.0.0.1:27021",
        "arbiterOnly" : false,
        "buildIndexes" : true,
        "hidden" : false,
        "priority" : 1,
        "tags" : {
        },
        "secondaryDelaySecs" : NumberLong(0),
        "votes" : 1
},
        "_id" : 2,
        "host": "127.0.0.1:27022",
        "arbiterOnly" : false,
        "buildIndexes" : true,
        "hidden" : false,
        "priority" : 1,
        "tags" : {
        },
```

2:15 PM

```
"tags" : {
                "secondaryDelaySecs" : NumberLong(0),
                "votes" : 1
"protocolVersion" : NumberLong(1),
"writeConcernMajorityJournalDefault" : true,
"settings" : {
        "chainingAllowed" : true,
        "heartbeatIntervalMillis" : 2000,
        "heartbeatTimeoutSecs" : 10,
        "electionTimeoutMillis" : 10000,
        "catchUpTimeoutMillis" : -1,
        "catchUpTakeoverDelayMillis" : 30000,
        "getLastErrorModes" : {
        },
        "getLastErrorDefaults" : {
                "w" : 1,
                "wtimeout" : 0
        "replicaSetId" : ObjectId("6153f3ef2f30b89a0b923317")
```

2:15 PM



Checking Status of the Replica Set

To check the status of replica set, command is rs.status()

```
rset:PRIMARY> rs.status()
                        "set" : "rset",
                        "date" : ISODate("2021-09-29T08:46:09.105Z"),
                        "myState" : 1,
                        "term" : NumberLong(4),
                        "syncSourceHost": "",
                        "syncSourceId" : -1,
                        "heartbeatIntervalMillis" : NumberLong(2000),
                        "majorityVoteCount" : 2,
                        "writeMajorityCount" : 2,
                        "votingMembersCount": 3,
                        "writableVotingMembersCount" : 3,
                        "optimes" : {
                                "lastCommittedOpTime" : {
                                        "ts" : Timestamp(1632905168, 1),
                                        "t" : NumberLong(4)
                                "lastCommittedWallTime" : ISODate("2021-09-29T08:46:08.299Z"),
                                "readConcernMajorityOpTime" : {
                                        "ts" : Timestamp(1632905168, 1),
                                        "t" : NumberLong(4)
NATIONALINST
                                                                                                              へ 口 40) 2:16 PM 9/29/2021
MMMUT CAMPL # P Type here to search
```



```
Administrator: Command Prompt - mongo --port 27020
                          "t" : NumberLong(1)
                 "priorityAtElection": 1
        "members" : [
                          " id" : 0,
                          "name": "127.0.0.1:27020",
                          "health" : 1,
                          "state" : 1,
                          "stateStr" : "PRIMARY",
                          "uptime" : 479,
                          "optime" : {
                                   "ts" : Timestamp(1632905168, 1),
                                   "t" : NumberLong(4)
                          "optimeDate" : ISODate("2021-09-29T08:46:08Z"),
                          "syncSourceHost" : "",
                          "syncSourceId" : -1,
                          "infoMessage" : "",
                          "electionTime" : Timestamp(1632905068, 1),
                          "electionDate" : ISODate("2021-09-29T08:44:28Z"),
                          "configVersion": 7,
                          "configTerm": 4,
                                                                                                    へ ED (3) 2:16 PM 9/29/2021
   Type here to search
```

Where,

state: 1->primary, 2->secondary for backup
health : 1->normal, 0->exception(means down)

It shows a lot of other details also.



Now create some database in this primary

```
Administrator: Command Prompt - mongo --port 27020
rset:PRIMARY>
rset:PRIMARY>
rset:PRIMARY> show dbs
admin 0.000GB
config 0.000GB
local
       0.000GB
rset:PRIMARY> use abc
switched to db abc
rset:PRIMARY> db.college.insert({name:"aaa",rollno:1,marks:56})
rset:PRIMARY> db.college.find()
 "_id" : ObjectId("61542918df4aa3588e8e1420"), "name" : "aaa", "rollno" : 1, "marks" : 56 }
rset:PRIMARY>
```



• Then connect to secondary and then this data should replicate in that

```
rset:SECONDARY> rs.secondaryOk()
rset:SECONDARY> show dbs
abc    0.000GB
admin    0.000GB
config    0.000GB
local    0.000GB
rset:SECONDARY>
```



Contents to be covered

- replication configuration
- Replication status
- Remove members
- Adding new port in replica set



Replication Configuration

Now if we check the configuration again, after adding the member

```
> rs.config()
{
    "_id": "indu",
    "version": 7,
    "protocolVersion": NumberLong(1),
    "writeConcernMajorityJournalDefault": true,
```

"members" : [

```
"_id":0,
    "host": "127.0.0.1:27020",
    "arbiterOnly": false,
    "buildIndexes": true,
    "hidden": false,
    "priority": 1,
    "tags" : {
     "slaveDelay": NumberLong(0),
    "votes": 1
},
```

```
"host": "127.0.0.1:27021",
 "arbiterOnly": false,
"buildIndexes": true,
"hidden": false,
"priority": 1,
"tags" : {
"slaveDelay": NumberLong(0),
"votes": 1
```

```
"_id" : 2,
"host": "127.0.0.1:27022",
"arbiterOnly": false,
"buildIndexes": true,
"hidden": false,
"priority": 1,
"tags" : {
"slaveDelay": NumberLong(0),
"votes": 1
```

},

```
"host": "127.0.0.1:27017",
    "arbiterOnly": false,
    "buildIndexes": true,
    "hidden": false,
    "priority": 1,
    "tags" : {
    "slaveDelay": NumberLong(0),
    "votes": 1
},
```

```
"_id" : 4,
  "host": "127.0.0.1:27023",
  "arbiterOnly": false,
  "buildIndexes": true,
  "hidden": false,
"priority": 1,
  "tags" : {
  "slaveDelay": NumberLong(0),
  "votes": 1
```



```
"settings" : {
"chainingAllowed": true,
"heartbeatIntervalMillis": 2000,
"heartbeatTimeoutSecs": 10,
"electionTimeoutMillis": 10000,
"catchUpTimeoutMillis": -1,
"catchUpTakeoverDelayMillis": 30000,
"getLastErrorModes" : {
"getLastErrorDefaults" : {
    "w" : 1,
    "wtimeout": 0
"replicaSetId": ObjectId("5ecffc87e7d9c40b106bbd3a")
```



We can see that two hosts has been added into the replica set.



Status of the Replica Set using rs.status()

rs.status() command is used to see the status of the replica set.

```
rs.status()
Lets check it now as we have added two hosts.
This will show the message like this:
{
    "set": "indu",
    "date": ISODate("2020-05-29T06:17:49.792Z"),
    "myState": 1,
    "term": NumberLong(6),
```



"syncingTo" : "",

```
"syncSourceHost": "",
"syncSourceId": -1,
"heartbeatIntervalMillis": NumberLong(2000),
"majorityVoteCount": 3,
"writeMajorityCount": 3,
"optimes" : {
    "lastCommittedOpTime" : {
        "ts": Timestamp(1590733061, 1),
        "t": NumberLong(6)
    "lastCommittedWallTime": ISODate("2020-05-29T06:17:41.788Z"),
    "readConcernMajorityOpTime" : {
        "ts": Timestamp(1590733061, 1),
```



"t" : NumberLong(6)

```
"readConcernMajorityWallTime":ISODate("2020-05-29T06:17:41.788Z"),
      "appliedOpTime" : {
        "ts": Timestamp(1590733061, 1),
        "t": NumberLong(6)
    "durableOpTime" : {
        "ts": Timestamp(1590733061, 1),
        "t": NumberLong(6)
    },
    "lastAppliedWallTime": ISODate("2020-05-29T06:17:41.788Z"),
    "lastDurableWallTime": ISODate("2020-05-29T06:17:41.788Z")
},
```



```
"lastStableRecoveryTimestamp": Timestamp(1590733061, 1),
"lastStableCheckpointTimestamp": Timestamp(1590733061, 1),
"electionCandidateMetrics" : {
"lastElectionReason": "electionTimeout",
"lastElectionDate": ISODate("2020-05-29T05:48:49.859Z"),
"electionTerm": NumberLong(6),
"lastCommittedOpTimeAtElection" : {
    "ts" : Timestamp(0, 0),
    "t": NumberLong(-1)
},
"lastSeenOpTimeAtElection" : {
   "ts": Timestamp(1590692225, 1),
    "t": NumberLong(4)
},
```



```
"numVotesNeeded": 2,

"priorityAtElection": 1,

"electionTimeoutMillis": NumberLong(10000),

"numCatchUpOps": NumberLong(0),

"newTermStartDate": ISODate("2020-05-29T05:48:51.633Z"),

"wMajorityWriteAvailabilityDate": ISODate("2020-05-29T05:48:52.045Z")
},
```



"members" : [

```
"_id" : 0,
"name": "127.0.0.1:27020",
"health": 1,
"state": 1,
"stateStr": "PRIMARY",
"uptime": 1755,
 "optime": {
    "ts": Timestamp(1590733061, 1),
    "t": NumberLong(6)
```



```
"optimeDate": ISODate("2020-05-29T06:17:41Z"),
    "syncingTo": ""
       "syncSourceHost": "",
    "syncSourceId": -1,
    infoMessage": "",
"electionTime": Timestamp(1590731330, 1),
"electionDate": ISODate("2020-05-29T05:48:50Z"),
"configVersion": 3,
    "self": true,
"lastHeartbeatMessage": ""
```

```
" id" : 1
"name": "127.0.0.1:27021",
"health" : 1,
"state" : 2,
"stateStr": "SECONDARY",
"uptime": 1742,
"optime" : {
    "ts": Timestamp(1590733061, 1),
    "t": NumberLong(6)
},
    "optimeDurable" : {
    "ts": Timestamp(1590733061, 1),
    "t": NumberLong(6)
```



```
"optimeDate": ISODate("2020-05-29T06:17:41Z"),
 "optimeDurableDate" : ISODate("2020-05-29T06:17:41Z"),
"lastHeartbeat": ISODate("2020-05-29T06:17:49.537Z"),
"lastHeartbeatRecv": ISODate("2020-05-29T06:17:49.077Z"),
"pingMs": NumberLong(0),
"lastHeartbeatMessage": "",
"syncingTo": "127.0.0.1:27020",
"syncSourceHost": "127.0.0.1:27020",
"syncSourceId": 0,
"infoMessage": "",
"configVersion": 3
```



{ "_id": 2,

```
"name" : "127.0.0.1:27022",
"health": 1,
"state" : 2,
"stateStr": "SECONDARY",
"uptime": 1728,
    "optime" : {
    "ts": Timestamp(1590733061, 1),
    "t" : NumberLong(6)
    "optimeDurable" : {
    "ts": Timestamp(1590733061, 1),
    "t": NumberLong(6)
    },
```



```
"optimeDate": ISODate("2020-05-29T06:17:41Z"),

"optimeDurableDate": ISODate("2020-05-29T06:17:41Z"),

"lastHeartbeat": ISODate("2020-05-29T06:17:49.541Z"),

"lastHeartbeatRecv": ISODate("2020-05-29T06:17:49.077Z"),

"pingMs": NumberLong(0),

"lastHeartbeatMessage": "",

"syncingTo": "127.0.0.1:27020",

"syncSourceHost": "127.0.0.1:27020",

"syncSourceId": 0,
```

```
"_id" : 3,
"name": "127.0.0.1:27017",
"health": 0,
"state": 8,
"stateStr": "(not reachable/healthy)",
"uptime" : 0,
"optime" : {
    "ts" : Timestamp(0, 0),
    "t": NumberLong(-1)
},
"optimeDurable" : {
    "ts" : Timestamp(0, 0),
     "t": NumberLong(-1)
},
```



```
"optimeDate": ISODate("1970-01-01T00:00:00Z"),
   "optimeDurableDate": ISODate("1970-01-01T00:00:00Z"),
    "lastHeartbeat": ISODate("2020-05-29T06:17:49.629Z"),
    "lastHeartbeatRecv": ISODate("1970-01-01T00:00:00Z"),
        "pingMs": NumberLong(0),
   "lastHeartbeatMessage": "not running with --replSet",
    "syncingTo": "",
    "syncSourceHost": "",
    "syncSourceId": -1,
    "infoMessage": "",
    "configVersion": -1
```

```
"name": "127.0.0.1:27023",
"health": 0,
"state": 8,
     "stateStr": "(not reachable/healthy)",
"uptime" : 0,
"optime" : {
    "ts" : Timestamp(0, 0),
    "t": NumberLong(-1)
},
"optimeDurable" : {
     "ts" : Timestamp(0, 0),
    "t": NumberLong(-1)
},
```



```
<u>"optimeDate" : ISODate("1970-01-01T00:00:00Z"),</u>
    "optimeDurableDate": ISODate("1970-01-01T00:00:00Z"),
    "lastHeartbeat": ISODate("2020-05-29T06:17:46.600Z"),
   "lastHeartbeatRecv": ISODate("1970-01-01T00:00:00Z"),
    "pingMs": NumberLong(0),
    "lastHeartbeatMessage": "Error connecting to 127.0.0.1:27023:: caused by::
No connection could be made because the target machine actively refused it.",
        "syncingTo": "",
    "syncSourceHost": "",
    "syncSourceId": -1,
    "infoMessage":"",
 "configVersion": -1
```



```
"ok": 1,
"$clusterTime": {
        "clusterTime": Timestamp(1590733061, 1),
"signature" : {
            "hash": BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAA="),
            "keyId": NumberLong(0)
    "operationTime": Timestamp(1590733061, 1)
```



- See the message for Port 27017 which was not started with replica Set –
 - "stateStr" : "(not reachable/healthy)",
 - "lastHeartbeatMessage": "not running with --replSet",

```
रा.इ.सू.प्रौ.सं
NIELIT
GORAKH
```

```
" id" : 3,
"name" : "127.0.0.1:27017",
"health" : 0,
"state" : 8,
"stateStr" : "(not reachable/healthy)",
"uptime" : 0,
"optime" :
        "ts" : Timestamp(0, 0),
        "t" : NumberLong(-1)
"optimeDurable" : {
        "ts" : Timestamp(0, 0),
        "t" : NumberLong(-1)
"optimeDate" : ISODate("1970-01-01T00:00:00Z"),
"optimeDurableDate" : ISODate("1970-01-01T00:00:00Z"),
"lastHeartbeat" : ISODate("2020-05-29T06:17:49.629Z"),
"lastHeartbeatRecv" : ISODate("1970-01-01T00:00:00Z"),
"pingMs" : NumberLong(0),
"lastHeartbeatMessage" : "not running with --replSet",
"syncingTo" : "",
"syncSourceHost" : "",
"syncSourceId" : -1,
"infoMessage" : "",
"configVersion" : -1
```



- If instance has not started even not created with replica Set, the messages will be like
- stateStr": "(not reachable/healthy)"
- "lastHeartbeatMessage": "Error connecting to 127.0.0.1:27023:: caused by:: No connection could be made because the target machine actively refused it.",

Command Prompt - mongo -port 27020

```
" id" : 4,
                        "name" : "127.0.0.1:27023",
                        "health" : 0,
                        "state" : 8,
                        "stateStr" : "(not reachable/healthy)",
                        "uptime" : 0,
                        "optime" : {
                                "ts" : Timestamp(0, 0),
                                "t" : NumberLong(-1)
                        "optimeDurable" : {
                                "ts" : Timestamp(0, 0),
                                "t" : NumberLong(-1)
                        "optimeDate" : ISODate("1970-01-01T00:00:00Z"),
                        "optimeDurableDate" : ISODate("1970-01-01T00:00:00Z"),
                        "lastHeartbeat" : ISODate("2020-05-29T06:17:46.600Z"),
                        "lastHeartbeatRecv" : ISODate("1970-01-01T00:00:00Z"),
                        "pingMs" : NumberLong(0),
                        "lastHeartbeatMessage" : "Error connecting to 127.0.0.1:27023 :: caused by :: No connection
could be made because the target machine actively refused it.",
                        "syncingTo" : "",
                        "syncSourceHost" : "",
                        "syncSourceId" : -1,
                        "infoMessage" : "",
                        "configVersion": -1
```

×



Remove Members

- we can remove any member(s) from the existing replica set using the rs.remove() command.
- Syntax
 - >rs.remove("HOST_NAME:PORT")
- **Example,** lets remove the member
 - >rs.remove("127.0.0.1:27023")



• If we run **rs.config()** command again, we will see that the particular port has been removed from the replica set.



Adding a new Port / instance in replica set: Steps and Cautions

- First create the folders for the database and log files
- Copy and Rename the configuration file
- Edit the configuration file with the db & log Path, Port number and bind IP, Replica name and size
- Start the MongoDB instance with Replica Set and Port no like

```
C:\Program Files\MongoDB\Server\4.2\bin>
C:\Program Files\MongoDB\Server\4.2\bin>mongod --port 27023 --dbpath "C:\data4\log\mongod.log" --replSet indu

C:\Program Files\MongoDB\Server\4.2\bin>mongod --port 27023 --dbpath "C:\data4" --logpath "C:\data4\log\mongod.log" --replSet indu

A
```



Adding a new Port / instance in replica set: Steps and Cautions

- Now Add the Host to the replica Set using rs.add () command >rs.add("127.0.0.1:27023")
- If desired, check the replica set configuration using rs.conf()
- Check the status of Replica Set with rs.status() again after adding the member properly. Now it shows that
 - "stateStr": "SECONDARY" -- Means that it is Secondary
 - "lastHeartbeatMessage": "" -- means no error and is active



Adding a new Port / instance in replica set: Steps and Cautions

```
Command Prompt - mongo -port 27020
                          "syncSourceId" : -1,
                         "infoMessage" : "",
"configVersion" : -1
                         "_id" : 6,
"name" : "127.0.0.1:27023",
                         "health" : 1,
                         "state" : 2,
                         "stateStr" : "SECONDARY",
                         "uptime" : 1969,
"optime" : {
    "ts" : Timestamp(1590736612, 1),
                                  "t" : NumberLong(6)
                         },
"optimeDurable" : {
                                  "ts" : Timestamp(1590736612, 1),
                                  "t" : NumberLong(6)
                         },
"optimeDate" : ISODate("2020-05-29T07:16:52Z"),
                         "optimeDurableDate" : ISODate("2020-05-29T07:16:52Z"),
                         "lastHeartbeat" : ISODate("2020-05-29T07:16:59.377Z"),
                         "lastHeartbeatRecv" : ISODate("2020-05-29T07:16:59.896Z").
                          "pingMs" : NumberLong(Θ),
                          "lastHeartbeatMessage" : ""
                         "syncingTo" : "127.0.0.1:27022",
                         "syncSourceHost": "127.0.0.1:27022",
                         "syncSourceId" : 2,
                         "infoMessage" : "",
                         "configVersion" : 7
        "$clusterTime" : {
                "clusterTime" : Timestamp(1590736612, 1),
                "signature" : {
                         "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAAAAAAAA."),
                         "kevId" : NumberLong(0)
       ),
"operationTime" : Timestamp(1590736612, 1)
indu:PRIMARY> __
```



Contents to be covered

- Terminologies used in replication,
- Help on replication
- Freezing a member
- Stepping down a member
- Printing the oplog of the replica set



Heartbeat

- Replication is used to keep multiple instances of the database simultaneously running usually on different systems of the same replica set.
- Heartbeat is the process to identify the current status of such servers (also called node servers) within the replica set. To do so, members of replica set send heartbeats (technically pings) to each other at every two seconds.
- In case, if a heartbeat does not return (i.e. no response) within 10 seconds, then the other members of replica set mark such delinquent member as an inaccessible member.



- A heartbeat request is a short message that checks everyone's current state. Heartbeat is the process to know the other member's state, like who's primary, from which member they need to sync from or which node is down.
- The most important activity of heartbeat is to check that the primary server live and reachable by all secondary servers. In case, if the most of the secondary servers can't reach to the primary server, then the process automatically demotes the primary server as a secondary server.



Elections

- Replica set members uses election to determine which set member will become primary as the Primary node is not accessible. If a primary becomes unavailable, elections allow the set to recover normal operations without manual intervention. Elections may also be called as the part of the failover process.
- Elections occur anytime when the primary becomes unavailable. It is important to know that the primary is the only member in the replica set which can accept write operations. And while an election is in process, the replica set has no primary and cannot accept writes. MongoDB avoids elections unless necessary.



Election Process

— During the heartbeat check, if a member can't reach the primary server then that particular member raises the **election flag** to the other members of the replica set so that other members of the replica set can't raise the same election flag within the process.



- If there is no objection against the election request, the other members will vote for the member seeking election.
- And if the member receives the majority votes from other members, then the election is successful and it will be promoted as the primary node.
- In case, if the member did not receive the majority of votes then it will remain as secondary node and maybe try to become a primary node in the future.



- Elections take time to complete .Elections are essential for independent operation of a replica set; however, while an election is in process, the replica set has no primary and cannot accept writes.
- MongoDB avoids elections unless necessary. If the network condition is healthy and most of the servers are up, then the entire elections process should happen very quickly.
- Two seconds is required to notify all the members that primary is down as heartbeat response has not been received yet and the election process starts immediately.



Oplog

- The oplog, i.e. operations log, is a special capped collection in MongoDB that keeps a rolling record of all operations that modify the data stored in your databases.
- All the database operations are applied to the primary and operations are recorded on the oplog of the primary. All the secondary members then copy and apply these operations in an asynchronous process.
- To maintain the current state of the database, all replica set members contain a copy of the oplog in the local.oplog.rs collection.
 Any secondary member may import oplog entries from any other member.



Syncing

- A log of operations called oplog is maintained by MongoDB having every write (transaction/operation) information onto the primary server. As the main aim of the replication process is to keep the same or an identical set of data on multiple servers of the Replica Set, the oplog helps in achieving this goal.
- Oplog is a capped collection and is there in the local database on the primary server. The secondary servers query this collection(oplog) to obtain the operation details so that they can replicate that data.



- Oplog is also maintained at every secondary server in which MongoDB stores each
 operation related to the replication process from the primary server.
- These log files allow any replica set members to use as a sync source for other members. Every time, the secondary server first fetch the information related to the pending operations from the primary member, then apply that operation to their own data set and then write down the logs about that operation into the oplog.
- In case, if the secondary server goes down and is going up after some time, then it starts the syncing process the last operation is done by itself according to its oplog file.
- As the operation first applied to the data and then it writes to the oplog, the secondary server may redo the operations that it has already applied to its data.



Help on Replication

• rs.help() command provides the basic help for various replica set function/ methods.



Help on Replication

```
C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
                                                                                                                           ×
rs.help()
      rs.status()
                                                   replSetGetStatus : 1 } checks repl set status
      rs.initiate()
                                                   replSetInitiate : null } initiates set with default settings
                                                   replSetInitiate : cfg } initiates set with configuration cfg
      rs.initiate(cfg)
      rs.conf()
                                                  get the current configuration object from local.system.replset
      rs.reconfig(cfg)
                                                  updates the configuration of a running replica set with cfg (disconnects)
      rs.add(hostportstr)
                                                  add a new member to the set with default attributes (disconnects)
      rs.add(membercfgobj)
                                                  add a new member to the set with extra attributes (disconnects)
                                                  add a new member which is arbiterOnly:true (disconnects)
      rs.addArb(hostportstr)
      rs.stepDown([stepdownSecs, catchUpSecs])
                                                 step down as primary (disconnects)
                                                 make a secondary sync from the given member
      rs.syncFrom(hostportstr)
      rs.freeze(secs)
                                                  make a node ineligible to become primary for the time specified
      rs.remove(hostportstr)
                                                  remove a host from the replica set (disconnects)
      rs.slaveOk()
                                                  allow queries on secondary nodes
                                                  check oplog size and time range
      rs.printReplicationInfo()
      rs.printSlaveReplicationInfo()
                                                  check replica set members and replication lag
      db.isMaster()
                                                  check who is primary
      reconfiguration helpers disconnect from the database so the shell will display
      an error, even if the command succeeds.
3/24/2020 5:24 AM Application
```

Freezing a member

 rs.freeze() method prevents the current member from seeking election as primary for a period of time specified in seconds.

```
Command Prompt - mongo --port 27021
indu:SECONDARY> rs.freeze(60)
        "ok" : 1,
        "$clusterTime" : {
                "clusterTime" : Timestamp(1591082665, 1),
                "signature" : {
                         "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAAAAAAAAA.").
                         "keyId" : NumberLong(0)
        "operationTime" : Timestamp(1591082665, 1)
indu:SECONDARY>
indu:SECONDARY> _
```



Stepping Down a Member

• The method runs only if the current member is a primary node and produce an error if runs on a non-primary member. rs.stepDown() instructs the current Primary node of the replica set to become a secondary which forces an election. After the primary steps down, eligible secondary nodes hold an election for primary.

Syntax

rs.stepDown(stepDownPeriod, secondaryCatchUpPeriod)



Stepping Down a Member

where,

• **stepDownPeriod-** Specifies the number of seconds to step down the primary, during this time the stepdown member is ineligible for becoming primary. By default it is 60 seconds. It must be greater than the secondaryCatchUpPeriod.



Stepping Down a Member

- **secondaryCatchUpPeriod** specifies the number of seconds that mongod will wait for an electable secondary to catch up to the primary. This is Optional. The default wait time is 10 seconds.
- rs.stepDown() method will not immediately step down the primary. If
 no electable secondary node is up to date with the primary, the primary
 waits up to secondaryCatchUpPeriod for a secondary node to catch up.
 Once an electable secondary is available, the method steps down the
 primary. Once stepped down, the original primary becomes a secondary
 and is ineligible from becoming primary again for the remainder of time
 specified by stepDownPeriod.



```
Command Prompt - mongo --port 27020
indu:PRIMARY>
indu:PRIMARY> rs.stepDown(90,20)
        "ok" : 1,
        "$clusterTime" : {
                "clusterTime" : Timestamp(1591082825, 1),
                "signature" : {
                         "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAAAAAAAAAA."),
                         "keyId" : NumberLong(0)
        "operationTime" : Timestamp(1591082825, 1)
indu:SECONDARY> _
```



Read Operation on Secondary Member

• rs.secondaryOk() method allows the current connection to allow read operations to run on *secondary* members. i.e. allow users to run commands in the secondary or arbiter node.

• Syntax

– rs.secondaryOk()



Read Operation on Secondary Member

```
rset:SECONDARY> rs.secondaryOk()
 rset:SECONDARY> show dbs
 admin
         0.000GB
config 0.000GB
₩local
         0.000GB
 student 0.000GB
 rset:SECONDARY> use student
 switched to db student
 rset:SECONDARY> db.school.find()
 { " id" : ObjectId("60d174b5a2e8fa78069d751b"), "name" : "sarita", "rollno" : 1, "marks" : 45 }
 { "_id" : ObjectId("60d17528a2e8fa78069d751c"), "name" : "geeta", "rollno" : 2, "marks" : 56 }
 rset:SECONDARY>
                                                                                           ♣ 31°C Light rain ヘ 및 切) ENG
                                    計
    Type here to search
```



Printing the oplog of the replica set

 rs.printReplicationInfo() method prints the report of the replica set member's oplog from the perspective of the Primary Member.

• Syntax

– rs.printReplicationInfo()



rs.printReplicationInfo()

```
Command Prompt - mongo --port 27022
indu:PRIMARY>
indu:PRIMARY> rs.printReplicationInfo()
configured oplog size: 4715.123436927795MB
log length start to end: 320secs (0.09hrs)
oplog first event time: Tue Jun 02 2020 13:05:11 GMT+0530 (India Standard Time)
oplog last event time: Tue Jun 02 2020 13:10:31 GMT+0530 (India Standard Time)
                         Tue Jun 02 2020 13:10:38 GMT+0530 (India Standard Time)
indu:PRIMARY>
```



rs.printReplicationInfo()

```
indu:SECONDARY> rs.printReplicationInfo()
configured oplog size: 4730.51523399353MB
log length start to end: 1494secs (0.42hrs)
oplog first event time: Tue Jun 02 2020 12:46:13 GMT+0530 (India Standard Time)
oplog last event time: Tue Jun 02 2020 13:11:07 GMT+0530 (India Standard Time)
now: Tue Jun 02 2020 13:11:08 GMT+0530 (India Standard Time)
indu:PRIMARY> _
```



rs.printSlaveReplicationInfo()

 rs.printSlaveReplicationInfo() method prints the report of the replica set member's oplog from the perspective of the Secondary Members of the Set.

Syntax

– rs.printSlaveReplicationInfo()



Stopping a Member

 db.shutdownServer() method is used to shut down or Stopping a member. It may be a secondary member. This command will run only with admin database.

Syntax

– db.shutdownServer()



Stopping a Member

```
indu:SECONDARY> db.shutdownServer() shutdown only works with the admin database; try 'use admin' indu:SECONDARY> use admin switched to db admin indu:SECONDARY> db.shutdownServer() 2020-06-02T13:17:04.383+0530 I NETWORK [js] DBClientConnection failed to receive message from 127.0.0.1:27020 - HostUn reachable: Connection reset by peer server should be down... 2020-06-02T13:17:04.441+0530 I NETWORK [js] trying reconnect to 127.0.0.1:27020 failed 2020-06-02T13:17:05.443+0530 I NETWORK [js] reconnect 127.0.0.1:27020 failed failed > ■
```



choosing the right type of database for your enterprise

 There are many databases to choose from but selecting the right kind of database is very essential. All have their strengths and weaknesses



1. What's the structure of the data?

- The data's structure will dictate how it's stored and retrieved. As most apps deal with data in a variety of formats, the process of selecting the database should include the appropriate data structures for storing and retrieving the data.
- If you fail to do this, your mobile application will be slow to retrieve the data from the database. Furthermore, it will also require more development time to work around the data issues that will certainly come up.



- Other data related elements that can play an important role are as follows:
 - Accessibility to the data
 - Size of the data you wish to store
 - Scope of multiple databases
 - Speed and scalability



2. Will you require a flexible data modeling solution?

- The flexibility of data modeling will dictate if you can appropriately and effectively modify your data model requirements for your mobile app.
- As mobile apps today evolve rapidly, model flexibly quickly becomes an important factor to consider. In this scenario, relational databases can be a good choice if you require strong data consistency. The same is true if the data will be highly relational.
- However, if the requirements are relaxed, NoSQL databases are the way forward as they offer enhanced flexibility.



3. What client platforms does it support?

- Are you planning on supporting iOS, Android, or both? Are looking to go beyond and support the Windows Phone as well? What about IoT devices and wearables?
- Even if you plan to do support more platforms at a later date, you have to take that into consideration now
- If you plan on going down the same path, it will be important to evaluate cloud options and databases based on the required platform support during the lifecycle of the application



4. How much data security will you need?

- Whether at rest or in motion, a high level of security must be maintained consistently. If the storage is decentralized and synchronized, it's also important to enable secure access and transmission.
- As a result, you will need to look at the following:
 - Authentication
 - Data in motion
 - Data at rest
 - Read/write access



- As far as authentication goes, it has to be flexible enough to allow the following authentication providers:
 - Custom
 - Public
 - Standard



- At the same time, anonymous access is also important, so for the data to both rest on the server and the client, you will need to be able to support both file system encryption and datalevel encryption.
- For the data that's in motion, all communication should be conducted on secure channels like TLS or SSL. Furthermore, for data read/write access, the database will need to provide granular control over what information can be accessed and modified by users.



5. How will it resolve data conflicts?

- Mobile platforms or platforms that generally use decentralized data writes can quickly experience conflicts as the same data can be simultaneously modified by multiple devices. As a result, you will need a robust support mechanism to resolve conflicts.
- What's important here is the flexibility of the conflict resolution mechanism. The one you choose should be able to seamlessly enable conflict resolution on the device, by a human, by an external system, in the cloud, automatically.



6. What are your partition requirements?

- To meet your partition requirements, you will need configurable synch topology support (like star). This will enable certain parts to function offline.
- Star topology is quite common because it enables devices to connect to a central hub utilizing a point-to-point connection (which also allows devices to function offline). There are other common topologies like tree and mesh which allow different parts of the system and devices to operate offline



7. Do you want to build or buy your synch capabilities?

 When it comes to adding synch to your mobile app, you will need to make a decision on whether you want to build your own synch solution or buy it from an established provider. For most apps, you're better off buying as building your own synch will be extremely difficult and expensive (because you have to deal with complexities associated with distributed computing).



- What's important here is to find a solution that's highly flexible.
 However, if you're determined to build your own synch, then you will
 need to be prepared to expand a significant amount of time and
 money to achieve it.
- When evaluating mobile synch and storage providers, be sure to go over everything listed above as it will be critical to developing a flexible, secure, manageable, and dependable mobile application.



Database Era

- Collecting and storing data is a very old ancient concept. Even after the evolution of computers and computing devices, data has been stored in various ways.
- Earlier databases had rather limited functionalities starting with plain text files called as **flat databases**. This means the data has to be of a textual format. Every new field is marked with a user defined delimiter a special character, a comma or a colon etc.
- As there are no relations between the fields, a flat database is hard to search and navigate. Still, this works well for a small amount of data that only needs to be read and or slight modification. An example of such databases can be CSV (Comma Separated Values) files.



Database Era

- IBM introduced hierarchical databases 60's in which the records are connected by a tree structure, based on parent-child relationships. One item can only have one parent, while one parent can have multiple children.
- This can be said as the first step toward relational databases. It
 has its own disadvantages and does not work well for all types
 of records. Network databases come after this and have a
 better organization of data. These have a tree structure, but
 children node may have multiple parents as well. Later the era
 of relational database has arrived followed by non-relational
 (NoSQL) databases.



Taking Decisions on Databases

- If compliance of ACID (Atomicity, Consistency, integrity and Durability) properties is essential, RDBMS is the best solution. For example, databases for financial transactions.
- If your input data is particularly heterogeneous and difficult to encapsulate according to a normalization schema, consider using a NoSQL DBMS. For example, data from Social Media.
- If your goal is to scale database vertically, a RDBMS is suggested and if you want to scale horizontally, a NoSQL DBMS may be preferred.
- If you have a massively distributed system and can settle for eventual consistency on some nodes/partitions, you might consider a wide column store



Normalized(SQL) Vs De-normalized Data(NoSQL)

 Normalization is a database business process to break up data into the smallest possible parts. Instead of storing first and last name in one bucket, or field, normalization requires that you store the first name separately from the last name.



Normalized(SQL) Vs De-normalized Data(NoSQL)

- This is helpful if you want to sort the data by last name or by first name. Relational databases are the most common database systems and require that data is **normalized**. These data is stored in columns and rows, which in turn make up tables like a spreadsheet.
- A set of tables makes up a schema. A number of schemas create a database. Many databases can be created on a single server. RDBMS system feature much better performance for managing data over desktop database programs. RDBMS allow multiple users (even thousands!) to work with the data at the same time, creating advanced security for access to the data.



Normalized(SQL) Vs De-normalized Data(NoSQL)

- Databases like NoSQL and object-oriented do not follow the table/row/column approach of RDBMS. Instead, they build bookshelves of elements and allow access per bookshelf.
- So, instead of tracking individual words in books, these databases narrow down the data being searched for by pointing to the bookshelf, then a mechanical assistant works with the books to identify the exact word you are looking for.
- NoSQL specifically attempts to simplify bookshelves by storing data in a denormalized way; this means storing it in large chunks



Contents to be covered

- Strengths and weaknesses of databases
- Free and open source database



 There are many databases to choose from but selecting the right kind of database is very essential. All databases are not equal and each of them has specific strengths and weaknesses. Let's have a view of some of the databases:

Relational database management systems

– It is over 4 decades when RDBMS were first come into the real world to handle the increasing flood of data being produced. And since then, RDBMS have a solid foundational theory and have influenced nearly every database system in use today.



– RDBMS store normalized data in form of data sets called relations i.e. tables with rows and columns where all information is stored as a value of a specific cell. SQL is used to manage data in an RDBMS system. SQL is standardized and provides a level of predictability and utility.

Strengths

- Relational databases have expertise in handling highly structured data and provide support for ACID (Atomicity, Consistency, Isolation, and Durability) transactions.
- The structure can be scaled up quickly because adding data without modifying existing data is simple.
- Data can be easily stored, modified and retrieved using SQL queries.



 Access rights can be assigned to certain users, like to only view the data, modify the data / data structure or full access etc. Due to these, relational databases are well-suited to applications that require tiered access.

Weaknesses

– RDBMS are expensive to set up and grow. Horizontal scaling, or scaling by adding more servers, is usually both faster and more economical than vertical scaling, which involves adding more resources to a server is supported by RDBMS. Further, the structure of relational databases complicates the process.



– RDBMS are good as they are handling structured data, but difficult to handle unstructured data. Representing real world entities in context is difficult in the bounds of an RDBMS. Now a day, a lot of unstructured data is playing in the real world.

- Name:aaa,designation:"manager",salary:13000,
- Name:abb,designation:"clerk",



- For better readability, the data has to be reorganized from tables into a more readable form, and this also impact the speed. The fixed schema of these databases doesn't support such changes in a positive way.
- Sharding is necessary to scale out a relational database. In sharding, data is horizontally partitioned and distributed across a collection of machines. But maintaining ACID compliance while Sharding in RDBMS is again a big issue.



Final verdict

- Based on the various parameters, strengths and weaknesses, a relational database may be used for
 - Situations where data integrity is absolutely necessary for example financial applications, banking application, military defense and security applications, private health information systems etc.
 - Applications having Highly structured data
 - Applications having automation of internal processes
 - Some example of RDBMS are Oracle, MySQL, MS SQL Server, Sybase, Informix, PostgreSQL.



NoSQL - Document store Databases

- Document-oriented databases, also known as document stores, are used to store and manage semi-structured or unstructured data. A document store is a nonrelational database that stores data in JSON, BSON, or XML documents.
- These databases supports a flexible schema i.e. the data does not adhere to a fixed structure, instead it forms its own structure. Due to the lack of a defined structure i.e. no schema of table is declared before inserting data, document stores don't enforce document structure.
- Documents can contain any data desired. Since the information of Document Stores cannot be arranged in tables, this data is not suitable for relational databases



NoSQL - Document store Databases

- A document database creates a key-value pair: A key is assigned to a specific document. The actual information is then located within this document, which may be formatted as an XML, JSON or YAML(YAML ain't markup language) file. Documents have key-value pairs but also embed attribute metadata (fields) to make querying easier.
- Since the document does not require a specific schema, different types of documents can also be integrated together in a document store. Further, any changes or modifications to the documents need not to be communicated to the database.



NoSQL - Document store Databases

Strengths

- Document stores are very flexible and can store any type of data.
 They can easily handle semi structured and unstructured data. Even large volumes of unstructured data can be accommodated in the database.
- Users don't need specify what data will be stored, so this is a good choice when it isn't clear in advance what types of data will be stored. For example: Social Media data.



- Easier to integrate new information
- Each record may have different structure, these records are called documents. Users
 can create their desired structure in a particular document without affecting all
 documents and subsequently modify it again to a new structure. Schema can be
 modified without causing downtime, which leads to high availability. Write speed is
 generally fast, as well.
- The information is not distributed over multiple linked tables. Everything is contained in a single location, and this can result in better performance.
- Document stores can be scaled horizontally easily. The sharding necessary for horizontal scaling is much more intuitive than with relational databases, so document stores scale out fast and efficiently.



NoSQL - Document store Databases

Weaknesses

- Document databases sacrifice ACID compliance for flexibility.
- Query can be done in a particular document not across multiple / different documents (not like multiple tables of RDBMS). So, not suitable for highly networked data volumes.
- No restrictions on data and any type of data can be feed into documents, this raises the potential for accidents & threats (both unintentionally or intentionally)



NoSQL - Document store Databases

- Final verdict
- Based on the various parameters, strengths and weaknesses, a Document Store database may be used for
 - Unstructured or semi structured data
 - Content management
 - In-depth data analysis
 - Some of the examples of Document Store databases are MongoDB,
 CouchDB, BaseX, Elasticsearch etc.



— A key-value store database is a type of nonrelational database that uses a simple key-value method to store data. In key-value, each value is associated with a specific key. A key-value database stores data as a collection of key-value pairs where the key serves as a unique identifier. Both keys and values can be anything, ranging from simple objects to complex compound objects. It just has a one-way mapping from the key to the value to store data.



- Student
- Name
- Rollno
- marks



- A key-value store uses an array of keys where each key is associated with only one value in a collection. It is quite similar to a dictionary or a map data structure. It's also known as an associative array.
- Key1:value1,key2:value2,key3:value3
- "key" is a unique identifier associated only with the value. Keys can be anything allowed by the DBMS.



- "Values" are stored as blobs and don't need predefined schema. They can take nearly any form: numbers, strings, counters, JSON, XML, HTML, PHP, binaries, images, short videos, lists, and even another key-value pair encapsulated in an object.

Strengths

- Are very flexible and able to handle a very wide array of data types easily.
- Have high performance as Keys are used to go straight to the value with no index searching or joins.
- High performance because the integrated caching feature allows users to store and retrieve data in the shortest time possible.



- Can be moved from one system to another without rewriting code. i.e. easily Portable
- Easy to scale without disrupting operations. Users can add and remove servers depending on their needs without causing undesirable disruptions.
- users can simply add new features when the need arises

Weaknesses

- No query language to retrieve data. Only have some simple operations such as get, put and delete.
- Almost impossible to query values, because they're stored as a blob and can only be returned as such.



- Difficult to do reporting or edit parts of values.
- Data querying (retrieving) usually handled manually at the application level.

Final Verdict

- Based on the various parameters, strengths and weaknesses, a Key-value Store database may be used for
 - Session management at high scale
 - User preference, profiles and settings
 - Unstructured data such as product reviews or blog comments
 - Product recommendations; latest items viewed on a retailer website drive future customer product recommendations



- Ad servicing; customer shopping habits result in customized ads, coupons, etc. for each customer in real-time
- Can effectively work as a cache for heavily accessed but rarely updated data i.e. Data that will be accessed frequently but not often updated
- Some of the examples of Key-Value Store databases are Amazon DynamoDB, MemcacheDB, InfnityDB, Redis, Aerospike, Oracle Berkeley DB, Riak KV, Voldemort etc



Strengths and weaknesses of databases NoSQL-Wide-column store

– Wide-column stores are dynamic column-oriented nonrelational databases. Wide column stores are database management systems that organize related facts into columns. Groups of these columns, called "column families", have multiple rows and rows may not have same number of columns.



Name	Details				
Aaa	Phone number	Eng	Hindi	Science	
	56456456456	52	62	52	
Bbb	address	Eng	Hindi	Science	Maths
	House no.23,bldg. no.34	85	95	95	45



Strengths and weaknesses of databases NoSQL-Wide-column store

- Wide-column stores use the concept of a keyspace instead of schemas. A keyspace encompasses column families (similar to tables but more flexible in structure), each of which contains multiple rows with distinct columns.
- Each row doesn't need to have the same number or type of column.
 The columns name as well as record keys are not fixed in these databases.



Student			Employee			Account			
Id:1 Name:aaa		ld:1	Name:ppp		Salary:333	ld:78	Designation:"asdsda"		
ld:3	Name:bbb Marks:34		ld:4		Salary:555		ld:56	Designatio n:clerk	Bonus:400



Strengths and weaknesses of databases NoSQL-Wide-column store

Strengths

This type of database has some benefits of both relational and nonrelational databases. It deals better with both structured and semi structured data than other nonrelational databases, and it's easier to update. Columnar databases compress better than rowbased systems. Also, large data sets are simple to explore.



NoSQL- Columnar databases

- Columnar databases are highly scalable horizontally and faster at scale because they store data in columns rather than rows. The columns are easy to scale so they can store large volumes of data. This feature also allows users to spread data across many computing nodes and data stores.
- Column stores are very efficient at data compression. Columnar databases are highly compressed compared to conventional relational databases that store data by row. They allow users to optimize storage size.
- Due to their structure, columnar databases perform particularly well with aggregation queries such as SUM, COUNT, AVG, etc.



NoSQL Columnar databases

- Database users can sort and manipulate data directly from a columnar database and there is no need to rely on the application.
- Columnar stores can be loaded extremely fast. A billion row table could be loaded within a few seconds. You can start analyzing almost immediately.

Weaknesses

- Wide-column stores are slower than relational databases when handling transactions.
- Good for bulk write and updates but it's a costly affair in lower quantity like individual records.
- Not suitable for Incremental data loading & Queries against only a few rows



NoSQL Columnar databases

- Final Verdict
- Based on the various parameters, strengths and weaknesses, a wide-column Store database may be used for
 - Big data analytics where speed is important
 - Reporting systems
 - Time Series Data



NoSQL Columnar databases

- Data warehousing on big data
- Extreme write speeds with relatively less velocity reads like Logging
- Sensor Logs [Internet of Things (IOT)]
- User preferences
- Geographic information
- Some of the examples of wide-column Store databases are Cassandra, HBase, Microsoft Azure Cosmos, Druid, Hypertable etc.



NoSQL-Search engine

- It may seem strange to include search engines in an article about database types. However, Elasticsearch has seen increased popularity in this sphere as developers look for innovative ways to cut down search lag.
- Elastisearch is a nonrelational, document-based data storage and retrieval solution specifically arranged and optimized for the storage and rapid retrieval of data.



NoSQL-Search engine

Strengths

- Elastisearch is very scalable. It features flexible schema and fast retrieval of records, with advanced search options including full text search, suggestions, and complex search expressions.
- One of the most interesting search features is stemming. Stemming analyzes the root form of a word to find relevant records even when another form is used. For example, a user searching an employment database for "paying jobs" would also find positions tagged as "paid" and "pay."



NoSQL-Search engine

Weaknesses

- Elastisearch is used more as an intermediary or supplementary store than a primary database. It has low durability and poor security. Also, Elastisearch doesn't support transactions.
- Use a search engine like Elastisearch for:
 - Improving user experience with faster search results
 - Logging



A few open source database solutions available in the market

MySQL

- MySQL has been around since 1995 and is now owned by Oracle. Apart from its open source version, there are also different paid editions available that offer some additional features, like automatic scaling and cluster geo-replication.
- We know that MySQL is an industry standard now, as it's compatible with just about every operating system and is written in both C and C++.
- This database solution is a great option for different international users, as the server can provide different error messages to clients in multiple languages, encompassing support for several different character sets.

MySQL

Pros

- It can be used even when there is no network available.
- It has a flexible privilege and password system.
- It uses host-based verification.
- It has security encryption for all the password traffic.
- It consists of libraries that can be embedded into different standalone applications.
- It provides the server as a separate program for a client/server networked environment.

MySQL

Cons

- Different members are unable to fix bugs and craft patches.
- Users feel that MySQL no longer falls under the category of a free OS.
- It's no longer community driven.
- It lags behind others due to its slow updates.



SQLite

SQLite

— SQLite is supposedly one of the most widely deployed databases in the world. It was developed in 2000 and, since then, it has been used by companies like Facebook, Apple, Microsoft and Google. Each of its releases is carefully tested in order to ensure reliability. Even if there are any bugs, the developers of SQLite are quite honest about the potential shortcomings by providing bug lists and the chronologies of different code changes for every release.

Pros

It has no separate server process.



SQLite

- The file format used is cross-platform.
- It has a compact library, which runs faster even with more memory.
- All its transactions are ACID compliant.
- Professional support is also available for this database.

SQLite

- Cons
- It's not recommended for:
 - Different client/server applications.
 - All high-volume websites.
 - High concurrency.
 - Large datasets.



MongoDB

- MongoDB was developed in 2007 and is well-known as the 'database for giant ideas.' It was developed by the people behind ShopWiki, DoubleClick, and Gilt Group. MongoDB is also backed by a large group of popular investors such as The Goldman Sachs Group Inc., Fidelity Investments, and Intel Capital.
- Since its inception, MongoDB has been downloaded over 15 million times and is supported by more than 1,000 partners. All its partners are dedicated to keeping this free and open source solution's code and database simple and natural.

MongoDB

Pros

- It has an encrypted storage engine.
- It enables validation of documents.
- Common use cases are mobile apps, catalogues, etc.
- It has real-time apps with an in-memory storage engine (beta).
- It reduces the time between primary failure and recovery.



MongoDB

Cons

- It doesn't fit applications which need complex transactions.
- It's a young solution—its software changes and evolves quickly



MariaDB

- MariaDB has been developed by the original developers of MySQL. It is widely used by tech giants like Facebook, Wikipedia and even Google. It's a database server that offers drop-in replacement functionality for MySQL.
- Security is one of the topmost concerns and priorities for MariaDB developers and in each of its releases, the developers also merge in all of MySQL's security patches, even enhancing them if required.



MariaDB

Pros

- It has high scalability with easier integration.
- It provides real-time access to data.
- It has the maximum core functionalities of MySQL (MariaDB is an alternative for MySQL).
- It has alternate storage engines, patches and server optimisations.



Maria DB

Cons

Password complexity plugin is missing.

1) What is MongoDB?
A. It is a data growth
B.It is a adminCommand
C.It is a combine objects
D.It is a document database
2) MongoDB was Initial released in
A. August 2009
B.August 2008
C.August 2010
D.None of the above
3) In mongodb, data is represented as a collection of
A. Tables
B.Images
C.Files
D.None of the above
4) A collection and a document in mongodb is equivalent to which of the sql concepts respectively?
A. Column and Row
B.Table and Row
C.Table and Column
D.Database and Table
5) Which of the following statements is correct about mongoose in mongodb?
A. it is Java library to connect with MongoDB.
B.It is a PHP library to connect with MongoDB.
C.It is Python library to connect with MongoDB

D.It is used for modeling your application data in node.js
6) A collection in mongodb is a group of
A. Rows
B.Schema
C.Databases
D.Related documents
7) Please choose the correct option?
A. MongoDB is a NoSQL database
B.MongoDB is column oriented data store
C.MongoDB uses XML more in comparison with JSON
D.None of the above
8) The concatenation of the collection name and database name is called
A. Replica
B.Sharding
C.MongoDB
D.Namespace
9) What kind of database MongoDB is?
A. NoSQL Database
B.SQL Database
C.Operational database
D.None of the above
10) Does MongoDB supports query joins between collections?
A. No
B.Yes

11) N	Mongodb supports search by field range queries regular expression searches.
	A. True
	B.False
12) V	Which among following is not a supported index type in mongodb?
	A. Unique
	B.TTL Index
	C.Neospatial
	D.None of the above
13) V	Which of the following format is supported by mongodb?
	A. SQL
	B.BSON
	C.XML
	D.None of the above
14) V	Which of the following language is mongodb written in?
	A. C
	B.C++
	C.Javascript
	D.All of the Above
	member is used to support dedicated functions, such as backup or rting.
	A. Primary
	B.Hidden
	C.ViewState
	D.None of the above
16) Ir	n mongodb sorting is not supported.

A. Heap	
B.Collation	
C.Collection	
D.None of the above	
17) mongodb is an example of	
A. Graph Database	
B.Key-Value pair Database	
C.Column-based Database	
D.Document-based Database	
18) MongoDB stores all documents in	
A. Rows	
B.Tables	
C.Collections	
D.None of the above	
19) Secondary indices are not available in MongoDB.	
A. True	
B.False	
20) MongoDB supports fixed-size collections called co	llections.
A. Capped	
B.Primary	
C.Secondary	
D.None of the above	
21) Which of the following is true about sharding?	
A. Sharding is enabled at the database level	

	B.We cannot change a shard key directly/automatically once it is set up
	C.Creating a sharded key automatically creates an index on the collection using that key
	D.All of the above
	longoDB has been adopted as software by a number of major ites and services.
	A. backend
	B.frontend
	C.proprietary
	D.None of the above
23) W	which of the following method is used to query documents in collections?
	A. findOne()
	B.selectOne()
	C.findOne1()
	D.None of the above
24) W	hich of the following is true about mongoDB?
	A. MongoDB is a cross-platform
	B.MongoDB provides high performance
	C.MongoDB is a document oriented database
	D.All of the above
25) W	hich of the following is a metapackage for enterprise?
	A. mongodb-enterprise
	B.mongodb-enterprise-server
	C.mongodb-enterprise-mongos
	D.None of the above

26) Mongodb stores all documents in
A. Tables
B.Files
C.Collections
D.Database
27. MongoDB Queries can return specific fields of documents which also include user-defined functions.
A. C
B. C++
C. Javascript
D. All of the mentioned
28. Dynamic schema in MongoDB makes easier for applications.
A. polymorphism
B. inheritance
C. encapsulation
D. None of the mentioned
29. With MongoDB supports a complete backup solution and full deployment monitoring.
A. AMS
B. DMS
C. MMS
D. CMS
30.MongoDB provides high with replica sets.
A. availability

B. perfori	nance	
C. scalabi	lity	
D. none o	of the mentioned	
31.MongoDl	3 scales horizontally using	for load balancing purpose.
A. Partitio	oning	
B. Shardi	ng	
C. Replica	ition	
D. None	of the mentioned	
32.MongoDl	3 supports fixed-size collections calle	d collections.
А. сар	ped	
B. sec	ondary	
C. prir	nary	
D. all o	of the mentioned	
33. Which of	the following sorting is not supporte	ed by MongoDB?
A. collect	ion	
B. heap		
C. collation	on	
D. none o	of the mentioned	
34.Which of	the following is not a NoSQL databas	se?
A. Cassan	dra	
B. Mongo	DB	
C. SQL Se	rver	
D. None	of the mentioned	
35. Which of	the following is a NoSQL Database T	ype

A. SQL
B. JSON
C. Document databases
D. All of the mentioned
36. Which of the following is a wide-column store?
A. Riak
B. Redis
C. Cassandra
D. MongoDB
37. Which of the following are the simplest NoSQL databases?
A. Wide-column
B. Document
C. Key-value
D. All of the mentioned
38. NoSQL databases is used mainly for handling large volumes of data.
A. structured
B. semi-structured
C. unstructured
D. all of the mentioned
39. What is the interactive shell for MongoDB called?
A. mongo
B. dbmong
C. mongodb
D. none of the mentioned

40. Which of the following is web-based client software for MongoDB?
A. Database Master
B. BI Studio
C. Fang of Mongo
D. Mongo3
41 provides statistics on the per-collection level.
A. mongotop
B. mongosniff
C. mongofiles
D. mongooplog
42. Which of the following network analyzer fully supports MongoDB?
A. Riakshark
B. Snort
C. Wireshark
D. Suricata
43. What is the maximum size of a MongoDB document?
A. 2 MB
B. 16 MB
C. 12 MB
D. There is no maximum size. It depends on the RAM.
44. Which of the following is a metapackage for enterprise?
A. mongodb-enterprise
B. mongodb-enterprise-mongos
C. mongodb-enterprise-server

	D. one of the mentioned
45.	Which of the following is not a part of mongodb-enterprise-tools?
	A. mongotop
1	B. mongodown
(C. mongodump
1	D. none of the mentioned
46.	MongoDB only provides Enterprise packages for Ubuntu LTS.
	A. 12.10
1	B. 12.04
(C. 13.04
	D. 13.10
47.) Which of the following is the Ubuntu package management tool?
	A. capt
	B. wat
(C. wapt
1	D. dpkg
48.	mongod process is stopped by issuing which of the following command?
	A. sudo service mongod restart
1	B. sudo service mongod start
(C. sudo service mongod stop
1	D. none of the mentioned
49.	is a diagnostic tool for inspecting BSON files.
	A. bsondump
	B. bsondumpjson

C. jsono	lump
D. all of	the mentioned
50.) mongo interfac	o looks for a database server listening on port 27017 on thee.
A. web	
B. web	host
C. locall	nost
D. all of	the mentioned
51.)	command display the list of databases.
A. displ	ay dbs
B. show	db
C. show	dbs
D. show	data
52. Which	of the following operation is used to switch to new database mydb?
A. use c	lb
B. use n	nydbs
C. use d	bs
D. use r	nydb
53. Which	of the following also returns a list of databases?
A. show	database
B. displa	ay dbs
C. show	databases
D. all of	the mentioned
54. Which	of the following method is used to query documents in collections?
A. repla	ce

E	3. shell
C	C. move
[D. find
55.	Point out the wrong statement.
A	A. A single MongoDB server typically has single databases
E	3. Database is a physical container for collections
(C. Collection is the equivalent of an RDBMS table
	D. None of the mentioned
56.1	MongoDB stores all documents in
A	A. collections
E	3. rows
(C. tables
[D. All of the mentioned
57.	Which of the following pipeline is used for aggregation in MongoDB?
Þ	A. information processing
E	3. knowledge processing
(C. data processing
[D. None of the mentioned
	In aggregation pipeline, the pipeline stage provides access to MongoDB queries.
A	A. \$match
E	3. \$batch
(C. \$catch
[D. All of the mentioned
59.	Which of the following functionality is used for aggregation framework?

A. \$proj	ectmatch
B. \$proj	ect
C. \$mat	ch
D. All of	the mentioned
60.Which	of the following is not a projection operator?
A. \$eler	nMatch
B. \$slice	
C. \$	
D. None	e of the mentioned
61 Which	of the following flag can be set by mongo shell?
A. no	oTimeout
B. Ti	me
C. Ti	meout
D. No	one of the mentioned
62 A query specification	consists of a combination of query, sort, and projection ons.
A. sh	ape
B. st	ats
C. pl	an
D. Al	I of the mentioned
63. Whi	ch of the following is wrong statement -
A. C N	MongoDB supports search by field, range queries, regular expression ches
	MongoDB can store the business subject in the minimal number of iments

	Secondary indices is not available in MongoDB
D. [©]	All of the above
64. A coll	ection and a document in MongoDB is equivalent to concepts
respectiv	ely.
A. [©]	Table and Column
B. C	Table and Row
C. C	Column and Row
D. [©]	Database and Table
65. in ho	w much time the MongDB writes are written to the journal?
A. C	100 s
в. С	60 s
C. C	1 s
D. [©]	100 ms
66. Does	MongoDB supports query joins between collections?
A. C	Yes
B. C	No
67 what i	s MongoDB?
A. C	data growth
B. C	document database
c. O	adminCommand
D. C	Combine objects

68. The concatenation of the collection name and database name is called a -				
A. C	Namespace			
B. C	MongoDB			
c. C	sharding			
D. C	replica			
69. What	is the good alternatives to MongoDB?			
A. C	Redis & CouchDB			
B. C	Cassandra			
c. C	Riak & Hbase			
D. C	All of the mentioned			
70. Which	n statements is correct about mongoose in MongoDB?			
A. C	it is Java library to connect with MongoDB			
в. С	It is used for modeling your application data in node.js			
c. C	It is Python library to connect with MongoDB			
D. C	It is a PHP library to connect with MongoDB			
71. Hidde	n member is used to support dedicated functions, such as backup or			
reporting.				
A. C	True			
в. С				
72. A coll	ection in MongoDB is a group of			

B. C Related documents C. Schema D. Rows 73. Which is not a supported index type in MongoDB? A. TTL Index B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON B. Ltxt	A. [©]	Databases				
73. Which is not a supported index type in MongoDB? A. TTL Index B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	в. С	Related documents				
73. Which is not a supported index type in MongoDB? A. TTL Index B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	c. [©]	Schema				
A. C TTL Index B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	D. C	Rows				
A. C TTL Index B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON						
B. Neospatial C. Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	73. Which	is not a supported index type in MongoDB?				
C. C Unique D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	A. C	TTL Index				
 D. None of These 74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON 	В. С	Neospatial				
74. Which of the following is correct option? A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON						
A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	D. C	None of These				
A. MongoDB uses XML more in comparison with JSON B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON						
B. MongoDB is column oriented data store C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	74. Which	of the following is correct option?				
C. MongoDB is a NoSQL database D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	A. [©]	MongoDB uses XML more in comparison with JSON				
 D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON 	В. С	MongoDB is column oriented data store				
 D. None of the above 75. Is MongoDB better than other NoSQL databases? A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON 	c. C	MongoDB is a NoSQL database				
A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON	_					
A. Yes B. No 76. In which format MongoDB represents document structure? A. BSON						
B. No 76. In which format MongoDB represents document structure? A. BSON	75. Is MongoDB better than other NoSQL databases?					
76. In which format MongoDB represents document structure? A. BSON	A. C	Yes				
A. BSON	В. С	No				
A. BSON						
	76. In which format MongoDB represents document structure?					
B. C .txt	A. C	BSON				
	в. С	.txt				

C. C .Do	сх
D. C Noi	ne of these
77. Howmany	byte counter in BSON is starting with a random value?
A. C 4	
B. C 2	
C. C 3	
D. C Noi	ne of the above
78. which field	d is always the first field in the document.?
A [©] _id	
B [©] Ob_io	d
C [©] id	
D [©] None	e of these
79. The applic client library,	ration, that communicates with application MongoDB by way of a is called
A. C Par	ent
B. C Driv	ver
C. C Rar	nk
D. C Noi	ne of the above
80	sorting is not supported by MongoDB.

A. Collection
B. Collation
C. C heap
D. $^{f C}$ none of the mentioned
81. When a relational expression is false, it has the value
A. C zero
B. C one
82. Which of the following is not a feature for NoSQL databases?A. Data can be easily held across multiple serversB. Relational DataC. ScalabilityD. Faster data access than SQL databases
83. Which of the following statement is correct with respect to mongoDB?
A. MongoDB is a NoSQL Database B. MongoDB used XML over JSON for data exchange C. MongoDB is not scalable D. All of the above
84. Which of the following is/are feature of mongoDB?
A. Vertically scalableB. Centralised databaseC. Dynamic schema is/are feature of mongoDB.D. All of the above
85. Which of the following represent column in mongoDB?

A. document B. database C. collection D. field
86. The system generated _id field is?
A. A 12 byte hexadecimal value B. A 16 byte octal value C. A 12 byte decimal value D. A 10 bytes binary value
87. Which of the following true about mongoDB?
A. MongoDB is a cross-platform B. MongoDB is a document oriented database C. MongoDB provides high performance D. All of the above
88. Collection is a group of MongoDB?
A. Database B. Document C. Field D. None of the above
89. A developer want to develop a database for LFC system where the data stored is mostly in similar manner. Which database should use?
A. Relational B. NoSQL C. Both A and B can be used D. None of the above
90. Documents in the same collection do not need to have the same set of fields or structure, and common fields in a collection's documents may hold different types of data is known as ?

A. dynamic schema B. mongod C. mongo D. Embedded Documents
91. Instead of Primary Key mongoDB use?
A. Embedded Documents B. Default key _id C. mongod D. mongo
92 are operations that process data records and return computed results. a) ReplicaAgg b) SumCalculation c) Aggregations d) None of the mentioned
93. Point out the wrong statement. a) Map-reduce cannot have a finalize stage to make final modifications to the result b) Map-reduce is less efficient and more complex than the aggregation pipeline c) Specifically, a user with the userAdmin role can grant itself any privilege in the database d) All of the mentioned
94. Running data aggregation on the instance simplifies application code and limits resource requirements. a) document

b) mongod
c) mongos
d) all of the mentioned
95. A set is a group of mongod instances that host the same data
set.
a) copy b) sorted
c) radii
d) replica
96Point out the wrong statement.
a) Replication provides redundancy and increases data availability
b) Replication allows you to recover from hardware failure and service
interruptions
c) With multiple copies of data on different database servers, replication
protects a database from the loss of a single server d) None of the mentioned
a) None of the mentioned
97. All other instances, secondaries, apply operations from the
so that they have the same data set.
a) center
b) secondary
c) primary
d) none of the mentioned
98. A replica set can have only primary.
a) One
b) Two
c) Three d) All of the mentioned
~, ···· ·· ··· ··· ··· ··· ··· ··· ··· ·

- 99. Point out the wrong statement.
- a) In all cases, you can use replication to increase read capacity
- b) Clients have the ability to send read and write operations to different servers
- c) You can also maintain copies in different data centers to increase the locality and availability of data for distributed applications
- d) None of the mentioned

100. To support replication,	the primary re	ecords all c	changes to	its data
sets in its				

- a) oplog
- b) adlog
- c) log
- d) none of the mentioned

1	d	41	Α
2	а	42	С
3	D	43	В
4	b	44	Α
5	D	45	В
6	D	46	В
7	Α	47	В
8	D	48	С
9	Α	49	Α
10	В	50	С
11	Α	51	С
12	С	52	D
13	В	53	С
14	D	54	D
15	В	55	Α
16	В	56	Α
17	D	57	С
18	С	58	А
19	В	59	b
20	Α	60	D

21	В	61	Α
22	Α	62	Α
23	Α	63	С
24	D	64	В
25	Α	65	D
26	С	66	В
27	С	67	b
28	Α	68	Α
29	С	69	D
30	Α	70	В
31	В	71	Α
32	а	72	В
33	С	73	В
34	С	74	С
35	С	75	Α
36	С	76	Α
37	С	77	С
38	С	78	Α
39	Α	79	В
40	Α	80	В
		81	Α
		82	В
		83	Α
		84	С
		85	D
		86	Α
		87	В
		88	В
		89	В
		90	Α
		91	В
		92	С
		93	Α
		94	В
		95	D
		96	D

	97	С
	98	Α
	99	Α
	100	Α

Assignment 1

- 1. Create a database employee with empno, name, designation, salary, incentive, sales, working hours
- 2. Show all the employees having designation manager
- 3. Show all the employees having salary 6000
- 4. Give incentive to employees having working hours greater than 9
- 5. Add 100 to incentive of employees doing sales greater than 1000.
- 6. Add 1000 to salary employee whose designation accountant.
- 7 create index on field empno
- 8.create compound index on fields empno, name
- 9.find sum of salaries of employee having designation clerk.
- 10.filter the employees having designation software engineer and find the minimum salary .

Assignment 2

create a database NIELIT.create collection student with rollno,name,class,section,marks. Input marks as embedded document like marks:{eng:45,hindi:67,science:90},grades .

- 1. calculate the total marks of all subjects of each student.
- 2. sort the records based on name in ascending order
- 3. filter the records based on subject English
- 4. increment the marks by 5 for students having total marks less than 150
- 5. show only two fields name, rollno

Assignment 3. Download movies.json file

Write and run the following queries:

- 1. Show the movies titled "Gladiator".
- 2. Show distinct genre values of movies.
- 3. Show the movies of "crime" or "drama" genre.
- 4. Show the list of movies directed by "Hitchcock", display only title and year and sort them by year.
- 5. Show the list of movies where "Cotillard" played.
- 6. Show the movies released between 1967 and 1995.
- 7. Show the list of the movies released between 1967 and 1995, by displaying only title, year, director's last name sorted by year.
- 8. Show the number of movies by country.
- 9. Show the number of movies by country and actor.

Assignment 4

Create a collection called 'games'. Add games to the database. Give each document the following properties:

name, genre, rating (out of 100)

- 1. Write a query that returns all the games
- 2. Write a query to find one of your games by name without using limit(). Use the findOne method.
- 3. Write a query that returns the 3 highest rated games.

- 4. Update your two favourite games to have two achievements called 'Game Master' and 'Speed Demon', each under a single key. Show two ways to do this. Do the first using update() and do the second using save().
- 5. Write a query that returns all the games that have both the 'Game Maser' and the 'Speed Demon' achievements.
- 6. Write a query that returns only games that have achievements. Not all of your games should have achievements.

Assignment 5

Download the restaurant.json file.import the json file

- 1. Write a MongoDB query to display all the documents in the collection restaurants.
- 2. Write a MongoDB query to display the fields restaurant_id, name, borough and cuisine for all the documents in the collection restaurant.
- Write a MongoDB query to display the fields restaurant_id, name, borough and cuisine, but exclude the field _id for all the documents in the collection restaurant.
- 4. Write a MongoDB query to display the fields restaurant_id, name, borough and zip code, but exclude the field _id for all the documents in the collection restaurant.
- 5. Write a MongoDB query to display all the restaurant which is in the borough Bronx.
- 6. Write a MongoDB query to display the first 5 restaurant which is in the borough Bronx.

- 7. Write a MongoDB query to display the next 5 restaurants after skipping first 5 which are in the borough Bronx.
- 8. Write a MongoDB query to find the restaurants who achieved a score more than 90.
- 9. Write a MongoDB query to find the restaurants that achieved a score, more than 80 but less than 100.
- 10. Write a MongoDB query to find the restaurants which locate in latitude value less than -95.754168.
- 11. Write a MongoDB query to find the restaurants that do not prepare any cuisine of 'American' and their grade score more than 70 and latitude less than -65.754168.
- 12. Write a MongoDB query to find the restaurants which do not prepare any cuisine of 'American' and achieved a score more than 70 and located in the longitude less than -65.754168.
- 13. Write a MongoDB query to find the restaurants which do not prepare any cuisine of 'American' and achieved a grade point 'A' not belongs to the borough Brooklyn. The document must be displayed according to the cuisine in descending order.
- 14. Write a MongoDB query to find the restaurants which belong to the borough Bronx and prepared either American or Chinese dish.
- 15. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which belong to the borough Staten Island or Queens or Bronxor Brooklyn.
- 16. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which are not belonging to the borough Staten Island or

Queens or Bronxor Brooklyn.

- 17. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which achieved a score which is not more than 10.
- 18. Write a MongoDB query to find the restaurant Id, name, borough and cuisine for those restaurants which prepared dish except 'American' and 'Chinese'
- 19. Write a MongoDB query to find the restaurant Id, name, and grades for those restaurants which achieved a grade of "A" and scored 11 on an ISODate "2014-08-11T00:00:00Z" among many of survey dates..
- 20. Write a MongoDB query to find the restaurant Id, name and grades for those restaurants where the 2nd element of grades array contains a grade of "A" and score 9 on an ISODate "2014-08-11T00:00:00Z".
- 21. Write a MongoDB query to find the restaurant Id, name, address and geographical location for those restaurants where 2nd element of coord array contains a value which is more than 42 and upto 52..
- 22. Write a MongoDB query to arrange the name of the restaurants in ascending order along with all the columns.
- 23. Write a MongoDB query to arrange the name of the restaurants in descending along with all the columns.
- 24. Write a MongoDB query to arrange the name of the cuisine in ascending order and for that same cuisine borough should be in descending order.
- 25. Write a MongoDB query to know whether all the addresses contains the street or not.